

Advanced Cryptography

(Provable Security)

Yi LIU

Security Against Chosen Plaintext Attacks

CPA security

- Our previous security definitions for encryption capture the case where a key is used to encrypt **only one** plaintext. Clearly it would be more useful to have an encryption scheme that allows **many** plaintexts to be encrypted under the same key.

Definition Let Σ be an encryption scheme. We say that Σ has security against chosen-plaintext attacks (CPA security) if $\mathcal{L}_{\text{cpa-L}}^{\Sigma} \approx \mathcal{L}_{\text{cpa-R}}^{\Sigma}$, where:

$\mathcal{L}_{\text{cpa-L}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{EAVESDROP}(m_L, m_R \in \Sigma.\mathcal{M}):$
$\frac{}{c := \Sigma.\text{Enc}(k, m_L)}$
return c

$\mathcal{L}_{\text{cpa-R}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{EAVESDROP}(m_L, m_R \in \Sigma.\mathcal{M}):$
$\frac{}{c := \Sigma.\text{Enc}(k, m_R)}$
return c

CPA security is often called “**IND-CPA**” security: indistinguishability of ciphertexts under chosen-plaintext attack

Limits of Deterministic Encryption

- For a block cipher, F corresponds to encryption, F^{-1} corresponds to decryption, and all outputs of F look **pseudorandom**. Use block cipher “as-is”, is it satisfy CPA security?
- Consider the following adversary \mathcal{A} , that tries to distinguish the $L_{\text{cpa}-*}$ libraries:

\mathcal{A}
$c_1 := \text{EAVESDROP}(\textcolor{red}{0}^{\text{blen}}, \textcolor{red}{0}^{\text{blen}})$
$c_2 := \text{EAVESDROP}(\textcolor{red}{0}^{\text{blen}}, \textcolor{red}{1}^{\text{blen}})$
return $c_1 \stackrel{?}{=} c_2$

Limits of Deterministic Encryption

- Consider the following adversary \mathcal{A} , that tries to distinguish the $\mathcal{L}_{\text{cpa-}*}$ libraries:

\mathcal{A}
$c_1 := \text{EAVESDROP}(\text{0}^{blen}, \text{0}^{blen})$
$c_2 := \text{EAVESDROP}(\text{0}^{blen}, \text{1}^{blen})$
return $c_1 \stackrel{?}{=} c_2$

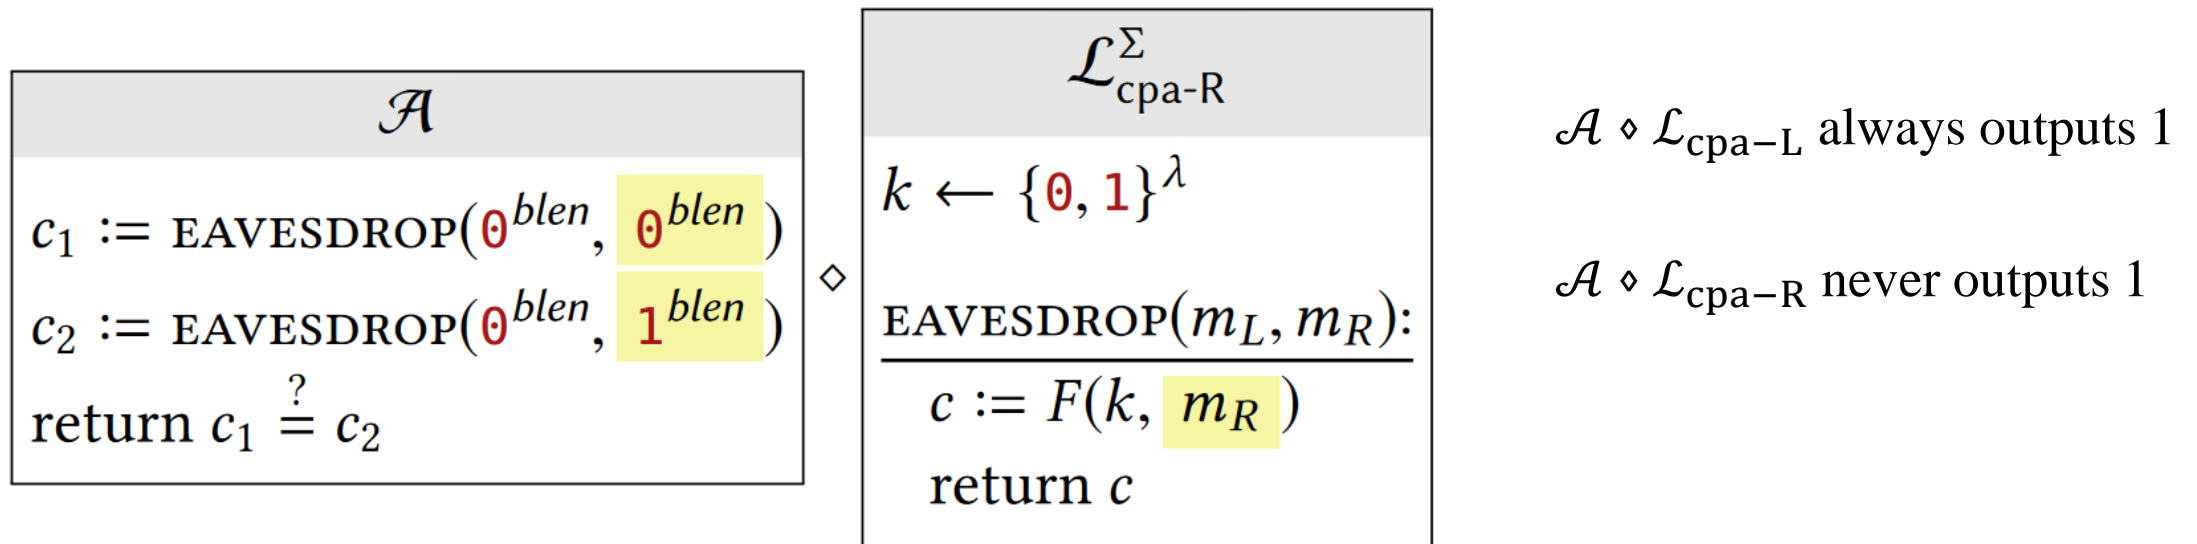
◇

$\mathcal{L}_{\text{cpa-L}}^\Sigma$
$k \leftarrow \{\text{0}, \text{1}\}^\lambda$
$\text{EAVESDROP}(m_L, m_R):$
$c := F(k, m_L)$
return c

$\mathcal{A} \diamond \mathcal{L}_{\text{cpa-L}}$ always outputs 1

Limits of Deterministic Encryption

- Consider the following adversary \mathcal{A} , that tries to distinguish the $\mathcal{L}_{\text{cpa-}*}$ libraries:



This adversary has advantage 1 in distinguishing the libraries, so the bare block cipher F is **not** a CPA-secure encryption scheme.

Limits of Deterministic Encryption

- The reason a bare block cipher does not provide CPA security is that it is **deterministic**. Calling $\text{Enc}(k, m)$ **twice** — with the **same** key and **same** plaintext — leads to the **same** ciphertext.
- Deterministic encryption can **never** be CPA-secure!
- It **leaks whether two ciphertexts encode the same plaintext**.

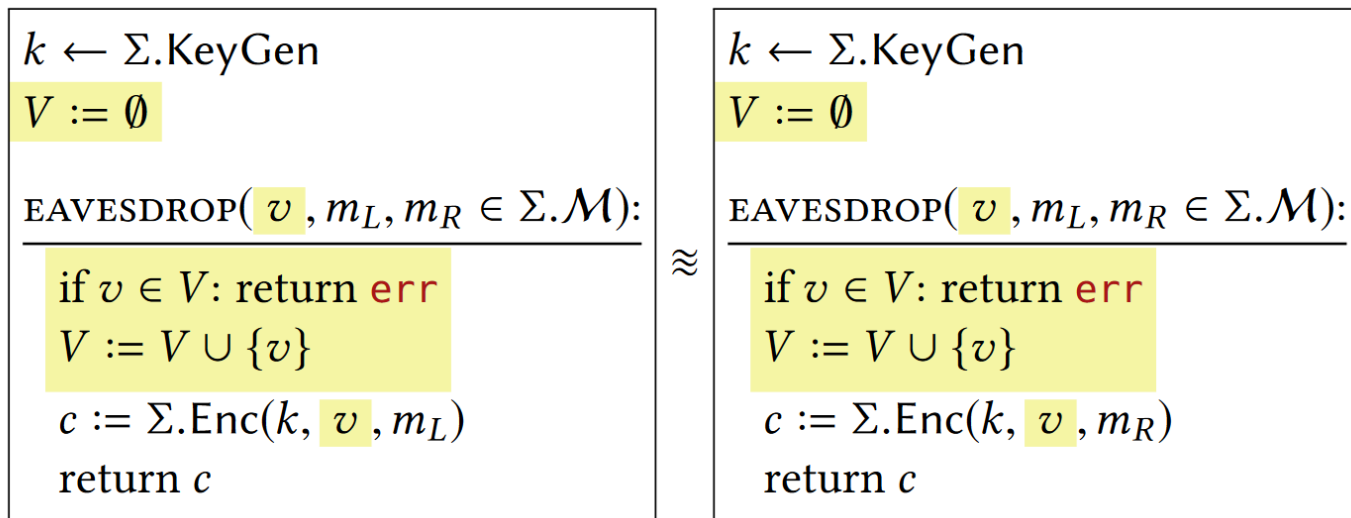
Avoiding Deterministic Encryption

We must design an Enc algorithm such that calling it **twice** with the **same** plaintext and key results in **different** ciphertexts.

- Encryption/decryption can be **stateful**, meaning that every call to Enc or Dec will actually **modify the value of k** . (symmetric ratchet construction)
- Encryption can be **randomized**. Each time a plaintext is encrypted, the Enc algorithm chooses **fresh, independent randomness** specific to that encryption.
- Encryption can be **nonce-based**. A “nonce” stands for “number used only once”.
 - A nonce **does not** need to be chosen randomly; it **does not** need to be secret; it **only** needs to be **distinct among all calls** made to Enc.

Avoiding Deterministic Encryption

- Encryption can be nonce-based. A “nonce” stands for “number used only once”.
 - A nonce only needs to be **distinct** among all calls made to Enc.
 - Nonce-based encryption **requires** a change to the interface of encryption, and therefore a change to the correctness & security definitions as well. The encryption/decryption algorithms syntax is updated to $\text{Enc}(k, v, m)$ and $\text{Dec}(k, v, c)$, where v is a nonce.
 - The correctness property is that $\text{Dec}(k, v, \text{Enc}(k, v, m)) = m$ for all k, v, m , so both encryption & decryption algorithms should use the same nonce.



Real-vs-Random Version of CPA

Definition Let Σ be an encryption scheme. We say that Σ has pseudorandom ciphertexts in the presence of chosen-plaintext attacks (CPA security) if $\mathcal{L}_{\text{cpa\$-real}}^{\Sigma} \approx \mathcal{L}_{\text{cpa\$-rand}}^{\Sigma}$, where:

$\mathcal{L}_{\text{cpa\$-real}}^{\Sigma}$	$\mathcal{L}_{\text{cpa\$-rand}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$	$\text{CTXT}(m \in \Sigma.\mathcal{M}):$
<hr/>	<hr/>
$c := \Sigma.\text{Enc}(k, m)$	$c \leftarrow \Sigma.C$
return c	return c

This definition is also called “**IND\$-CPA**”, meaning “indistinguishable from random under chosen plaintext attacks.”

Real-vs-Random Version of CPA

Definition Let Σ be an encryption scheme. We say that Σ has pseudorandom ciphertexts in the presence of chosen-plaintext attacks (CPA\$ security) if $\mathcal{L}_{\text{cpa\$-real}}^{\Sigma} \approx \mathcal{L}_{\text{cpa\$-rand}}^{\Sigma}$.

$\mathcal{L}_{\text{cpa\$-real}}^{\Sigma}$	$\mathcal{L}_{\text{cpa\$-rand}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$	
$\text{CTXT}(m \in \Sigma.\mathcal{M}):$	$\text{CTXT}(m \in \Sigma.\mathcal{M}):$
$\frac{c := \Sigma.\text{Enc}(k, m)}{\text{return } c}$	$\frac{c \leftarrow \Sigma.C}{\text{return } c}$

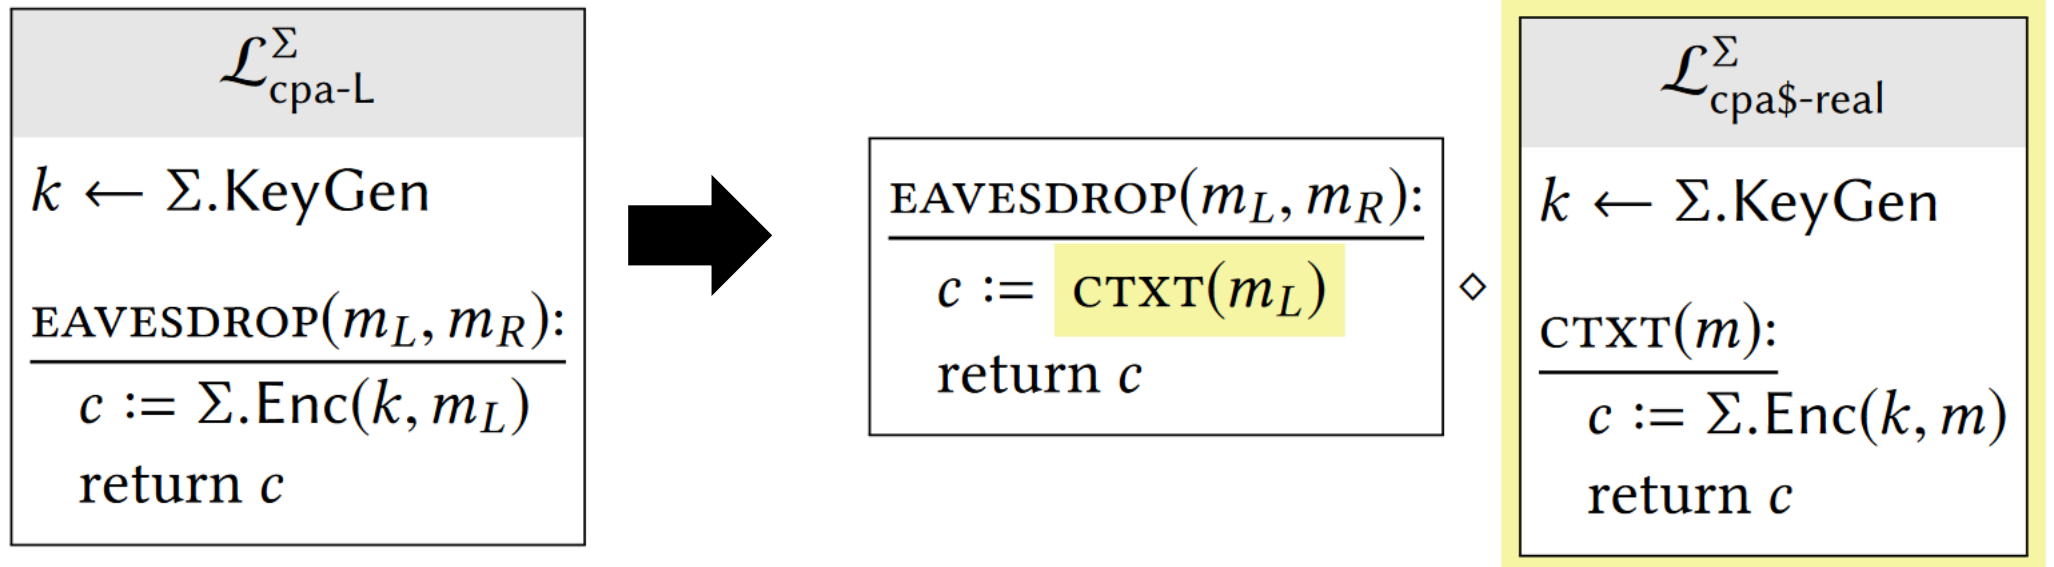
- It is **easier** to prove CPA\$ security than to prove CPA security.
- CPA\$ security **implies** CPA security.
- **Most** of the schemes we will consider achieve CPA\$ anyway.

CPA\$ security implies CPA security

Claim If an encryption scheme has CPA\$ security, then it also has CPA security.

proof

We want to prove that $\mathcal{L}_{\text{cpa-L}}^\Sigma \approx \mathcal{L}_{\text{cpa-R}}^\Sigma$, using the assumption that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$.

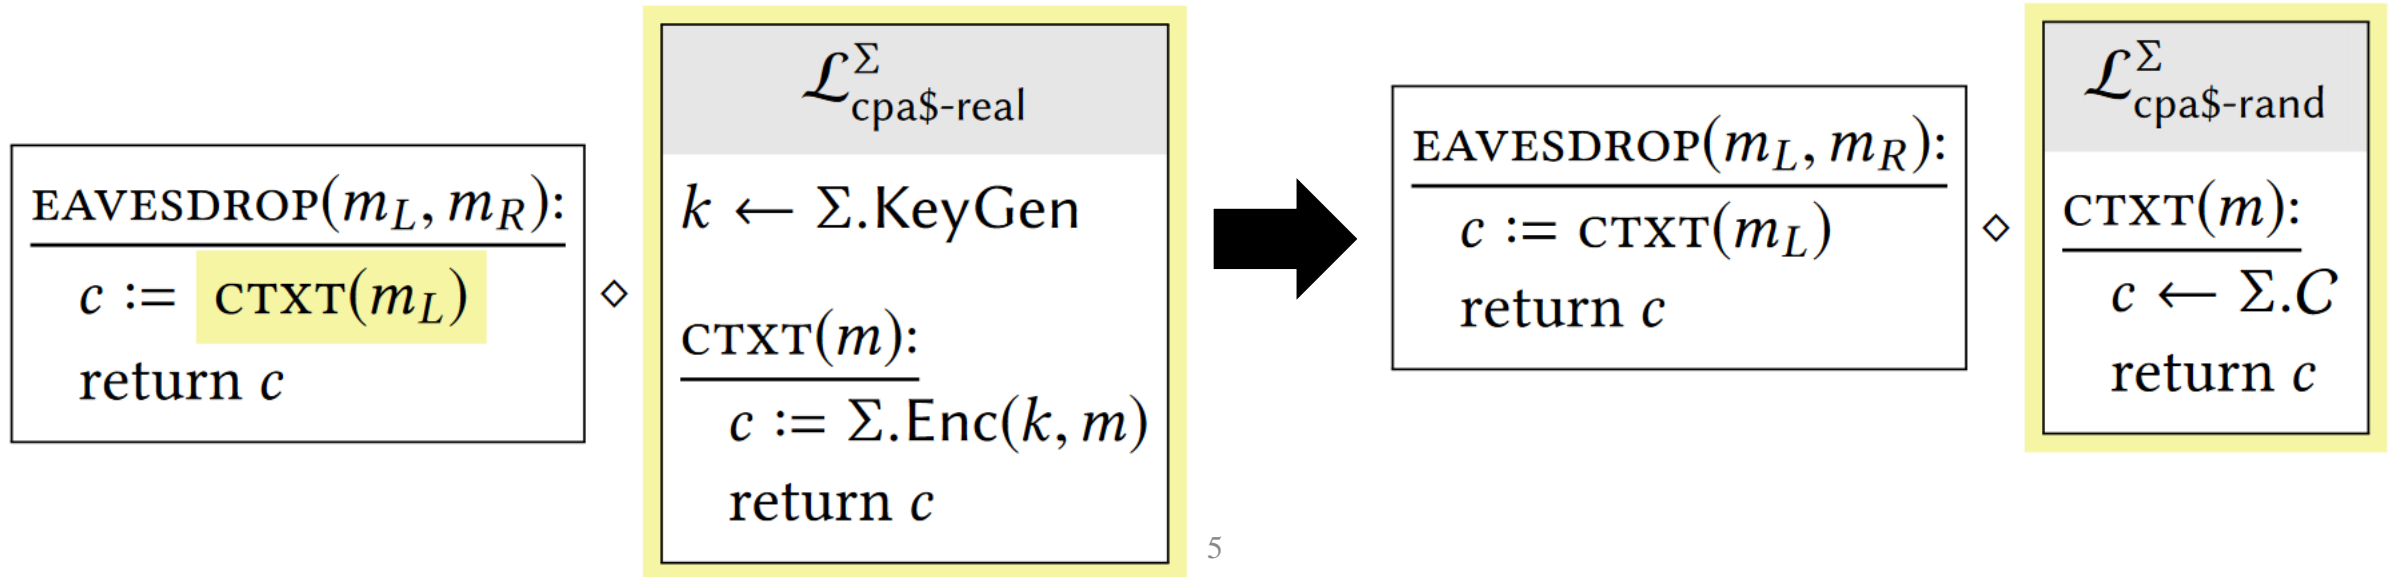


CPA\$ security implies CPA security

Claim If an encryption scheme has CPA\$ security, then it also has CPA security.

proof

We want to prove that $\mathcal{L}_{\text{cpa-L}}^\Sigma \approx \mathcal{L}_{\text{cpa-R}}^\Sigma$, using the assumption that $\mathcal{L}_{\text{cpa$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa$-rand}}^\Sigma$.




CPA\$ security implies CPA security

Claim If an encryption scheme has CPA\$ security, then it also has CPA security.

proof

We want to prove that $\mathcal{L}_{\text{cpa-L}}^\Sigma \approx \mathcal{L}_{\text{cpa-R}}^\Sigma$, using the assumption that $\mathcal{L}_{\text{cpa$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa$-rand}}^\Sigma$.

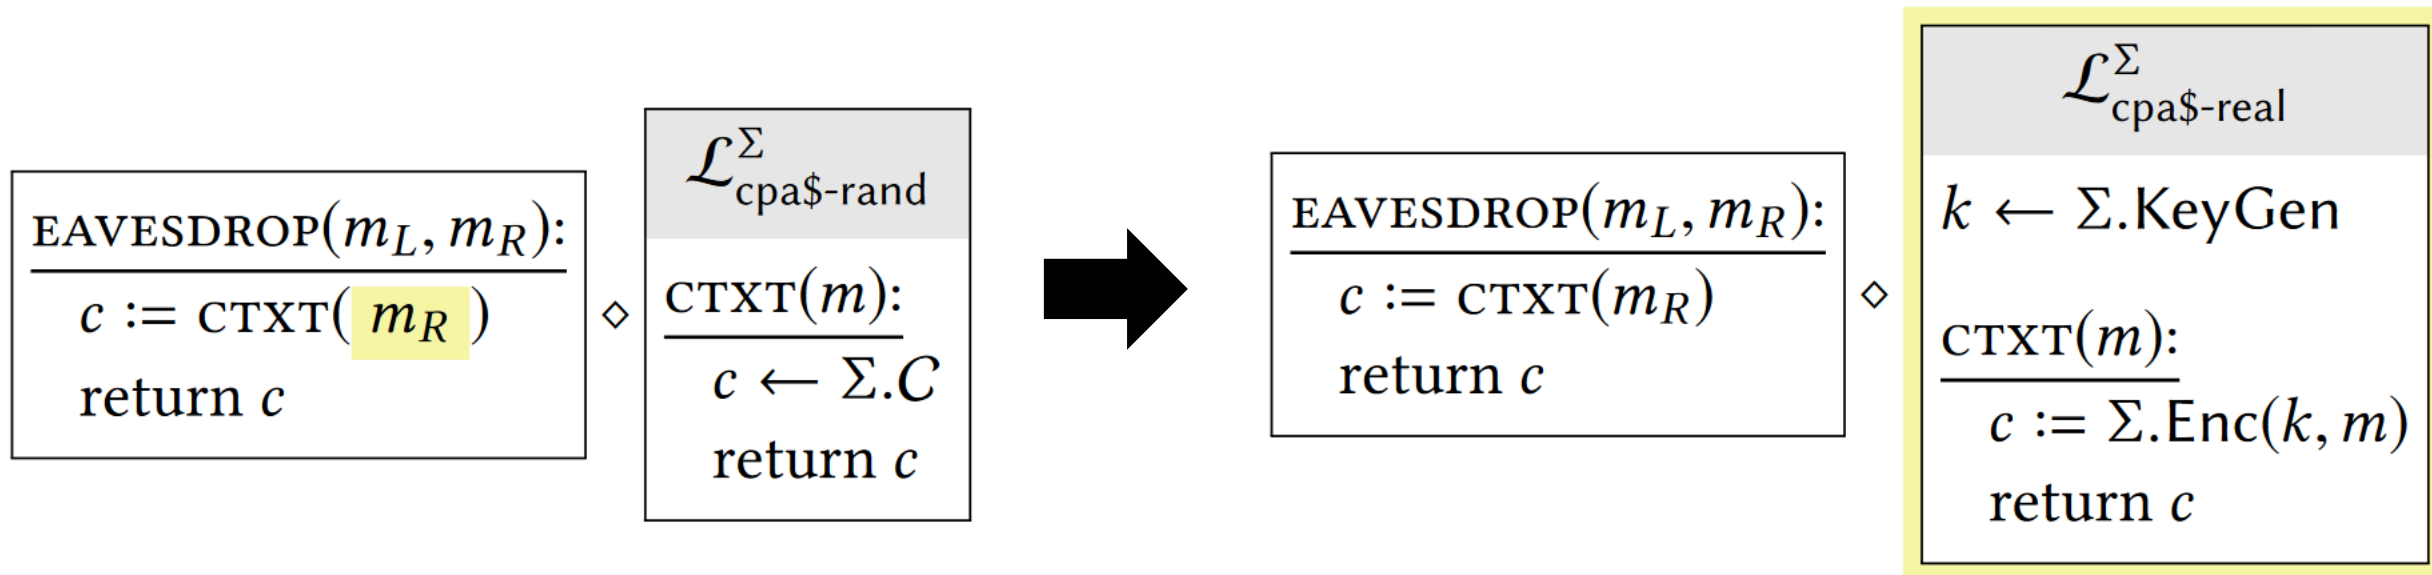


CPA\$ security implies CPA security

Claim If an encryption scheme has CPA\$ security, then it also has CPA security.

proof

We want to prove that $\mathcal{L}_{\text{cpa-L}}^\Sigma \approx \mathcal{L}_{\text{cpa-R}}^\Sigma$, using the assumption that $\mathcal{L}_{\text{cpa$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa$-rand}}^\Sigma$.

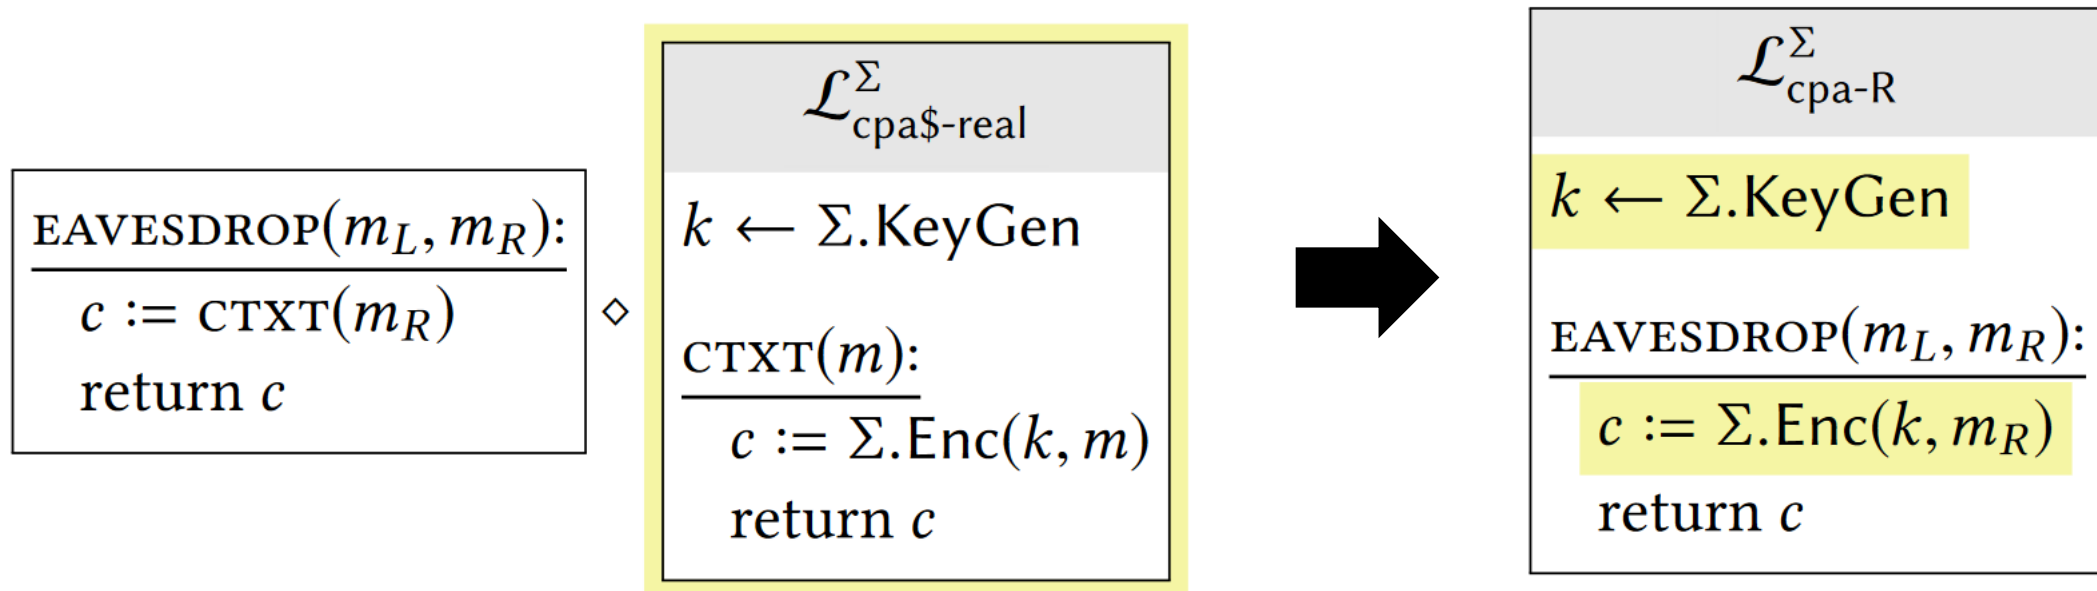


CPA\$ security implies CPA security

Claim If an encryption scheme has CPA\$ security, then it also has CPA security.

proof

We want to prove that $\mathcal{L}_{\text{cpa-L}}^\Sigma \approx \mathcal{L}_{\text{cpa-R}}^\Sigma$, using the assumption that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$.



CPA-Secure Encryption Based On PRFs

- On the one hand, we need an encryption method that is **randomized**, so that **each plaintext m is mapped to a large number of potential ciphertexts**.
- On the other hand, the decryption method must be able to **recognize all** of these various ciphertexts as being **encryptions of m** .
- If Alice and Bob share a **huge table T** initialized with uniform data, then Alice can encrypt a plaintext m to Bob by saying something like “this is encrypted with one-time pad, using key #674696273” and sending $T[674696273] \oplus m$. Seeing the number 674696273 doesn’t help the eavesdropper know what $T[674696273]$ is.
- Use PRF!

CPA-Secure Encryption Based On PRFs

Let F be a secure PRF with $in = \lambda$. Define the following encryption scheme based on F

$$\begin{aligned}\mathcal{K} &= \{0, 1\}^\lambda \\ \mathcal{M} &= \{0, 1\}^{out} \\ \mathcal{C} &= \{0, 1\}^\lambda \times \{0, 1\}^{out}\end{aligned}$$

$$\begin{array}{l} \text{KeyGen:} \\ \hline k \leftarrow \{0, 1\}^\lambda \\ \text{return } k \end{array}$$

$$\begin{array}{l} \text{Enc}(k, m): \\ \hline r \leftarrow \{0, 1\}^\lambda \\ x := F(k, r) \oplus m \\ \text{return } (r, x) \end{array}$$


$$\begin{array}{l} \text{Dec}(k, (r, x)): \\ \hline m := F(k, r) \oplus x \\ \text{return } m \end{array}$$

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

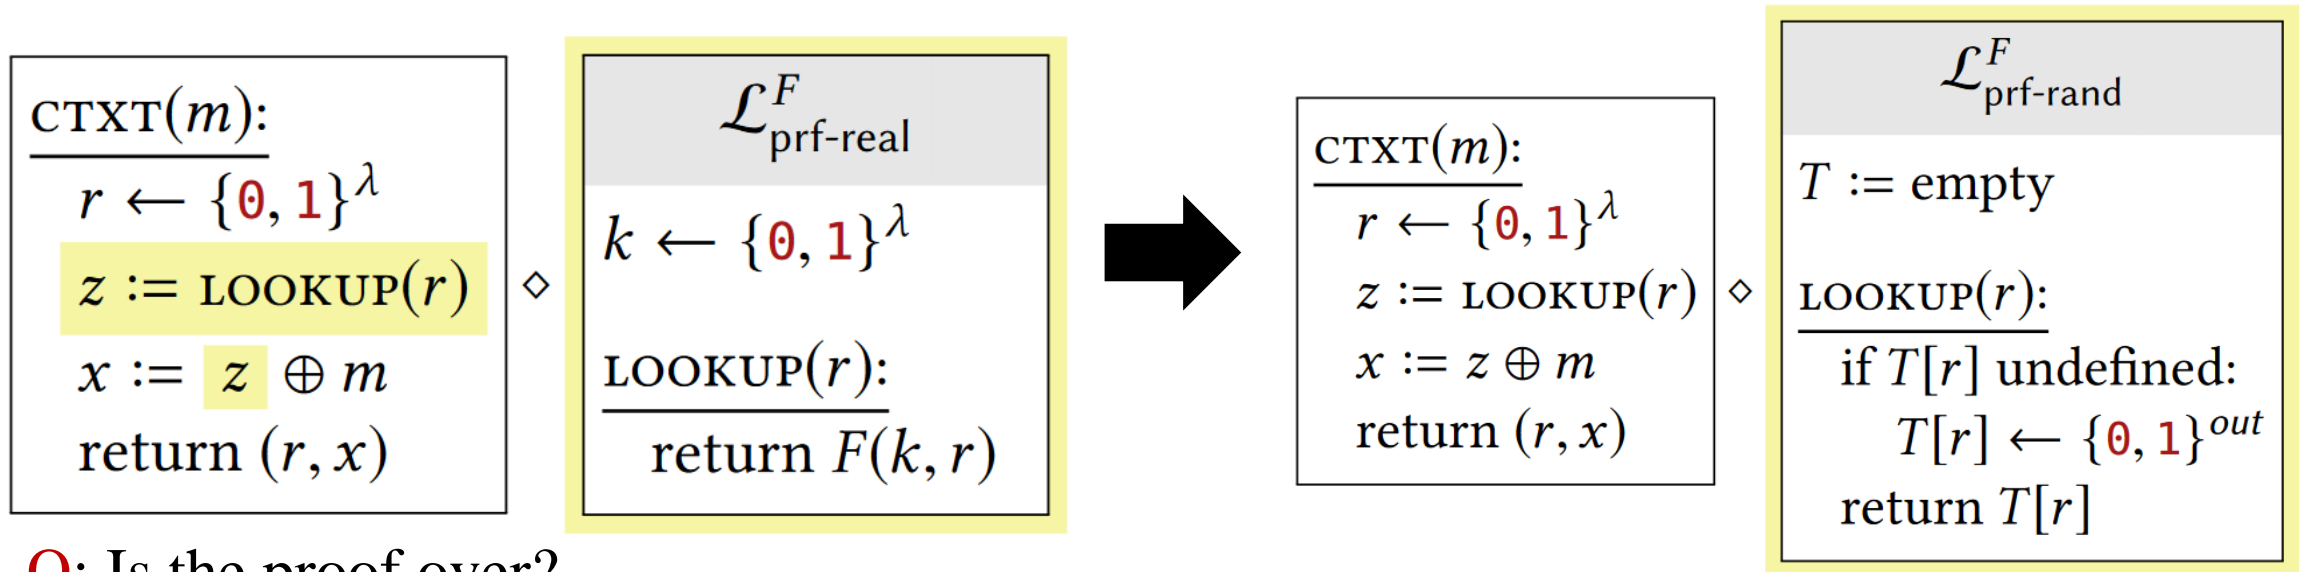
Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique




Q: Is the proof over?

r may be repeated! Our proof must explicitly contain reasoning about **why PRF inputs are unlikely to be repeated**.

CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.


Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

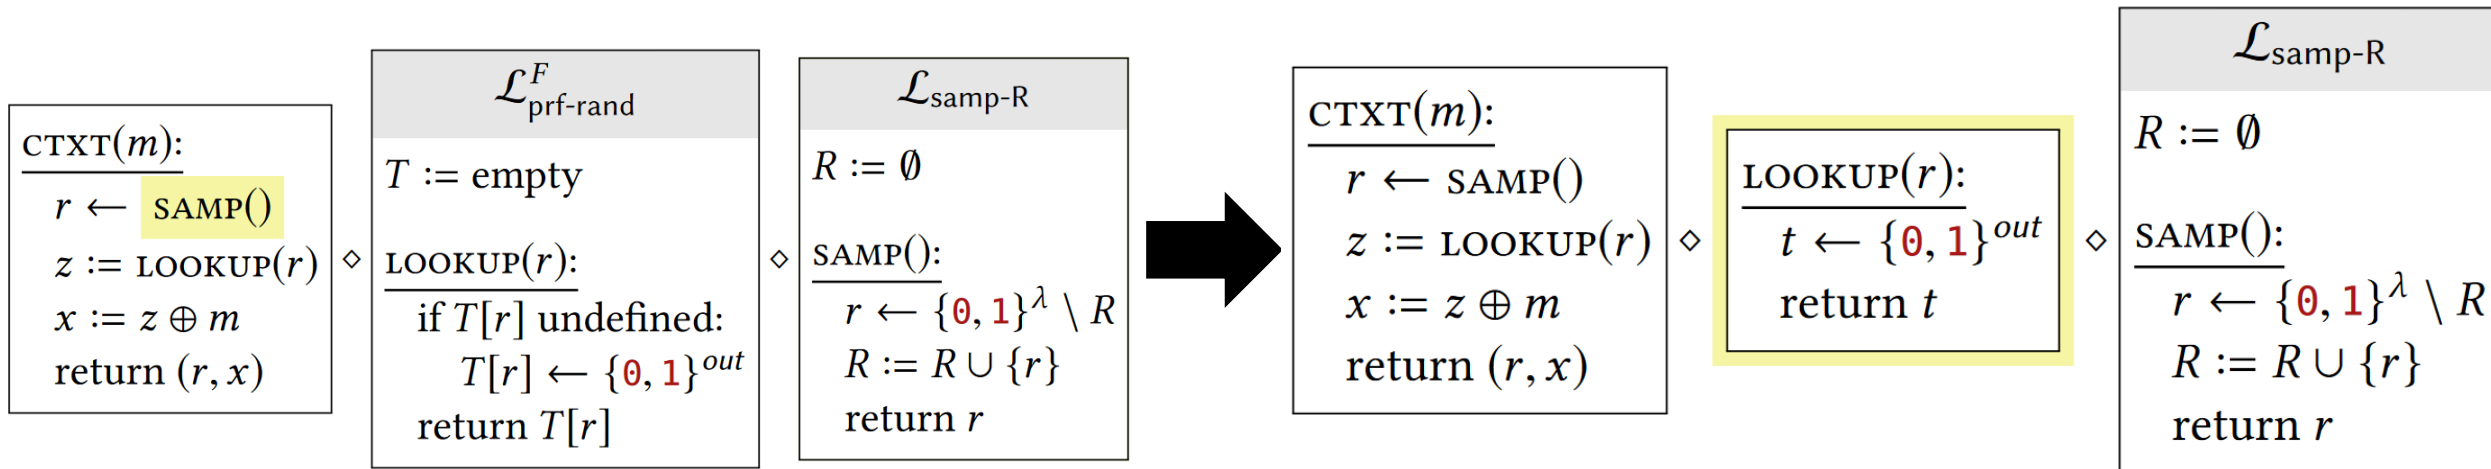
Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique

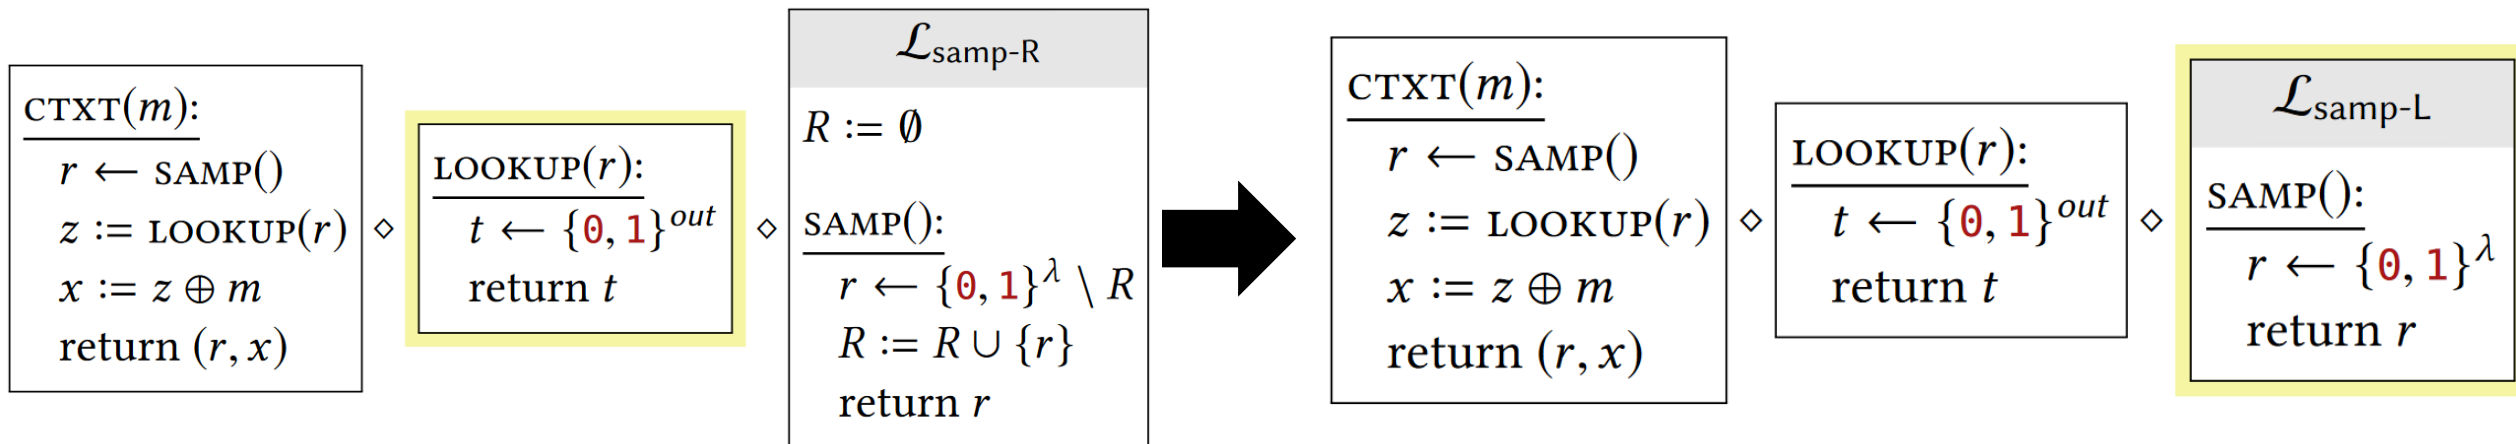


Recall that our goal is to arrive at a hybrid in which the outputs of CTXT are chosen **uniformly**. These outputs include the value r , **but** now r is no longer being chosen uniformly!

CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA security (and therefore CPA security) if F is a secure PRF.

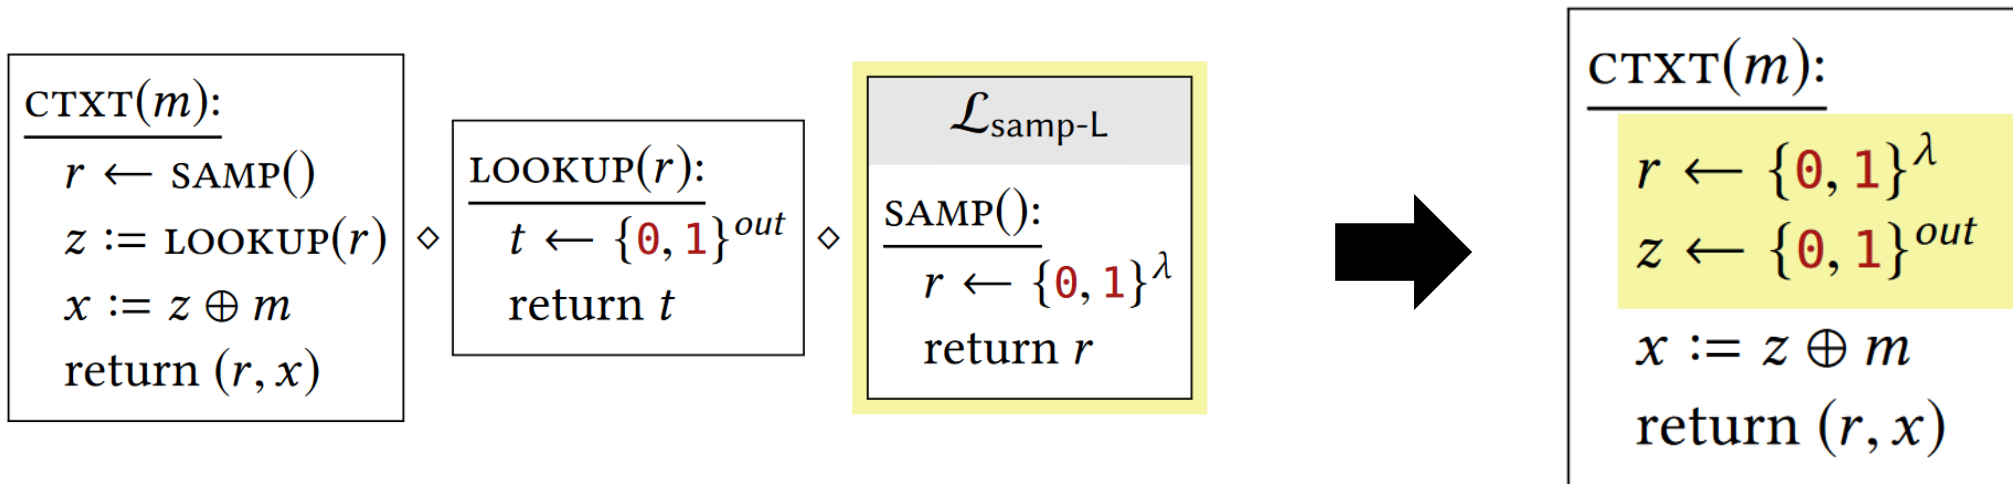
Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA security (and therefore CPA security) if F is a secure PRF.

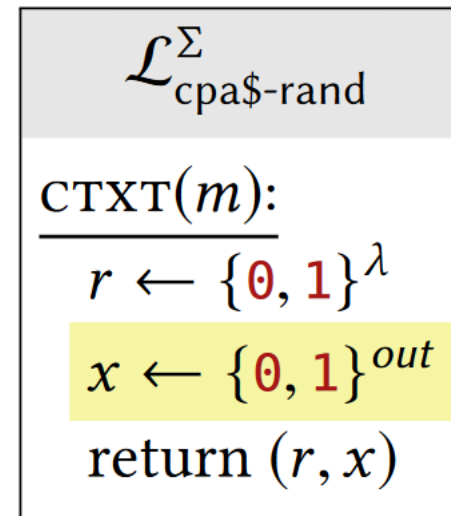


Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



CPA-Secure Encryption Based On PRFs

Claim This scheme has CPA\$ security (and therefore CPA security) if F is a secure PRF.

Proof We prove that $\mathcal{L}_{\text{cpa\$-real}}^\Sigma \approx \mathcal{L}_{\text{cpa\$-rand}}^\Sigma$ using the hybrid technique



Block Cipher Modes of Operation

Block Cipher Modes of Operation

- One of the drawbacks of the previous CPA-secure encryption scheme is that its ciphertexts are λ bits longer than its plaintexts.
- Is there any way to encrypt data (especially lots of it) without requiring such a significant overhead?
- $blen$ will denote the blocklength of a block cipher F . We'll write F_k as shorthand for $F(k, \cdot)$.
- When m is the plaintext, we will write $m = m_1 || m_2 || \cdots || m_\ell$, where each m_i is a single block

ECB: Electronic Codebook (**never never** use this! **Never!!**)

Enc($k, m_1 \parallel \dots \parallel m_\ell$):

for $i = 1$ to ℓ :

$c_i := F(k, m_i)$

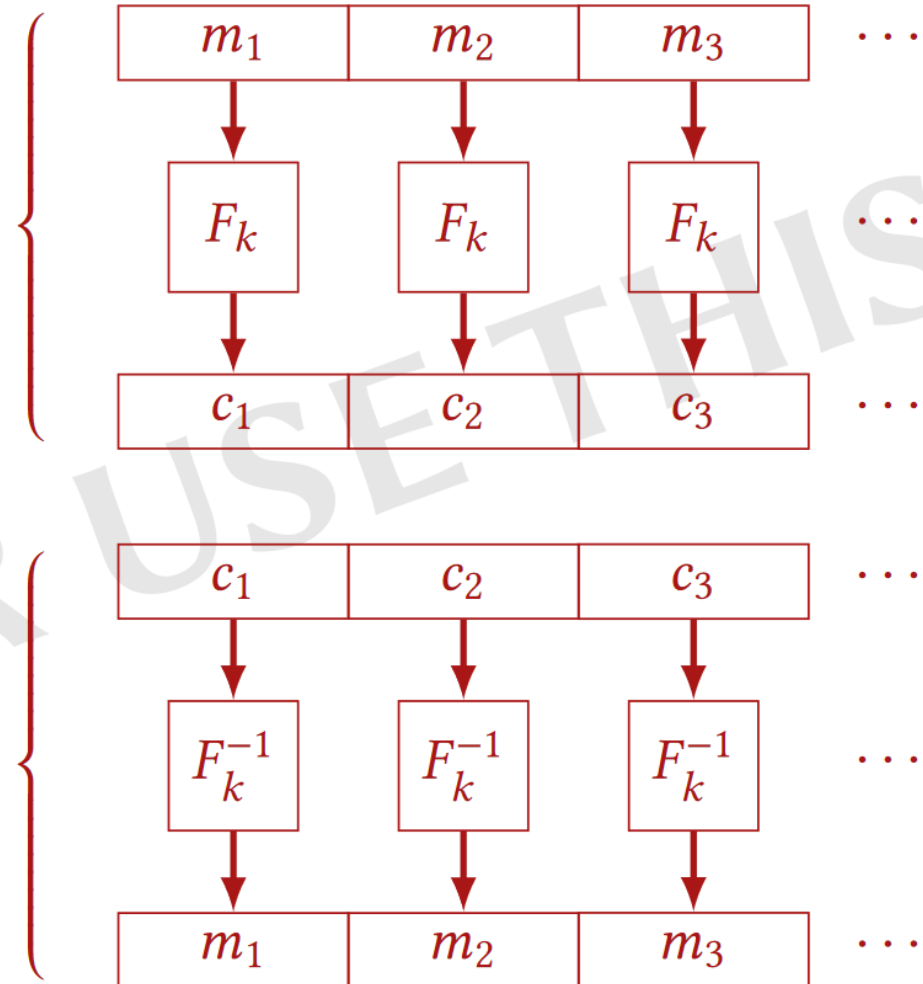
return $c_1 \parallel \dots \parallel c_\ell$

Dec($k, c_1 \parallel \dots \parallel c_\ell$):

for $i = 1$ to ℓ :

$m_i := F^{-1}(k, c_i)$

return $m_1 \parallel \dots \parallel m_\ell$



CBC: Cipher Block Chaining

Enc($k, m_1 \parallel \dots \parallel m_\ell$):

$c_0 \leftarrow \{\mathbf{0}, \mathbf{1}\}^{blen};$

for $i = 1$ to ℓ :

$c_i := F(k, m_i \oplus c_{i-1})$

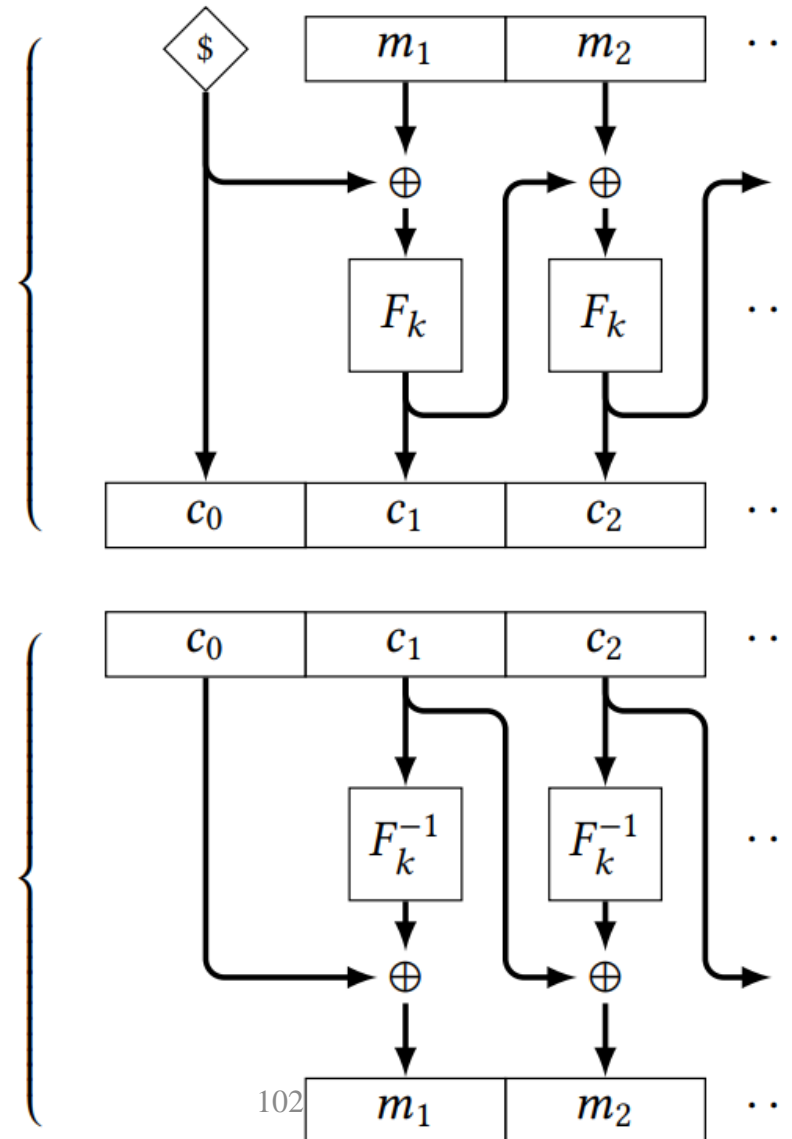
return $c_0 \parallel c_1 \parallel \dots \parallel c_\ell$

Dec($k, c_0 \parallel \dots \parallel c_\ell$):

for $i = 1$ to ℓ :

$m_i := F^{-1}(k, c_i) \oplus c_{i-1}$

return $m_1 \parallel \dots \parallel m_\ell$



CBC: Cipher Block Chaining

- The **most common** mode in practice.
- The CBC encryption of an ℓ -block plaintext is $\ell + 1$ blocks long.
- The first ciphertext block is called an **initialization vector (IV)**.
- Here we have described CBC mode as a **randomized** encryption, with the IV of each ciphertext being chosen **uniformly**.
- CBC mode provides **CPA security**.

CTR: Counter

- The next most common mode in practice is **counter mode**.
- Just like CBC mode, it involves an **additional IV block r** that is chosen **uniformly**. The idea is to then use the sequence

$$F(k, r); F(k, r + 1); F(k, r + 2); \dots$$

CTR: Counter

$\text{Enc}(k, m_1 \| \dots \| m_\ell):$

$r \leftarrow \{\mathbf{0}, \mathbf{1}\}^{blen}$


$c_0 := r$

for $i = 1$ to ℓ :

$c_i := F(k, r) \oplus m_i$

$r := r + 1 \% 2^{blen}$

return $c_0 \| \dots \| c_\ell$



OFB: Output Feedback

- OFB (output feedback) mode is **rarely** used in practice. But it has the **easiest security proof**.

$\text{Enc}(k, m_1 \| \dots \| m_\ell):$

$r \leftarrow \{0, 1\}^{blen}$

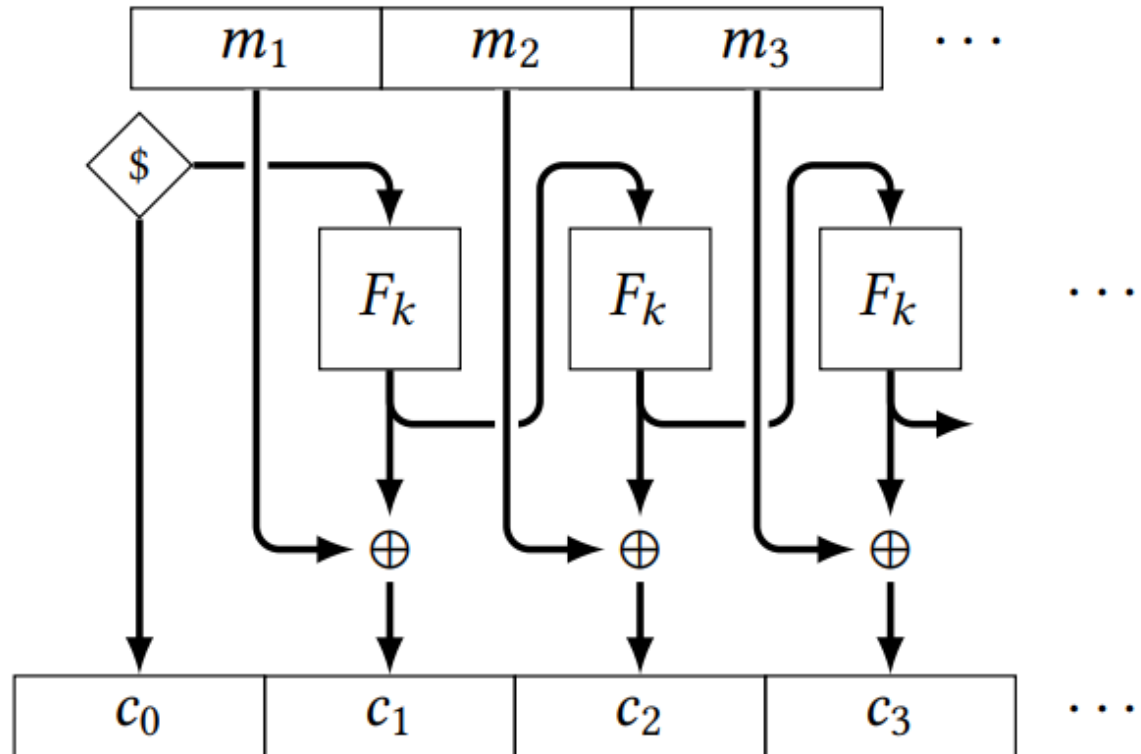
$c_0 := r$

for $i = 1$ to ℓ :

$r := F(k, r)$

$c_i := r \oplus m_i$


return $c_0 \| \dots \| c_\ell$



Compare & Contrast

CBC and CTR modes are essentially the **only two** modes that are ever considered **in practice for CPA security**. Both provide the **same security guarantees**.

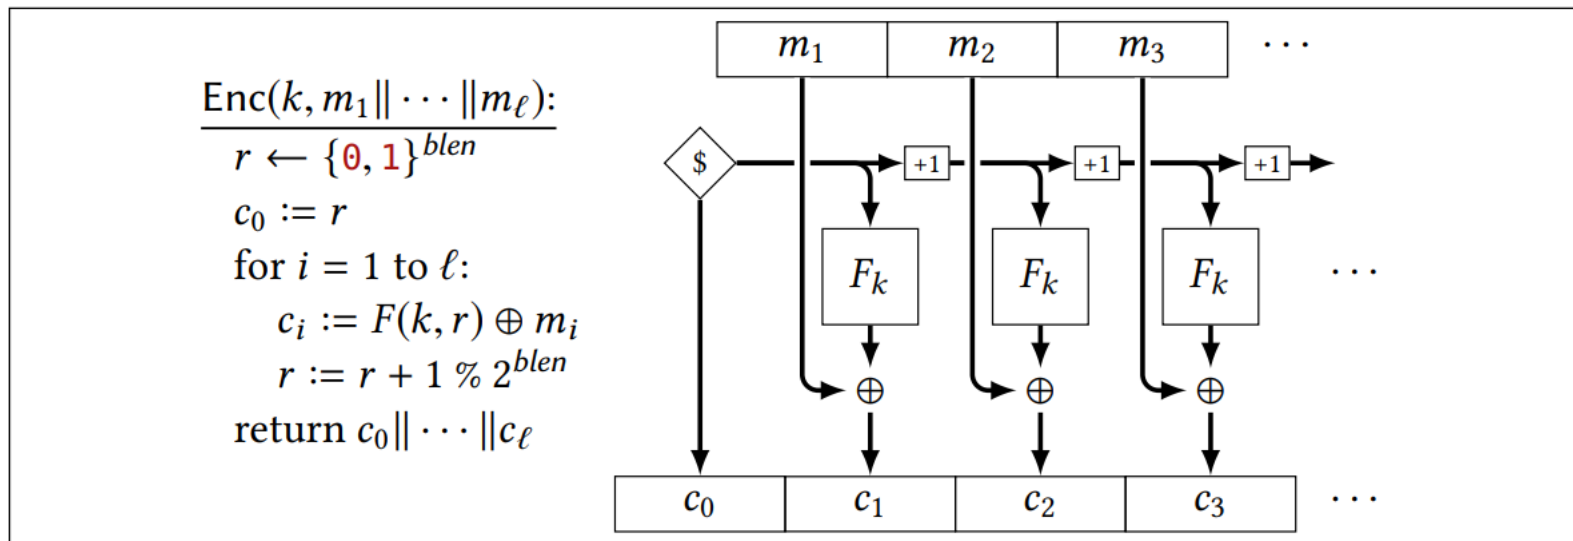
- CTR mode **does not** even use the block cipher's inverse F^{-1} . CTR mode can be instantiated from a **PRF**; it **doesn't need a PRP**. However, in practice it is rare to encounter an efficient PRF that is not a PRP.



Compare & Contrast

CBC and CTR modes are essentially the **only two** modes that are ever considered **in practice for CPA security**. Both provide the **same security guarantees**.


- CTR mode encryption can be **parallelized**. CBC mode **cannot**.
- If calls to the block cipher are expensive, it might be desirable to **pre-compute and store** them before the plaintext is known. CTR mode can, CBC mode cannot.



Compare & Contrast

CBC and CTR modes are essentially the **only two** modes that are ever considered **in practice for CPA security**. Both provide the **same security guarantees**.


- It is relatively **easy** to modify CTR to support plaintexts that are **not** an **exact multiple of the blocklength**.



Compare & Contrast

CBC and CTR modes are essentially the **only two** modes that are ever considered **in practice for CPA security**. Both provide the **same security guarantees**.

- It is common for implementers to **misunderstand** the security implications of the IV in these modes. **Many careless implementations allow an IV to be reused**. The effects of IV-reuse in CTR mode are quite devastating to message privacy.



Compare & Contrast

CBC and CTR modes are essentially the **only two** modes that are ever considered **in practice for CPA security**. Both provide the **same security guarantees**.

- In CBC mode, **reusing an IV can actually be safe**, if the two plaintexts **have different first blocks**!


$\text{Enc}(k, m_1 \| \dots \| m_\ell):$

$c_0 \leftarrow \{\mathbf{0}, \mathbf{1}\}^{blen};$

for $i = 1$ to ℓ :

$c_i := F(k, m_i \oplus c_{i-1})$

return $c_0 \| c_1 \| \dots \| c_\ell$



CPA Security and Variable-Length Plaintexts

- In CBC mode, a plaintext consisting of ℓ blocks is encrypted into a ciphertext of $\ell + 1$ blocks. In other words, **the ciphertext leaks the number of blocks in the plaintext**.
- Cannot achieve the CPA security we have defined.

$\mathcal{L}_{\text{cpa-L}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{EAVESDROP}(m_L, m_R \in \Sigma.\mathcal{M}):$
$c := \Sigma.\text{Enc}(k, m_L)$
return c

$\mathcal{L}_{\text{cpa-R}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{EAVESDROP}(m_L, m_R \in \Sigma.\mathcal{M}):$
$c := \Sigma.\text{Enc}(k, m_R)$
return c

CPA Security and Variable-Length Plaintexts

- Suppose we **don't** really **care about hiding the length of plaintexts**. Is there a way to make a security definition that says: *ciphertexts hide everything about the plaintext, except their length*?
- When discussing encryption schemes that support **variable-length** plaintexts, CPA security will refer to the following updated libraries.

$\mathcal{L}_{\text{cpa-L}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
<u>$\text{CTXT}(m_L, m_R \in \Sigma.\mathcal{M})$:</u>
if $ m_L \neq m_R $ return err
$c := \Sigma.\text{Enc}(k, m_L)$
return c

$\mathcal{L}_{\text{cpa-R}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
<u>$\text{CTXT}(m_L, m_R \in \Sigma.\mathcal{M})$:</u>
if $ m_L \neq m_R $ return err
$c := \Sigma.\text{Enc}(k, m_R)$
return c

CPA Security and Variable-Length Plaintexts

- Then when discussing encryption schemes supporting **variable-length** plaintexts, CPA\$ security will refer to the following libraries:

$\mathcal{L}_{\text{cpa\$-real}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{CHALLENGE}(m \in \Sigma.\mathcal{M}):$
$c := \Sigma.\text{Enc}(k, m)$
return c

$\mathcal{L}_{\text{cpa\$-rand}}^\Sigma$
$\text{CHALLENGE}(m \in \Sigma.\mathcal{M}):$
$c \leftarrow \Sigma.C(m)$
return c

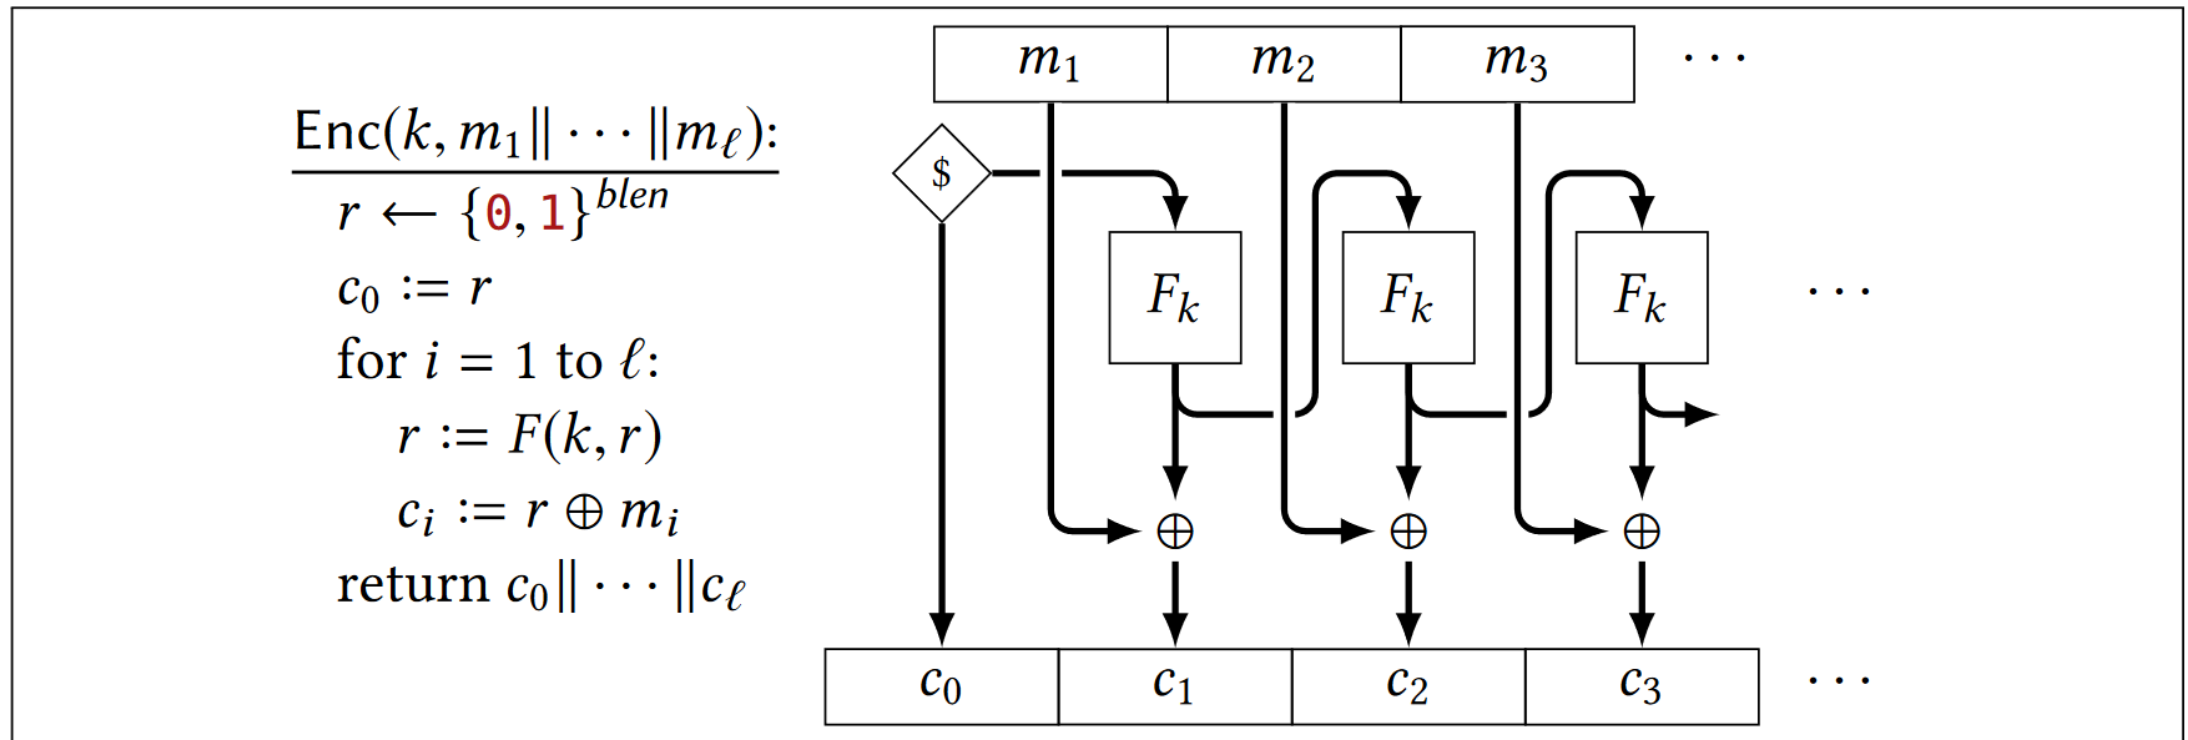
- With respect to these updated security definitions, **CPA\$ security implies CPA security** as before.

Don't Take Length-Leaking for Granted!

- We have just gone from requiring encryption to leak no partial information to casually **allowing some specific information to leak**.
- If we want to **truly support plaintexts of arbitrary length**, then leaking the length is in fact **unavoidable**.
- By observing only the length of encrypted network traffic, many serious attacks are **possible**.
 - Google maps
 - Variable-bit-rate (VBR) :the changes in bit rate are reflected as changes in packet length
 - Netflix, Youtube...
 - Voice chat programs (who was speaking, the language being spoken)

Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.



Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.

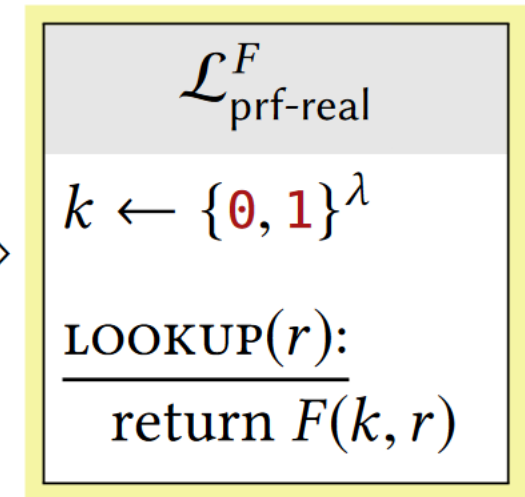
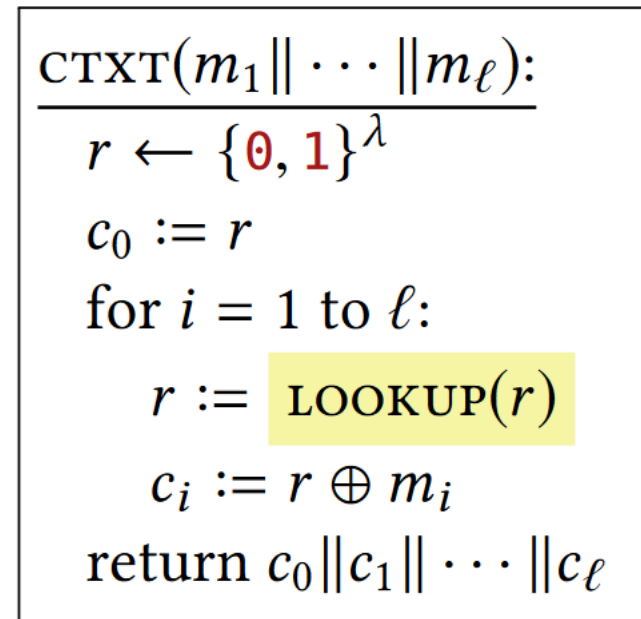

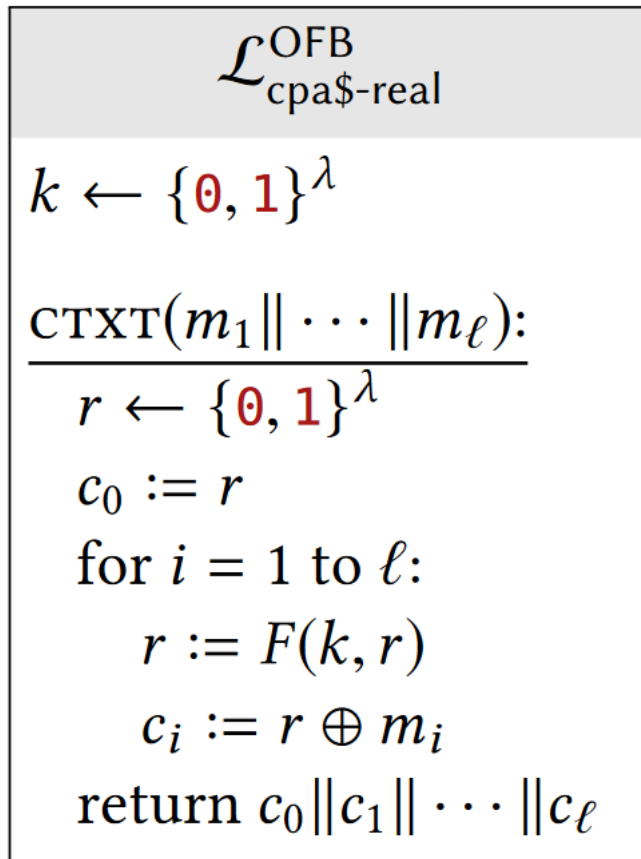
proof

- Each ciphertext block (apart from the IV) is computed as $c_i := r \oplus m_i$. By the one-time pad rule, it suffices to show that the r values are **independently pseudorandom**.
- Each r value is the result of a call to the PRF. These PRF outputs will be **independently pseudorandom only if all of the inputs to the PRF are distinct**.

Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.

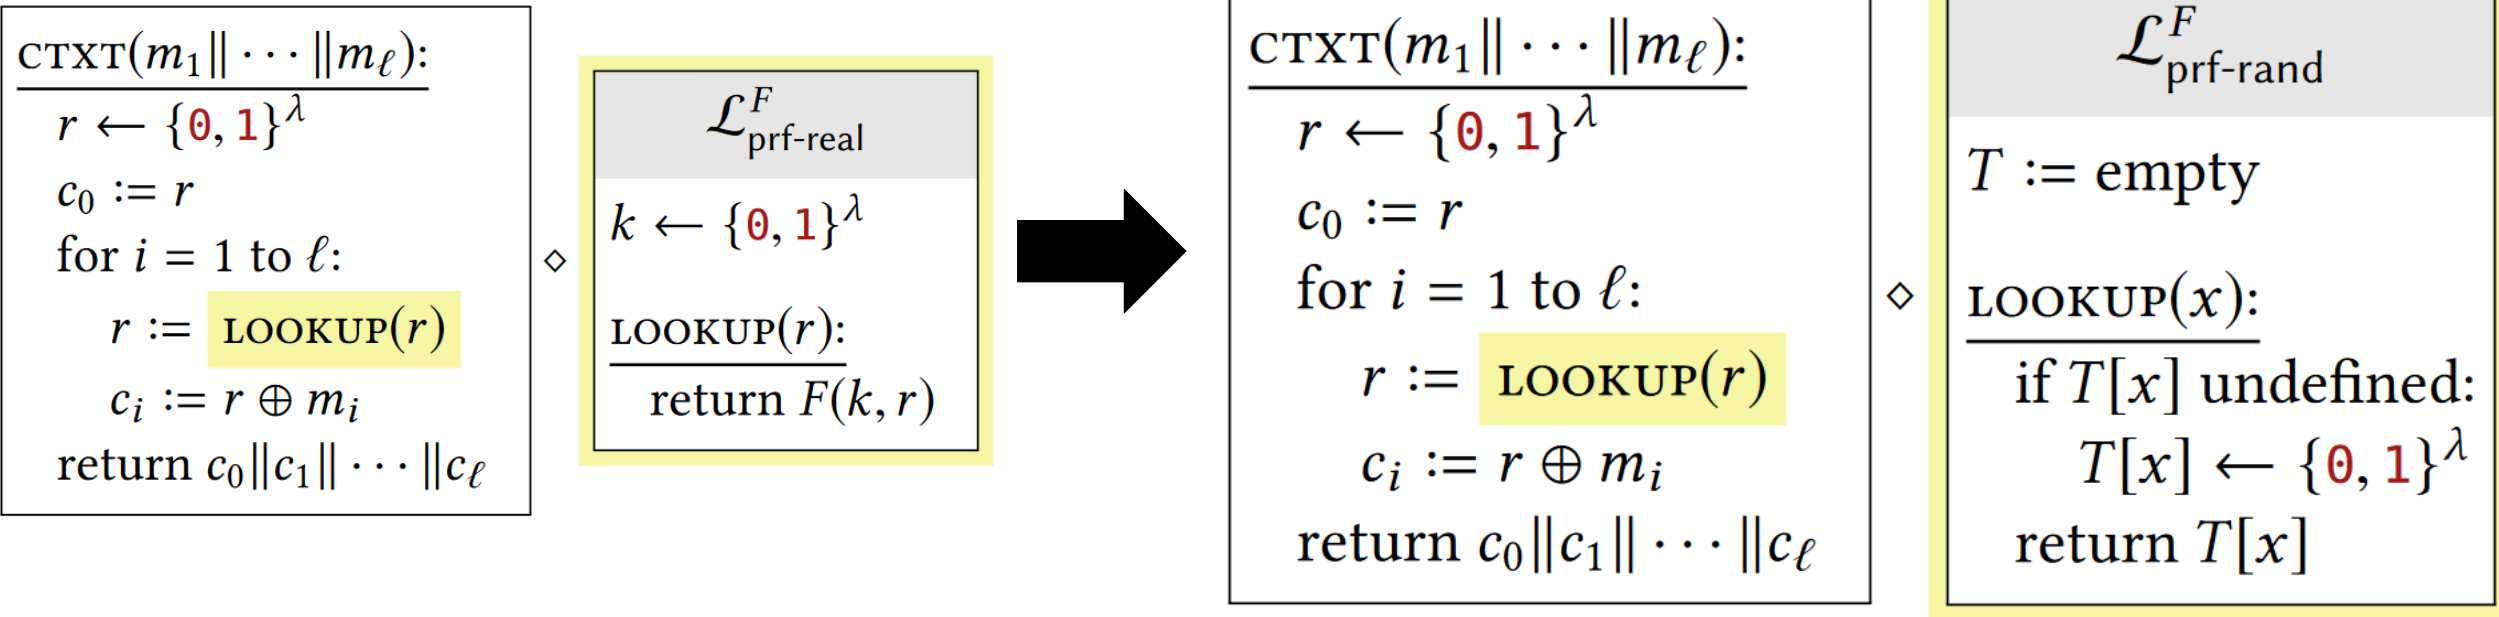
proof



Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.

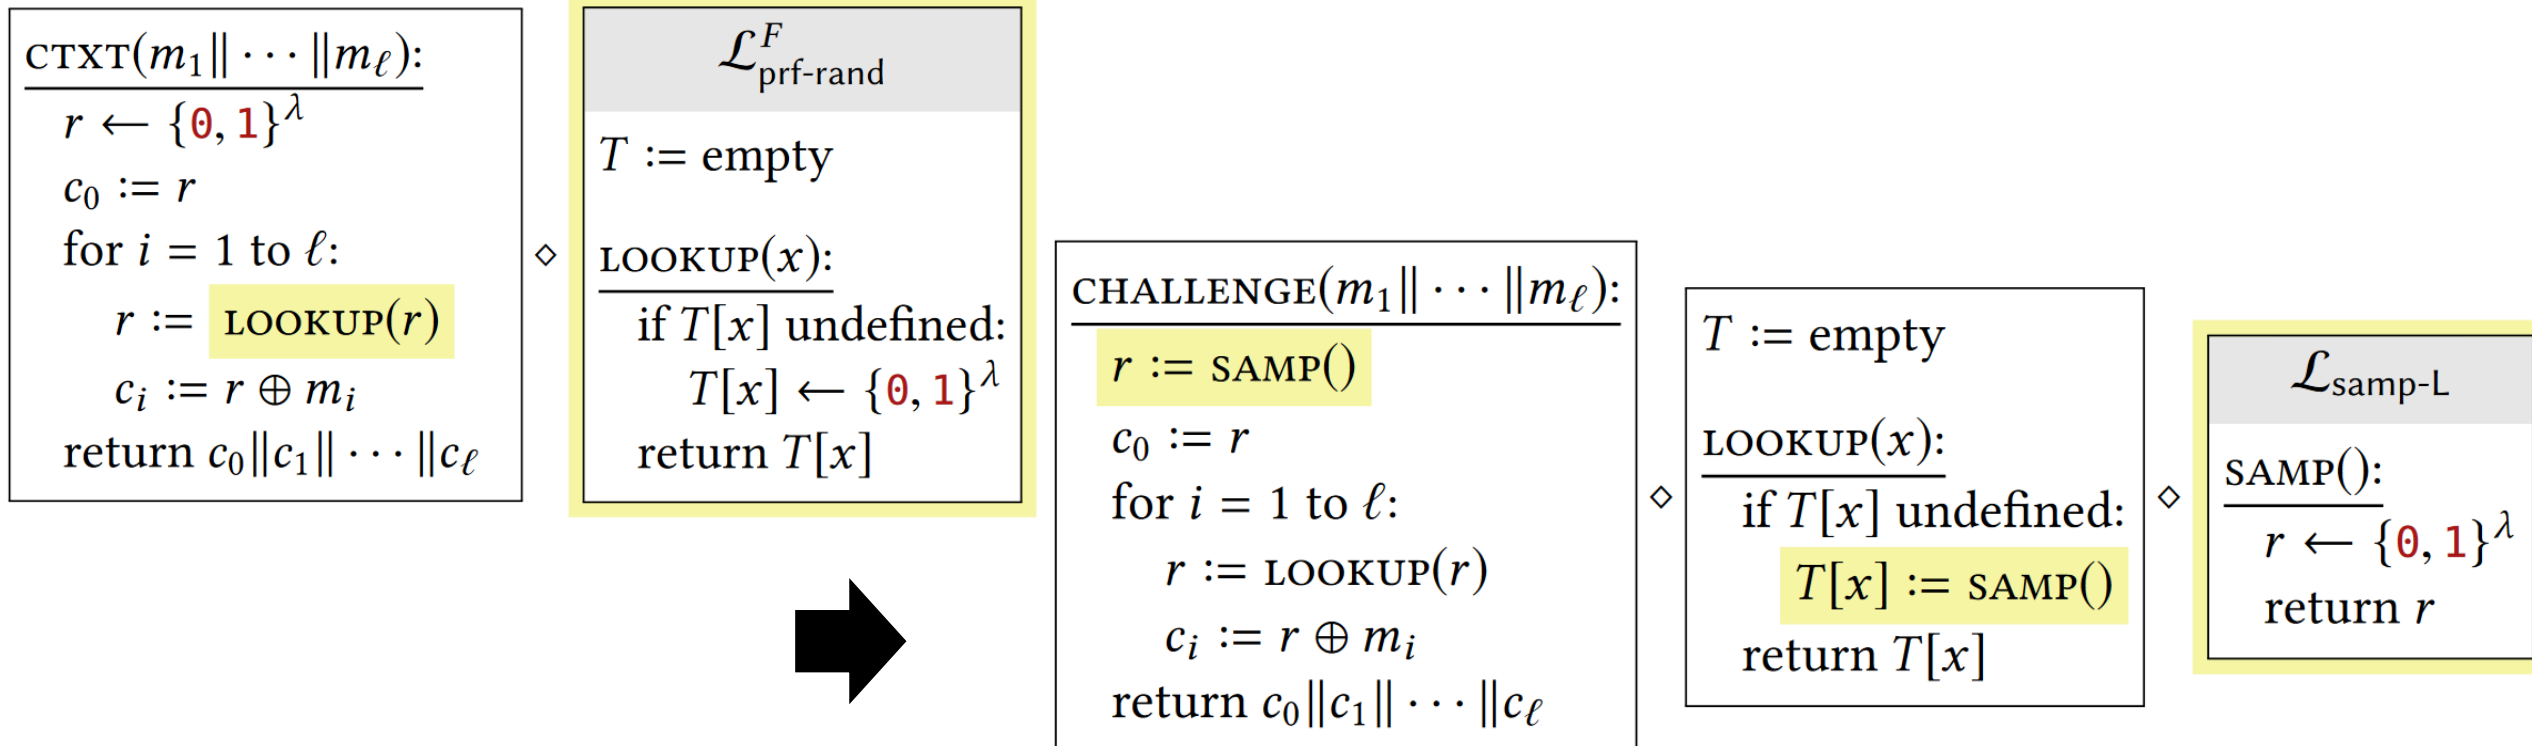
proof



Security of OFB Mode

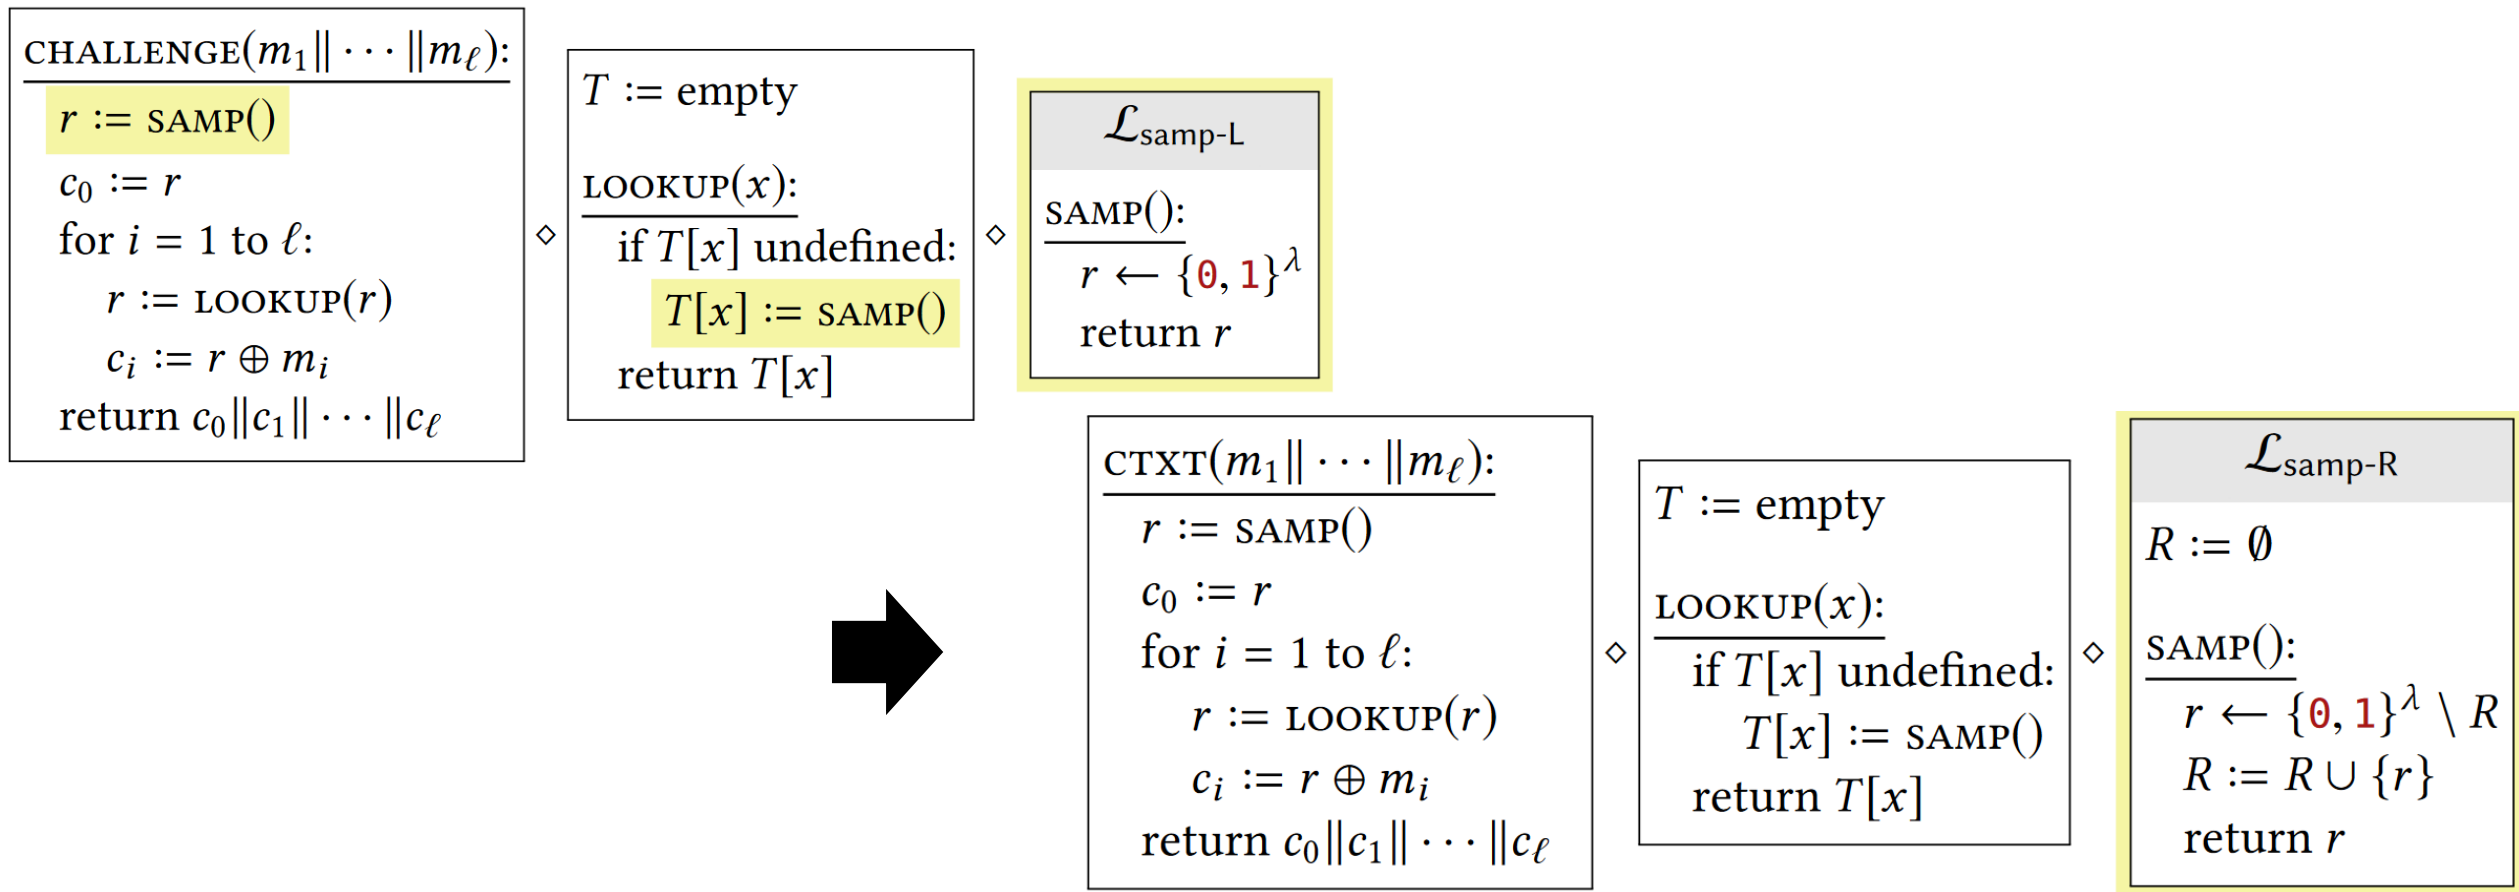
Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.

proof



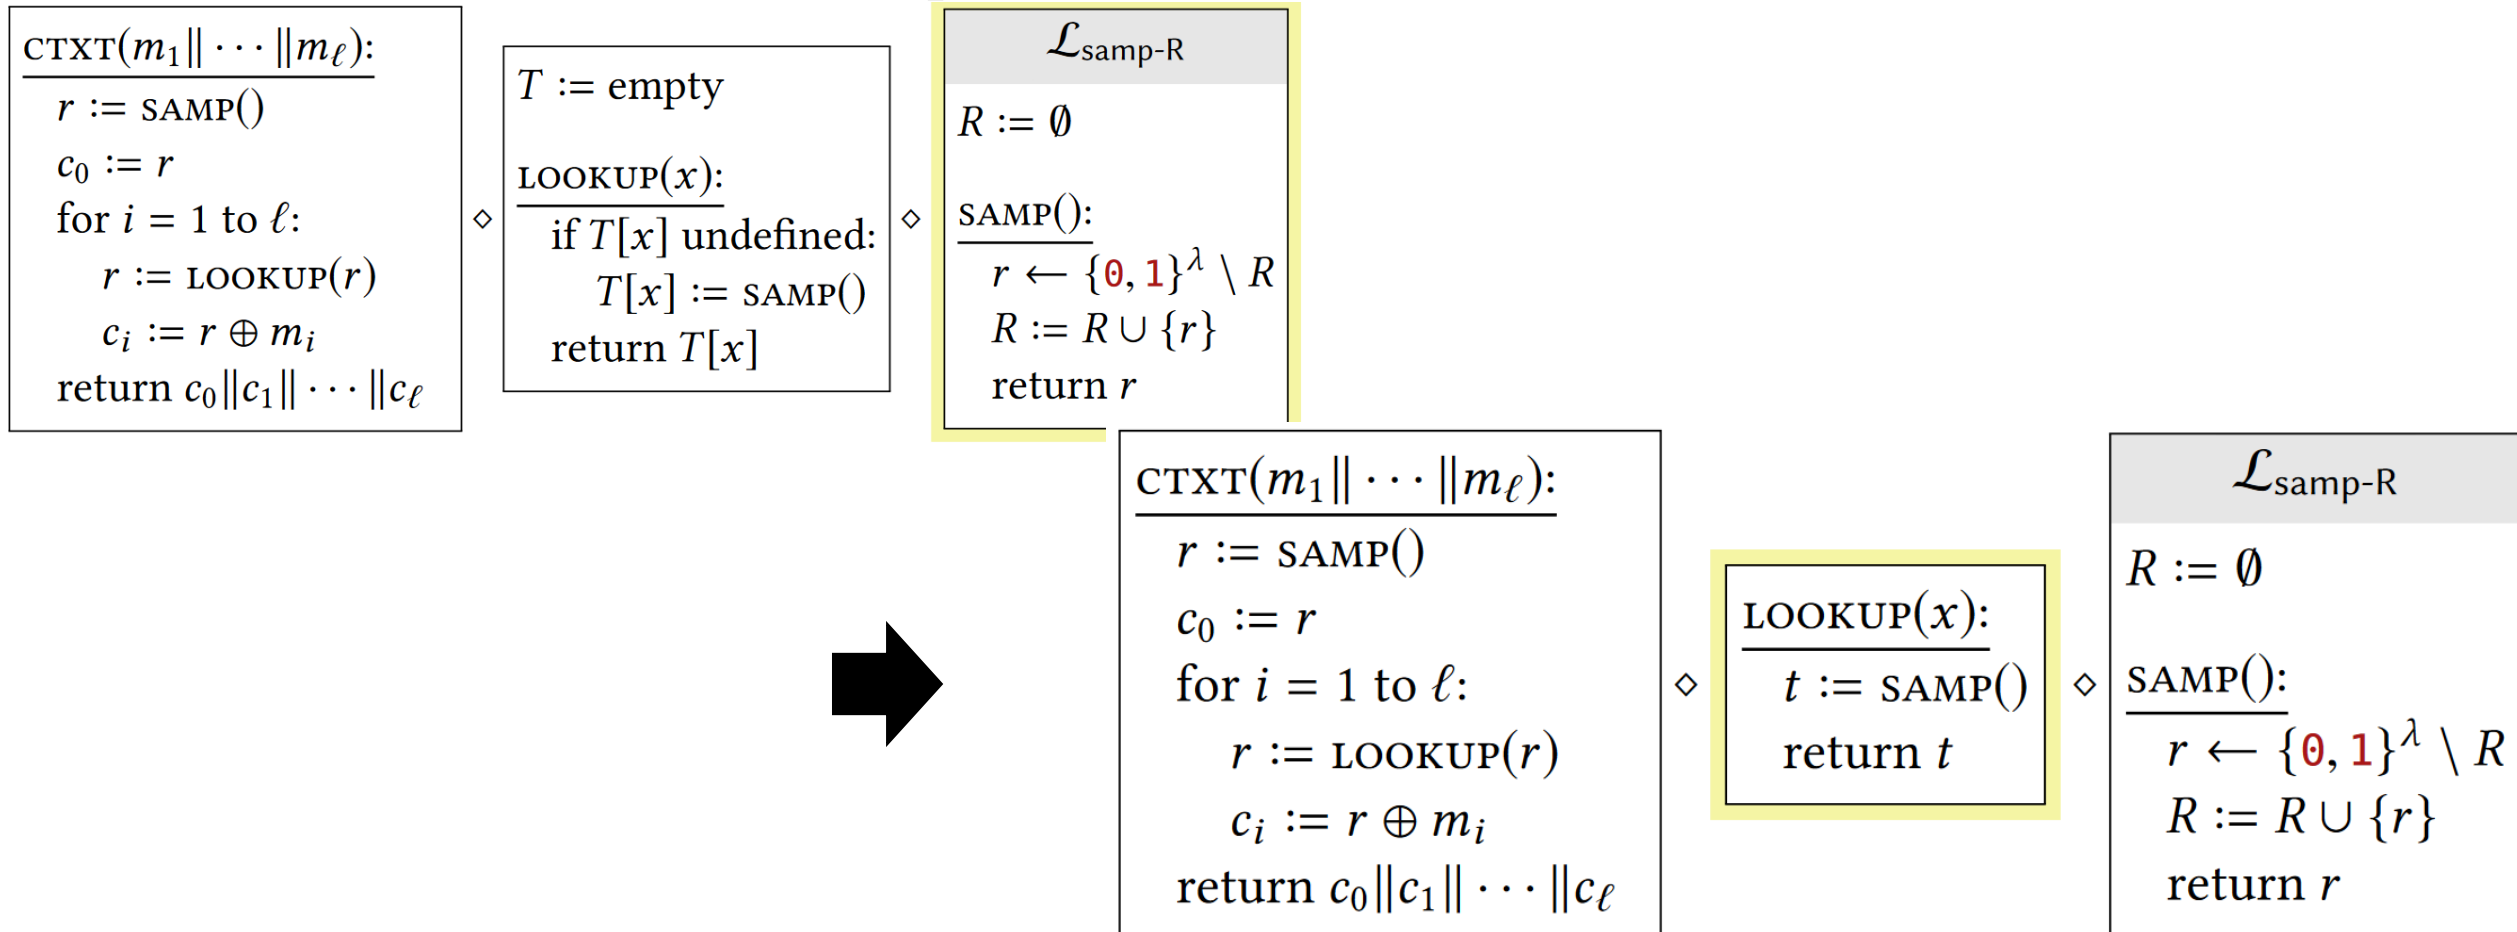
Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.



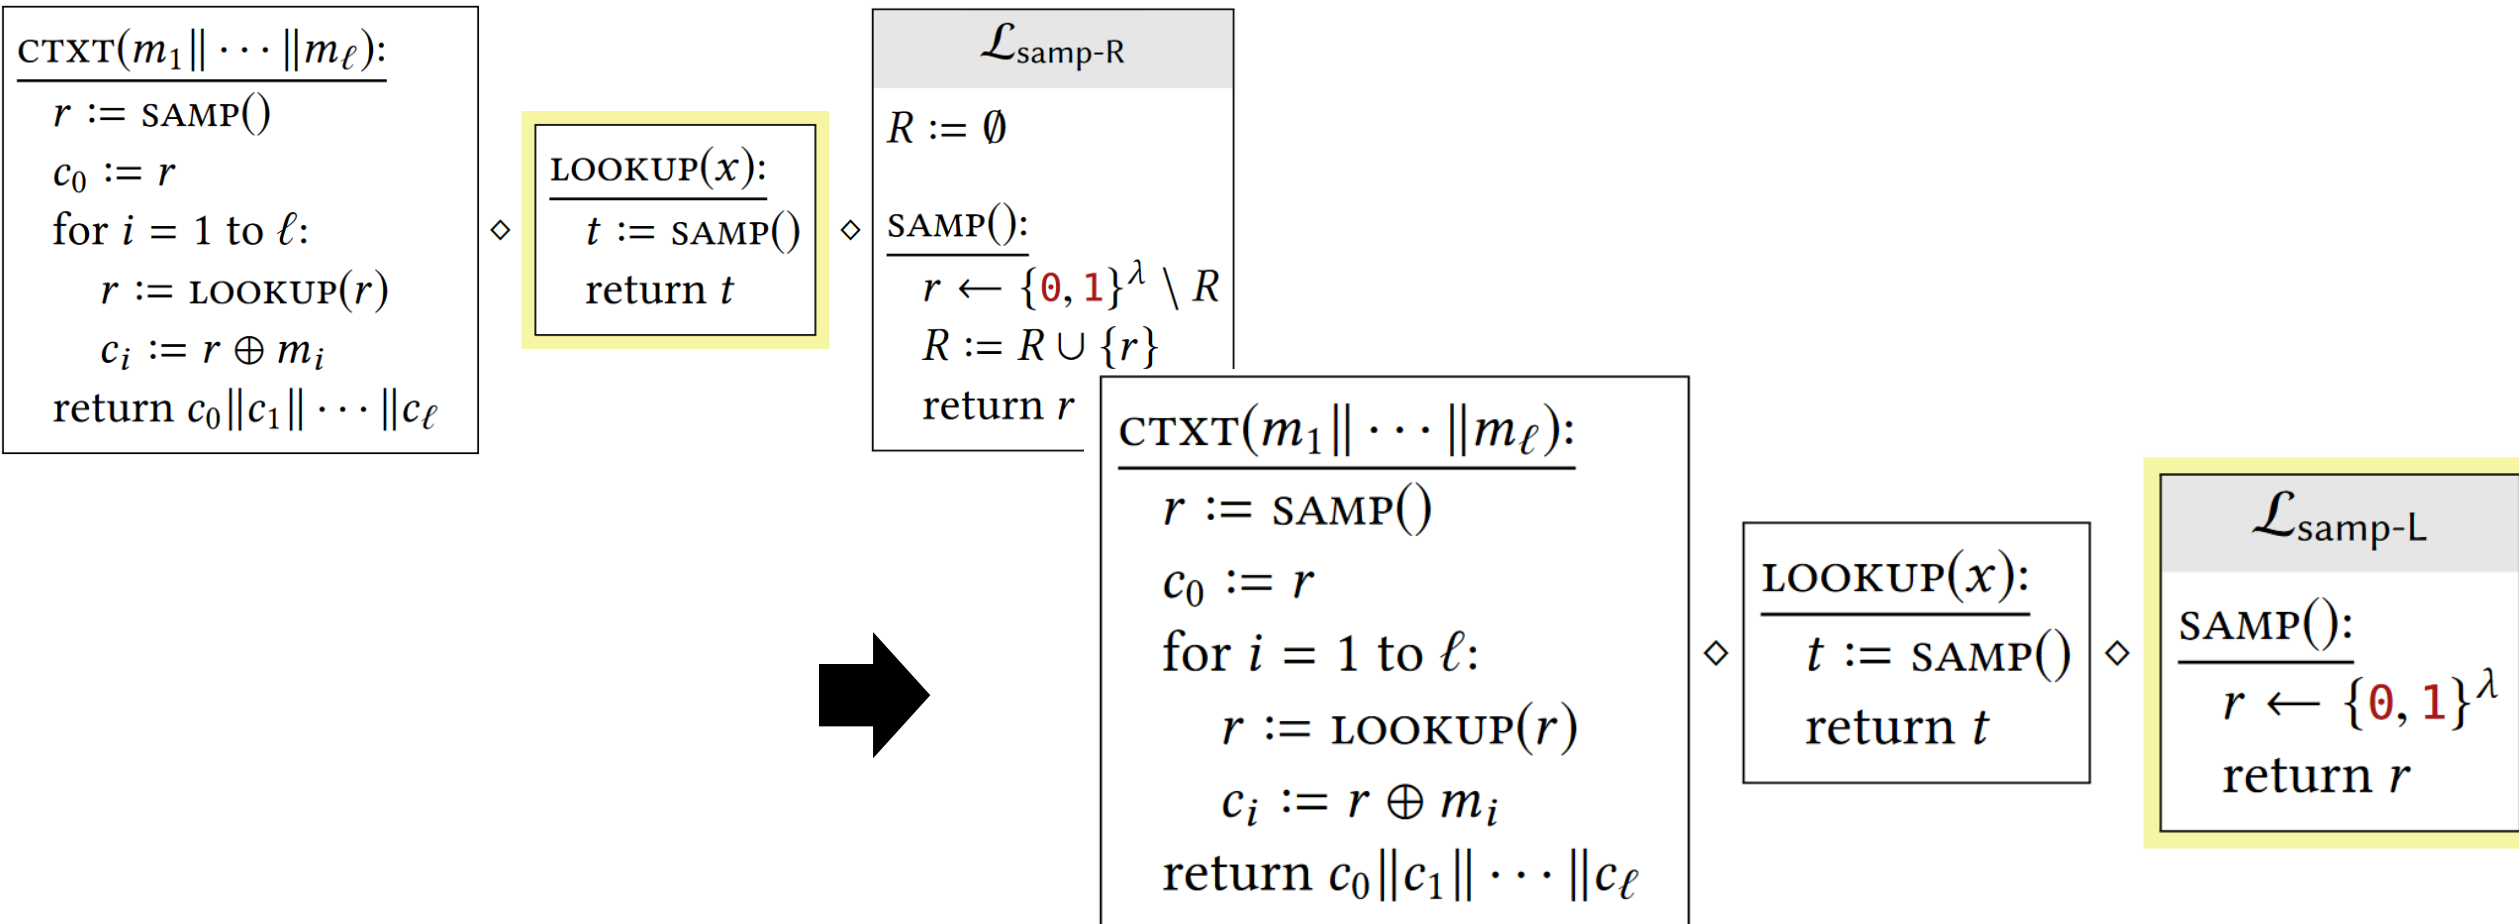
Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.



Security of OFB Mode

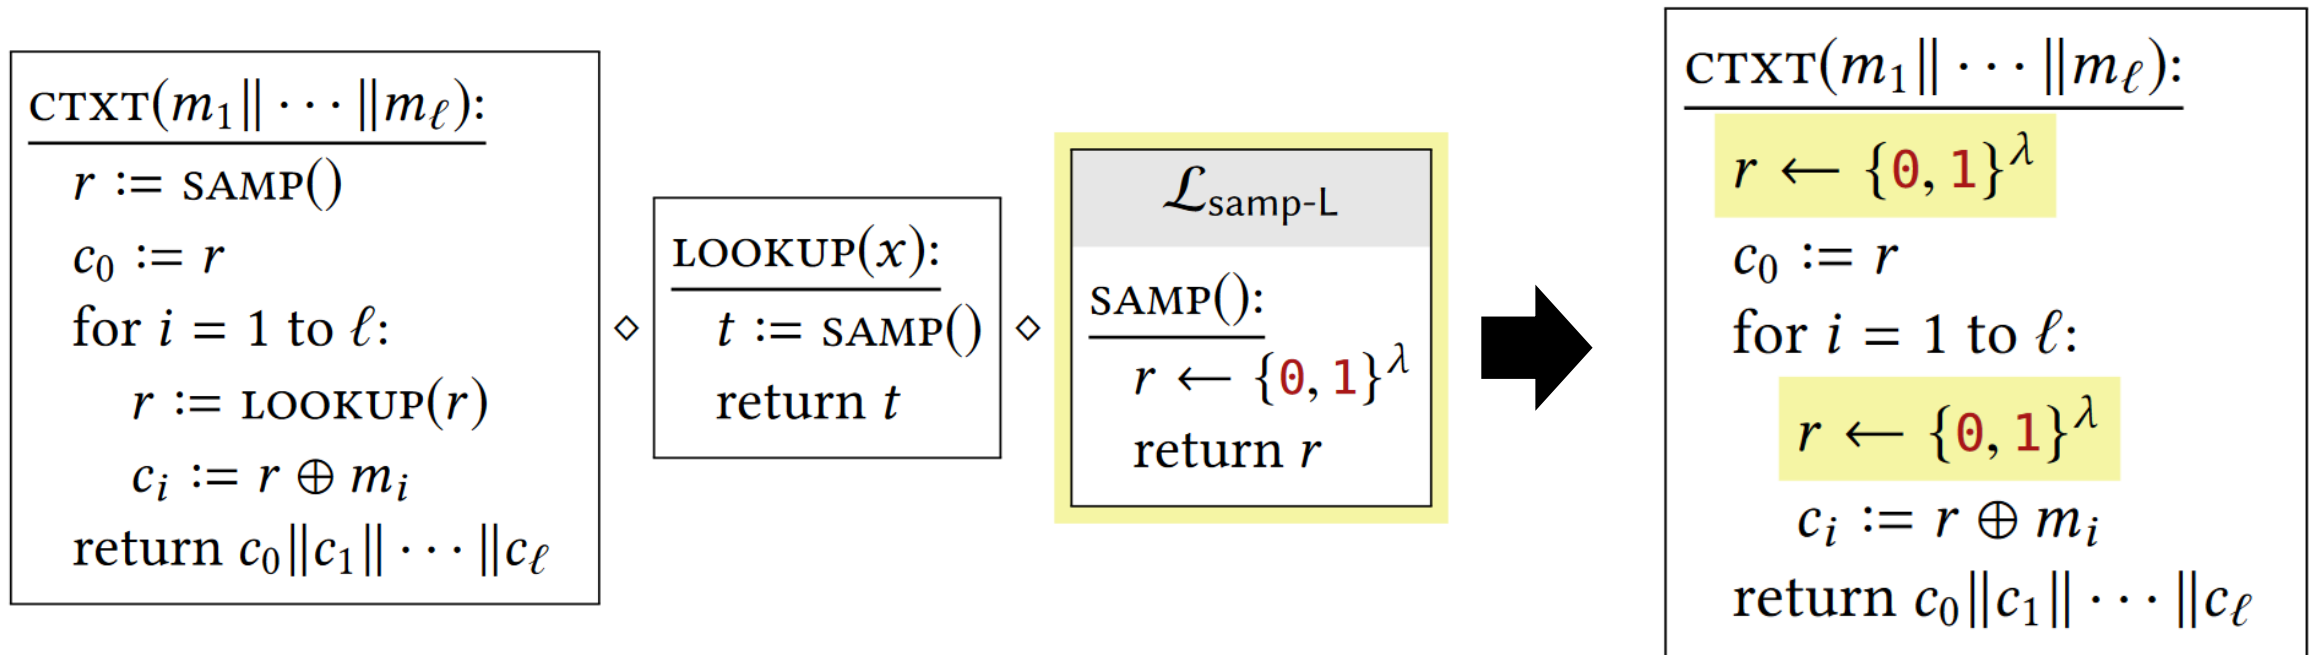
Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.



Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.


proof



Security of OFB Mode

Claim OFB mode has CPA\$ security, if its underlying block cipher F is a secure PRF (no need for PRP) with parameters $in = out = \lambda$.

proof



Padding & Ciphertext Stealing

- So far we have assumed that all plaintexts are **exact multiples of the blocklength**.
- How are block ciphers used in practice with data that has **arbitrary length**?
- Padding
- Ciphertext Stealing

Padding

- Encode arbitrary-length data into data that is a **multiple of the blocklength**.
- A padding scheme should consist of two algorithms:
 - pad: takes as input a string of **any length**, and outputs a string whose length is a **multiple of the blocklength**
 - unpad: the inverse of pad. We require that $\text{unpad}(\text{pad}(x)) = x$ for all strings x .
- In the real world, data almost always comes in **bytes** and **not bits**
- Padding schemes are **not a security feature!**

Padding

- Null padding
 - Fill the final block with **null bytes** (00)
 - It is **not always reversible**
 - pad(31 41 59) and pad(31 41 59 00) will **give the same result**

Padding

- ANSI X.923 standard
 - Data is padded with **null bytes**, except for the last byte of padding which indicates **how many padding bytes there are**. (tell the receiver how many bytes to remove to recover the original message)
 - If the original unpadded data is already a multiple of the block length, then an **entire extra block of padding must be added**.

01 34 11 d9 81 88 05 57 1d 73 c3 00 00 00 00 05 \Rightarrow *valid*

95 51 05 4a d6 5a a3 44 af b3 85 00 00 00 00 03 \Rightarrow *valid*

71 da 77 5a 5e 77 eb a8 73 c5 50 b5 81 d5 96 01 \Rightarrow *valid*

5b 1c 01 41 5d 53 86 4e e4 94 13 e8 7a 89 c4 71 \Rightarrow *invalid*

d4 0d d8 7b 53 24 c6 d1 af 5f d6 f6 00 c0 00 04 \Rightarrow *invalid*

Padding

- PKCS#7 standard
 - If b bytes of padding are needed, then the data is padded **not** with null bytes but with **b** bytes.

01 34 11 d9 81 88 05 57 1d 73 c3 05 05 05 05 05 \Rightarrow *valid*

95 51 05 4a d6 5a a3 44 af b3 85 03 03 03 03 03 \Rightarrow *valid*

71 da 77 5a 5e 77 eb a8 73 c5 50 b5 81 d5 96 01 \Rightarrow *valid*

5b 1c 01 41 5d 53 86 4e e4 94 13 e8 7a 89 c4 71 \Rightarrow *invalid*

d4 0d d8 7b 53 24 c6 d1 af 5f d6 f6 04 c0 04 04 \Rightarrow *invalid*

Padding

- ISO/IEC 7816-4 standard
 - The data is padded with a **80** byte followed by null bytes. To remove the padding, remove all trailing null bytes and ensure that the last byte is **80** (and then remove it too).
 - The significance of **80** is clearer when you write it in **binary** as 10000000

01 34 11 d9 81 88 05 57 1d 73 c3 **80** **00** **00** **00** **00** \Rightarrow *valid*

95 51 05 4a d6 5a a3 44 af b3 85 03 03 **80** **00** **00** \Rightarrow *valid*

71 da 77 5a 5e 77 eb a8 73 c5 50 b5 81 d5 96 **80** \Rightarrow *valid*

5b 1c 01 41 5d 53 86 4e e4 94 13 e8 7a 89 c4 **71** \Rightarrow *invalid*


d4 0d d8 7b 53 24 c6 d1 af 5f d6 f6 **c4** **00** **00** **00** \Rightarrow *invalid*

Ciphertext Stealing

- The main idea behind ciphertext stealing is to use a standard block-cipher mode that **only supports full blocks (e.g., CBC mode)**, and then simply **throw away some bits of the ciphertext**, in such a way that decryption is still possible.

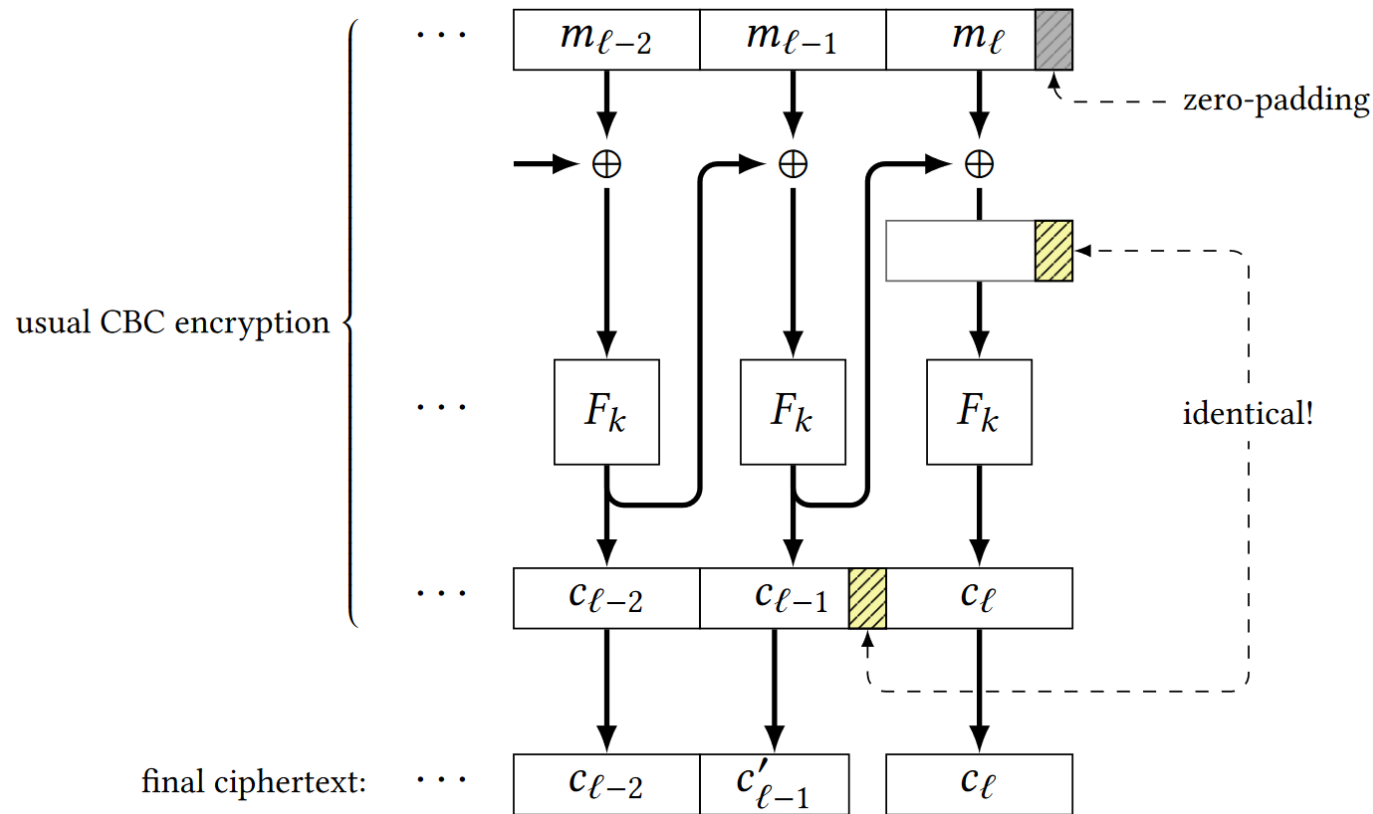
Ciphertext Stealing (for CBC mode)

- We start by extending m_ℓ with j zeroes and performing CBC as usual. Then **throw away the last j bits of $c_{\ell-1}$** (the next-to-last block of the CBC ciphertext).



Ciphertext Stealing (for CBC mode)

- For decryption: Those missing j bits were **redundant**, because there is another way to compute them. Let $x^* := c_{\ell-1} \oplus m_\ell$. Since the last j bits of m_ℓ are 0's, the last j bits of x^* are the last j bits of $c_{\ell-1}$ — the missing bits.



Ciphertext Stealing (for CBC mode)

- In practice, the last two blocks of the ciphertext are often **interchanged**.

Enc($k, m_1 \parallel \dots \parallel m_\ell$):

*// each m_i is $blen$ bits,
// except possibly m_ℓ*

$j := blen - |m_\ell|$

$m_\ell := m_\ell \parallel \mathbf{0}^j$

$c_0 \leftarrow \{\mathbf{0}, \mathbf{1}\}^{blen}$:

for $i = 1$ to ℓ :

$c_i := F(k, m_i \oplus c_{i-1})$

if $j \neq 0$:

remove final j bits of $c_{\ell-1}$

swap $c_{\ell-1}$ and c_ℓ

return $c_0 \parallel c_1 \parallel \dots \parallel c_\ell$

Dec($k, c_0 \parallel \dots \parallel c_\ell$):

*// each c_i is $blen$ bits,
// except possibly c_ℓ*

$j := blen - |c_\ell|$

if $j \neq 0$:

swap $c_{\ell-1}$ and c_ℓ

$x :=$ last j bits of $F^{-1}(k, c_\ell)$

$c_{\ell-1} := c_{\ell-1} \parallel x$

for $i = 1$ to ℓ :

$m_i := F^{-1}(k, c_i) \oplus c_{i-1}$

remove final j bits of m_ℓ

return $m_1 \parallel \dots \parallel m_\ell$

The marked lines correspond to plain CBC mode.