

Advanced Cryptography

(Provable Security)

Yi LIU

Construction of PRG

Theorem Let f be a one-way permutation with hard-core predicate hc . Then algorithm G defined by $G(s) = f(s) || hc(s)$ is a pseudorandom generator with stretch $\ell = 1$.

Since f is a **permutation**, $f(s)$ itself is uniformly distributed.

We can also construct PRGs based on one-way function.

Construction of PRG

Theorem Let f be a one-way permutation with hard-core predicate hc . Then algorithm G defined by $G(s) = f(s) || hc(s)$ is a pseudorandom generator with stretch $\ell = 1$.

proof

$\mathcal{L}_{\text{prg-real}}^G$
$\text{QUERY}():$ $s \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$ return $G(s)$

$\mathcal{L}_{\text{prg-rand}}^G$
$\text{QUERY}():$ $r \leftarrow \{\mathbf{0}, \mathbf{1}\}^{\lambda+\ell}$ return r

$\mathcal{L}_{\text{prg-real}}^G$
$\text{QUERY}():$ $s \leftarrow \{0,1\}^\lambda$ $c := f(s) hc(s)$ return c

$\text{QUERY}():$ $s \leftarrow \{0,1\}^\lambda$ $\mathbf{b} \leftarrow \{\mathbf{0}, \mathbf{1}\}$ $c := f(s) \mathbf{b}$ return c
--

$\text{QUERY}():$ $s \leftarrow \{0,1\}^\lambda$ $b \leftarrow \{0,1\}$ $c := \mathbf{s} \mathbf{b}$ return c
--

$\mathcal{L}_{\text{prg-rand}}^G$
$\text{QUERY}():$ $r \leftarrow \{\mathbf{0}, \mathbf{1}\}^{\lambda+\ell}$ return r

Message Authentication Codes

Message Authentication Codes

- The challenge of CCA-secure encryption is dealing with ciphertexts that were **generated by an adversary**. Imagine there was a way to “certify” that a **ciphertext was not adversarially generated** — i.e., it was generated by someone who knows the secret key.
- What we are asking for is **not to hide the ciphertext but to authenticate it**: message authentication code.

Message Authentication Codes



Definition A message authentication code (MAC) scheme for message space \mathcal{M} consists of the following algorithms:

- KeyGen: sample a key.
- MAC: take a key k and a message $m \in \mathcal{M}$ as input, and output a tag t .

The mac algorithm is **deterministic**.

How to Think About Authenticity Properties

- “an adversary should not be able to guess a uniformly chosen λ -bit value.”



- It's **hard** for an adversary to find/generate this special value!

The MAC Security Definition

- A more useful property is:
even if the adversary knows valid MAC tags corresponding to various messages, she cannot produce a valid MAC tag for a different message.
- We call it a **forgery** if the adversary can produce a “new” valid MAC tag.

The MAC Security Definition

Definition Let Σ be a MAC scheme. We say that Σ is a secure MAC if

$\mathcal{L}_{\text{mac-real}}^{\Sigma} \approx \mathcal{L}_{\text{mac-fake}}^{\Sigma}$, where:

$\mathcal{L}_{\text{mac-real}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
$\text{GETTAG}(m \in \Sigma.\mathcal{M}):$ <hr/> return $\Sigma.\text{MAC}(k, m)$
$\text{CHECKTAG}(m \in \Sigma.\mathcal{M}, t):$ <hr/> return $t \stackrel{?}{=} \Sigma.\text{MAC}(k, m)$

$\mathcal{L}_{\text{mac-fake}}^{\Sigma}$
$k \leftarrow \Sigma.\text{KeyGen}$
$\mathcal{T} := \emptyset$
$\text{GETTAG}(m \in \Sigma.\mathcal{M}):$ <hr/> $t := \Sigma.\text{MAC}(k, m)$ $\mathcal{T} := \mathcal{T} \cup \{(m, t)\}$ return t
$\text{CHECKTAG}(m \in \Sigma.\mathcal{M}, t):$ <hr/> return $(m, t) \stackrel{?}{\in} \mathcal{T}$

MAC Applications

- Although MACs are less embedded in public awareness than encryption, they are **extremely useful**. A frequent application of MACs is to **store some information in an untrusted place**, where we don't intend to hide the data, only **ensure that the data is not changed**.
 - Browser cookie: Imagine a webserver that stores a cookie when a user logs in, containing that user's account name. Adding a MAC tag of the cookie data (using a key known only to the server) ensures that **an attacker cannot modify their cookie to contain a different user's account name**.
 - ...

A PRF is a MAC

- The definition of a PRF says (more or less) that even if you've seen the output of the PRF on several chosen inputs, all other outputs look **independently & uniformly random**.
- Furthermore, uniformly chosen values are **hard to guess**, as long as they are sufficiently long (e.g., λ bits).
- In other words, after seeing some outputs of a PRF, any other PRF output will be **hard to guess**.

A PRF is a MAC

Claim The following two libraries are indistinguishable:

$\mathcal{L}_{\text{guess-L}}$
$T := \text{empty assoc. array}$
$\text{GUESS}(m \in \{\mathbf{0}, \mathbf{1}\}^{in}, g \in \{\mathbf{0}, \mathbf{1}\}^\lambda):$
<hr/>
if $T[m]$ undefined:
$T[m] \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$
return $g \stackrel{?}{=} T[m]$
$\text{REVEAL}(m \in \{\mathbf{0}, \mathbf{1}\}^{in}):$
<hr/>
if $T[m]$ undefined:
$T[m] \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$
return $T[m]$

$\mathcal{L}_{\text{guess-R}}$
$T := \text{empty assoc. array}$
$\text{GUESS}(m \in \{\mathbf{0}, \mathbf{1}\}^{in}, g \in \{\mathbf{0}, \mathbf{1}\}^\lambda):$
<hr/>
<i>// returns false if $T[m]$ undefined</i>
return $g \stackrel{?}{=} T[m]$
$\text{REVEAL}(m \in \{\mathbf{0}, \mathbf{1}\}^{in}):$
<hr/>
if $T[m]$ undefined:
$T[m] \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$
return $T[m]$

A PRF is a MAC

Claim The two libraries are indistinguishable.

proof

Let q be the number of queries that the calling program makes to guess. We will show that the libraries are indistinguishable with a hybrid sequence of the form:

$$\mathcal{L}_{\text{guess-L}} \equiv \mathcal{L}_{\text{hyb-0}} \approx \mathcal{L}_{\text{hyb-1}} \approx \cdots \approx \mathcal{L}_{\text{hyb-}q} \equiv \mathcal{L}_{\text{guess-R}}$$

A PRF is a MAC

Claim The two libraries are indistinguishable.

proof

$\mathcal{L}_{\text{hyb-}h}$
$count := 0$ $T := \text{empty assoc. array}$ <u>GUESS(m, g):</u> $count := count + 1$ if $T[m]$ undefined and $count > h$: $T[m] \leftarrow \{0, 1\}^\lambda$ return $g \stackrel{?}{=} T[m]$ <i>// returns false if $T[m]$ undefined</i> <u>REVEAL(m):</u> if $T[m]$ undefined: $T[m] \leftarrow \{0, 1\}^\lambda$ return $T[m]$

In $\mathcal{L}_{\text{hyb-}0}$, the clause “ $count > 0$ ” is **always true** so this clause **can be removed from the if-condition**. This modification results in $\mathcal{L}_{\text{guess-L}}$, so we have $\mathcal{L}_{\text{guess-L}} \equiv \mathcal{L}_{\text{hyb-}0}$.

In $\mathcal{L}_{\text{hyb-}q}$, the clause “ $count > q$ ” in the if-statement is **always false** since the calling program makes only q queries. Removing the unreachable if-statement it results in $\mathcal{L}_{\text{guess-R}}$, so we have $\mathcal{L}_{\text{guess-R}} \equiv \mathcal{L}_{\text{hyb-}q}$. It remains to show that $\mathcal{L}_{\text{guess-}h} \equiv \mathcal{L}_{\text{hyb-}(h+1)}$ for all h . We can rewrite the libraries.

A PRF is a MAC

It remains to show that $\mathcal{L}_{\text{guess}-h} \equiv \mathcal{L}_{\text{hyb}-(h+1)}$ for all h .

Claim The two libraries are indistinguishable.

proof

$\mathcal{L}_{\text{hyb}-h}$
$count := 0$ $T := \text{empty assoc. array}$
<u>GUESS(m, g):</u> $count := count + 1$ if $T[m]$ undefined and $count > h$: $T[m] \leftarrow \{0, 1\}^\lambda$ if $g = T[m]$ and $count = h + 1$: $bad := 1$ return $g \stackrel{?}{=} T[m]$ <i>// returns false if $T[m]$ undefined</i>
<u>REVEAL(m):</u> if $T[m]$ undefined: $T[m] \leftarrow \{0, 1\}^\lambda$ return $T[m]$

$\mathcal{L}_{\text{hyb}-(h+1)}$
$count := 0$ $T := \text{empty assoc. array}$
<u>GUESS(m, g):</u> $count := count + 1$ if $T[m]$ undefined and $count > h$: $T[m] \leftarrow \{0, 1\}^\lambda$ if $g = T[m]$ and $count = h + 1$: $bad := 1$; return false return $g \stackrel{?}{=} T[m]$ <i>// returns false if $T[m]$ undefined</i>
<u>REVEAL(m):</u> if $T[m]$ undefined: $T[m] \leftarrow \{0, 1\}^\lambda$ return $T[m]$

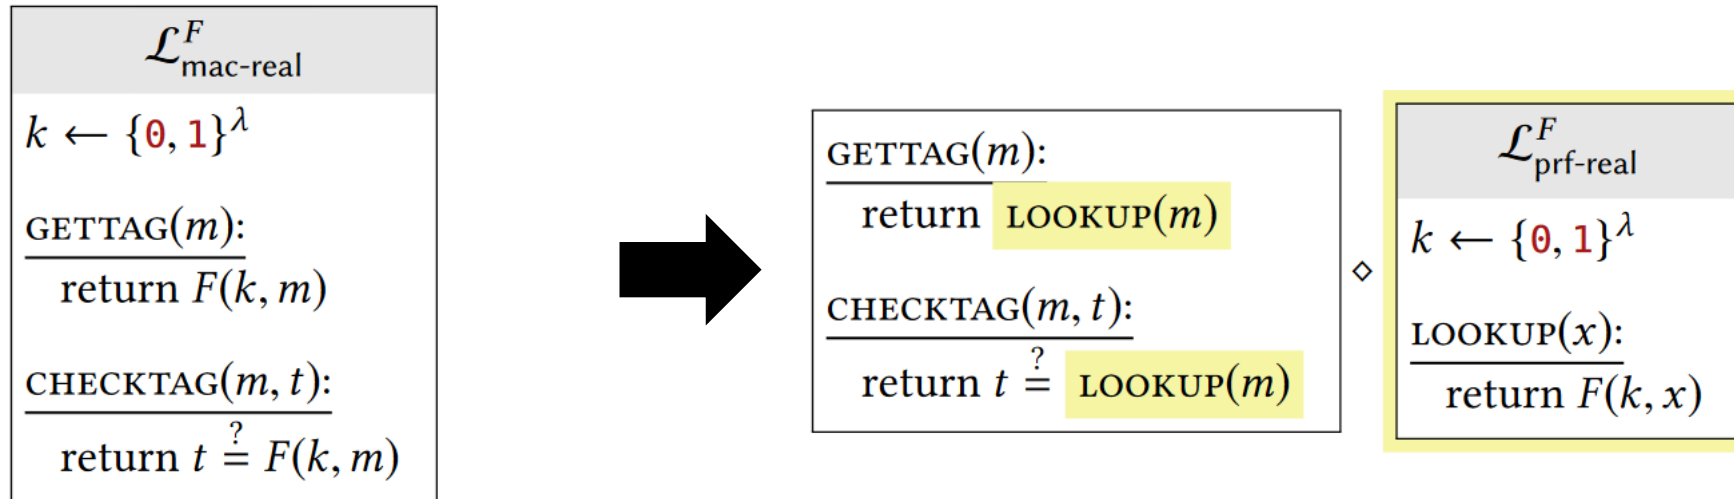
They **differ** only in code that is reachable when $bad = 1$.
From the **bad-event lemma**, we know that these two libraries are indistinguishable if $\Pr[bad = 1]$ is negligible. This happens with probability $1/2^\lambda$, which is indeed negligible.

A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$

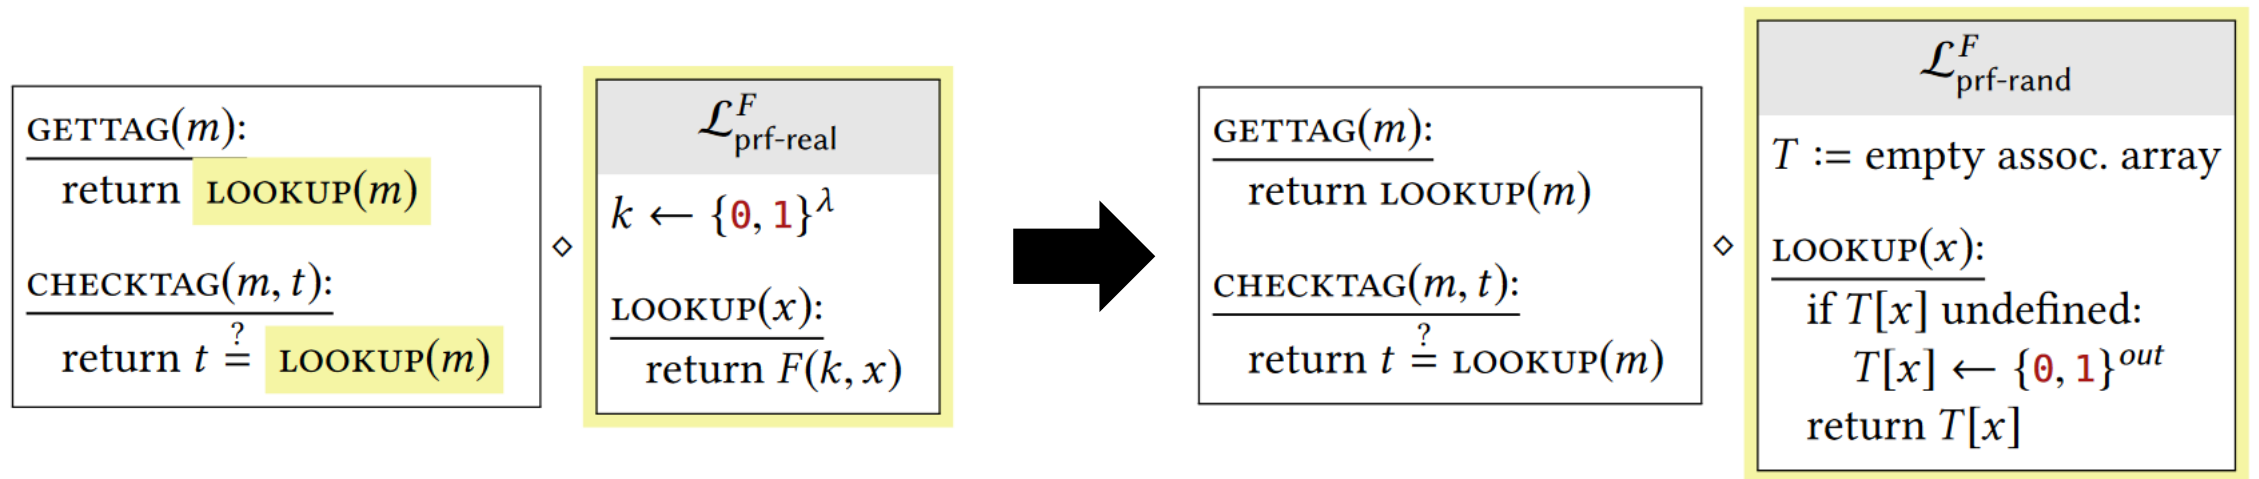


A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $MAC(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$

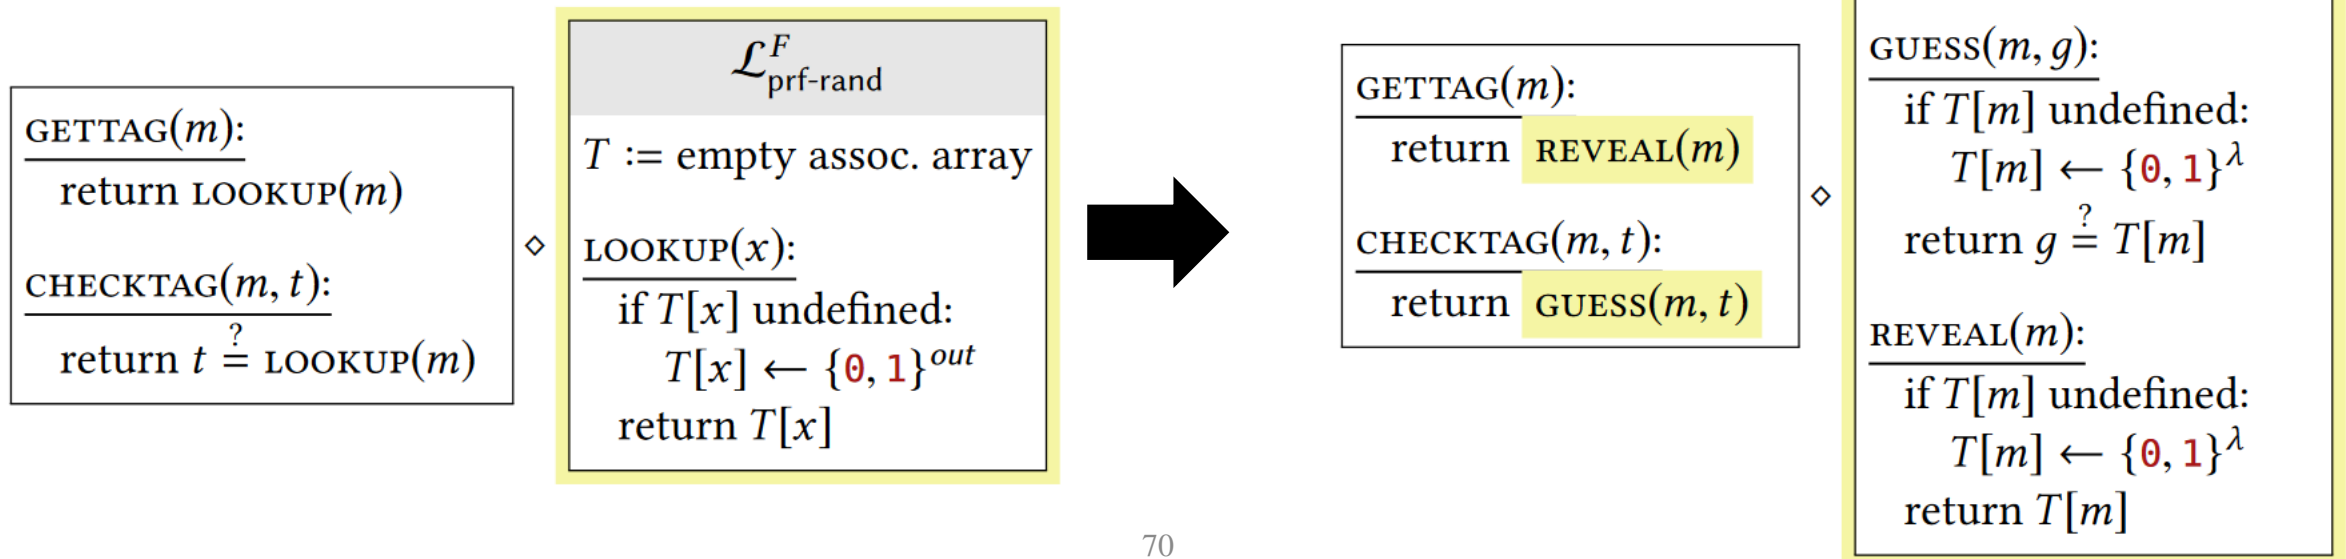


A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

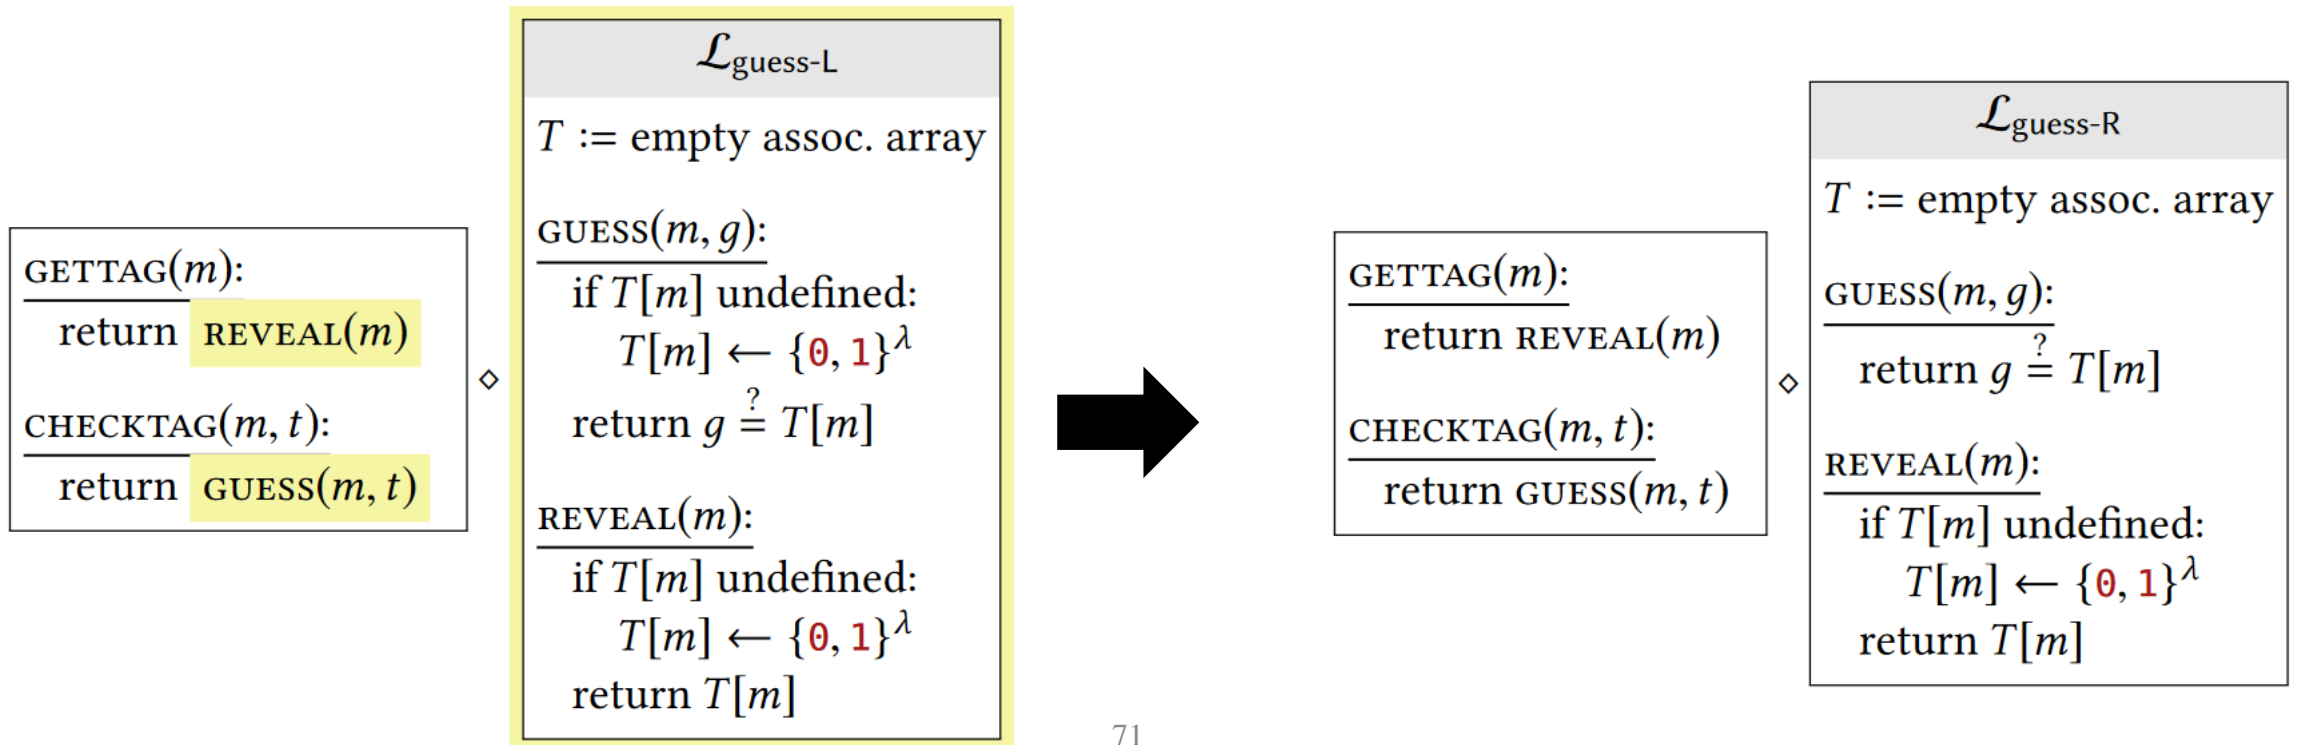
We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$



A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $MAC(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

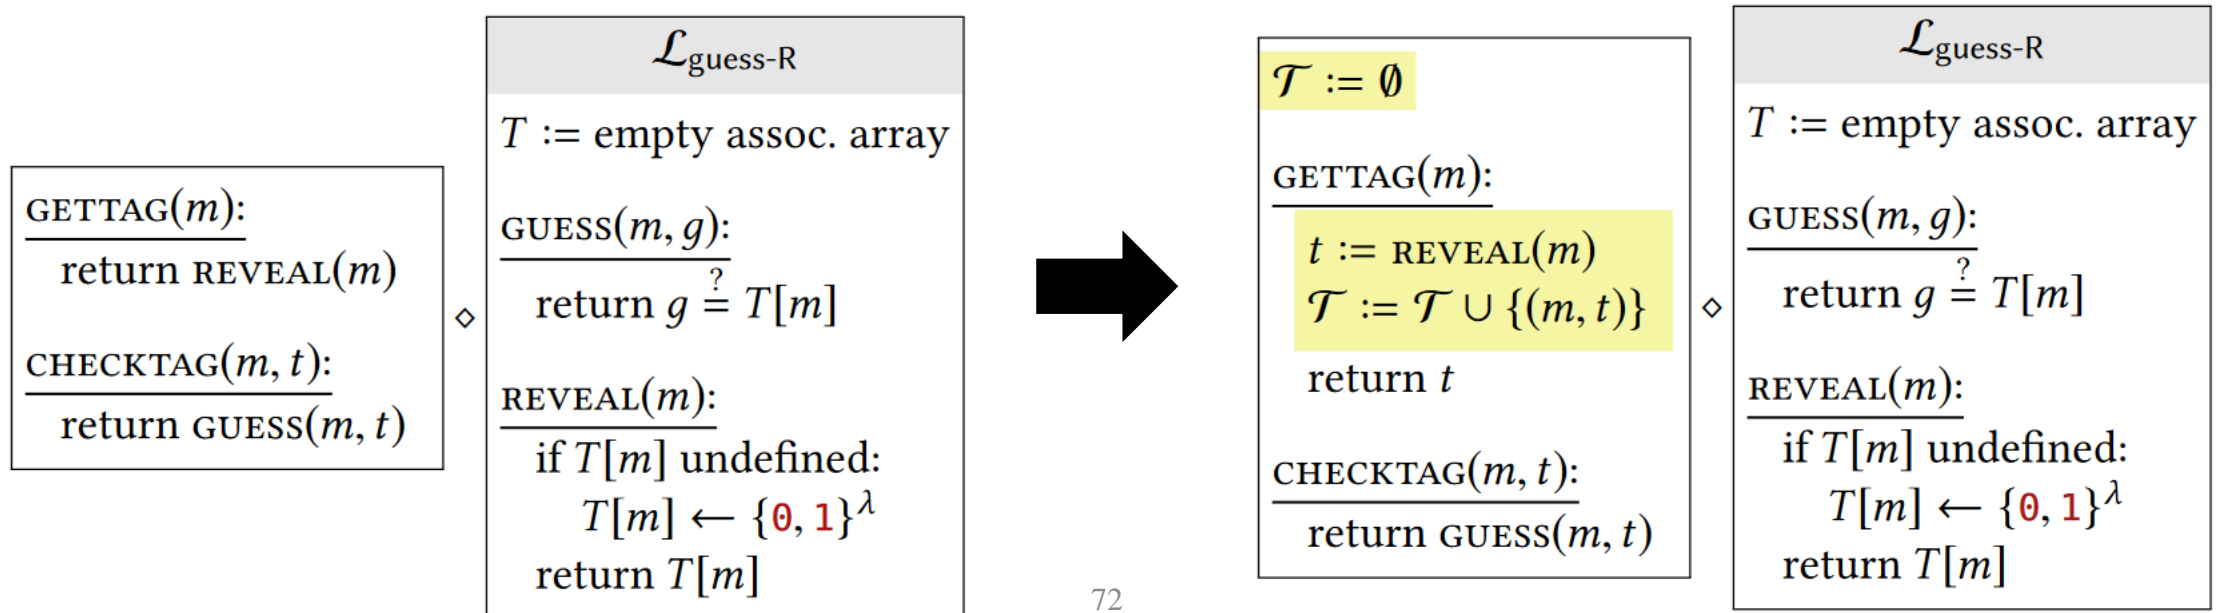


A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$

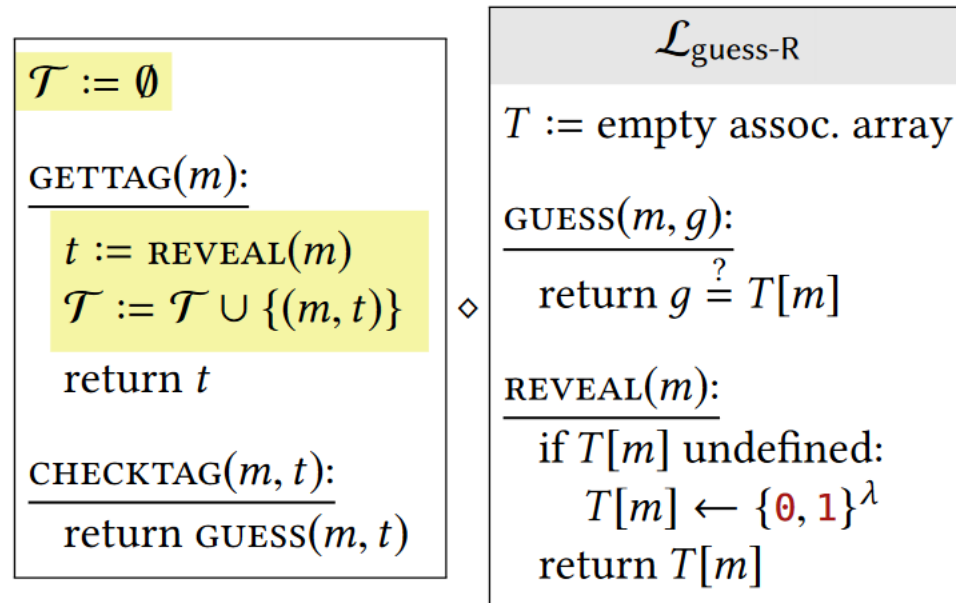


A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $MAC(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$



Suppose the calling program makes a call to CHECKTAG(m, t).

- Case 1: there was a **previous** call to GETTAG(m). $T[m]$ is defined and $(m, T[m])$ **already exists in \mathcal{T}** . CHECKTAG(m, t) will be $t =_? T(m)$.
- Case 2: there was **no previous** call to GETTAG(m). There is no value of the form $(m, ?)$ in \mathcal{T} . The call to GUESS(m, t) will return **false**.

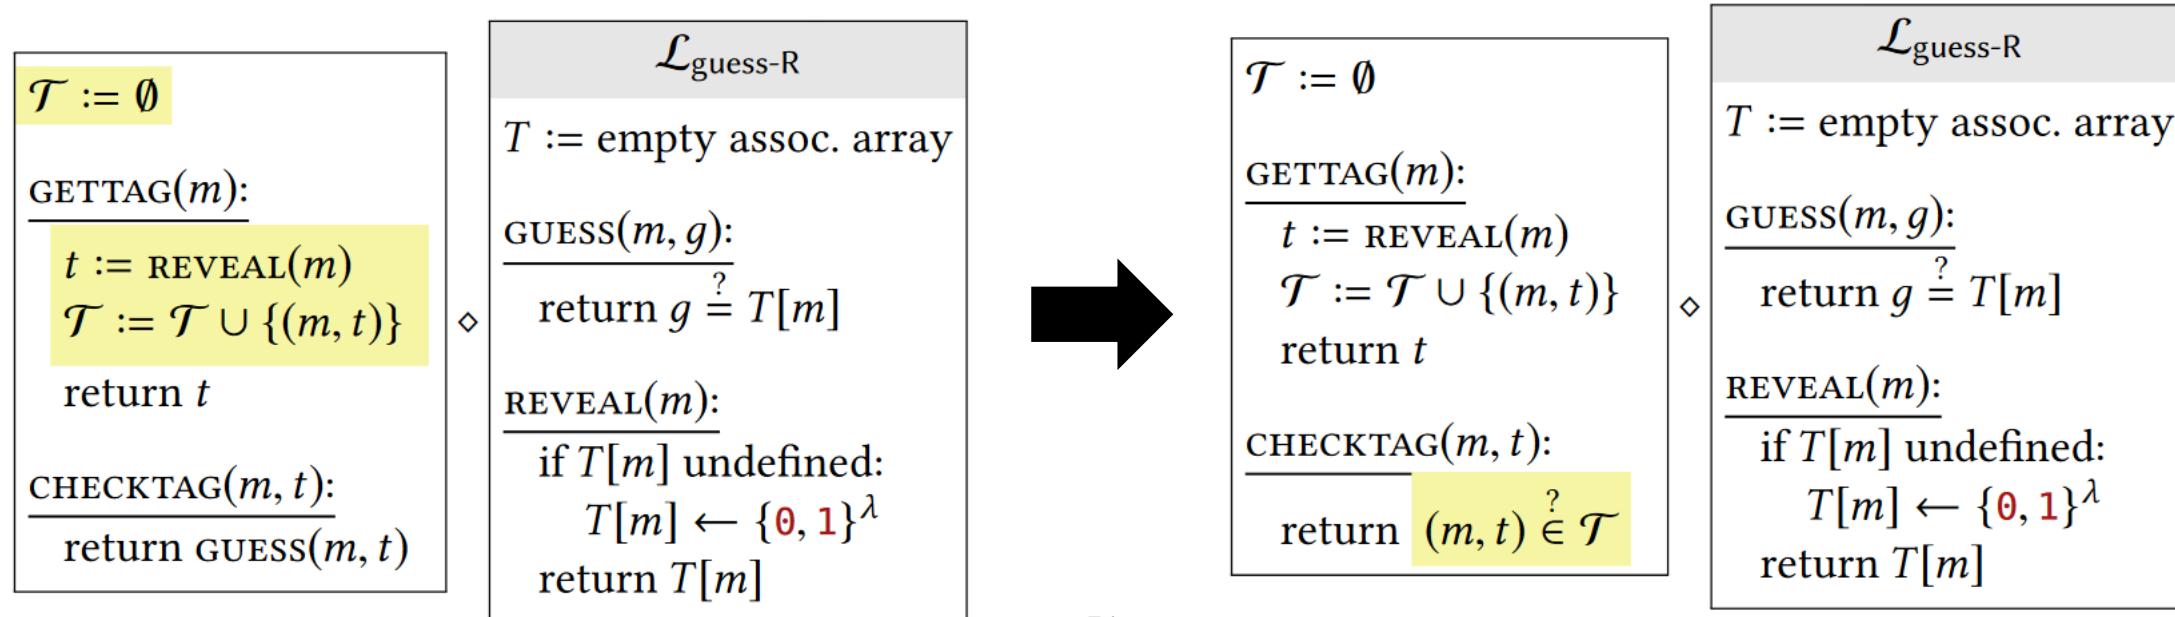
In both cases, the result of CHECKTAG(m, t) is true if and only if $(m, t) \in \mathcal{T}$.

A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$

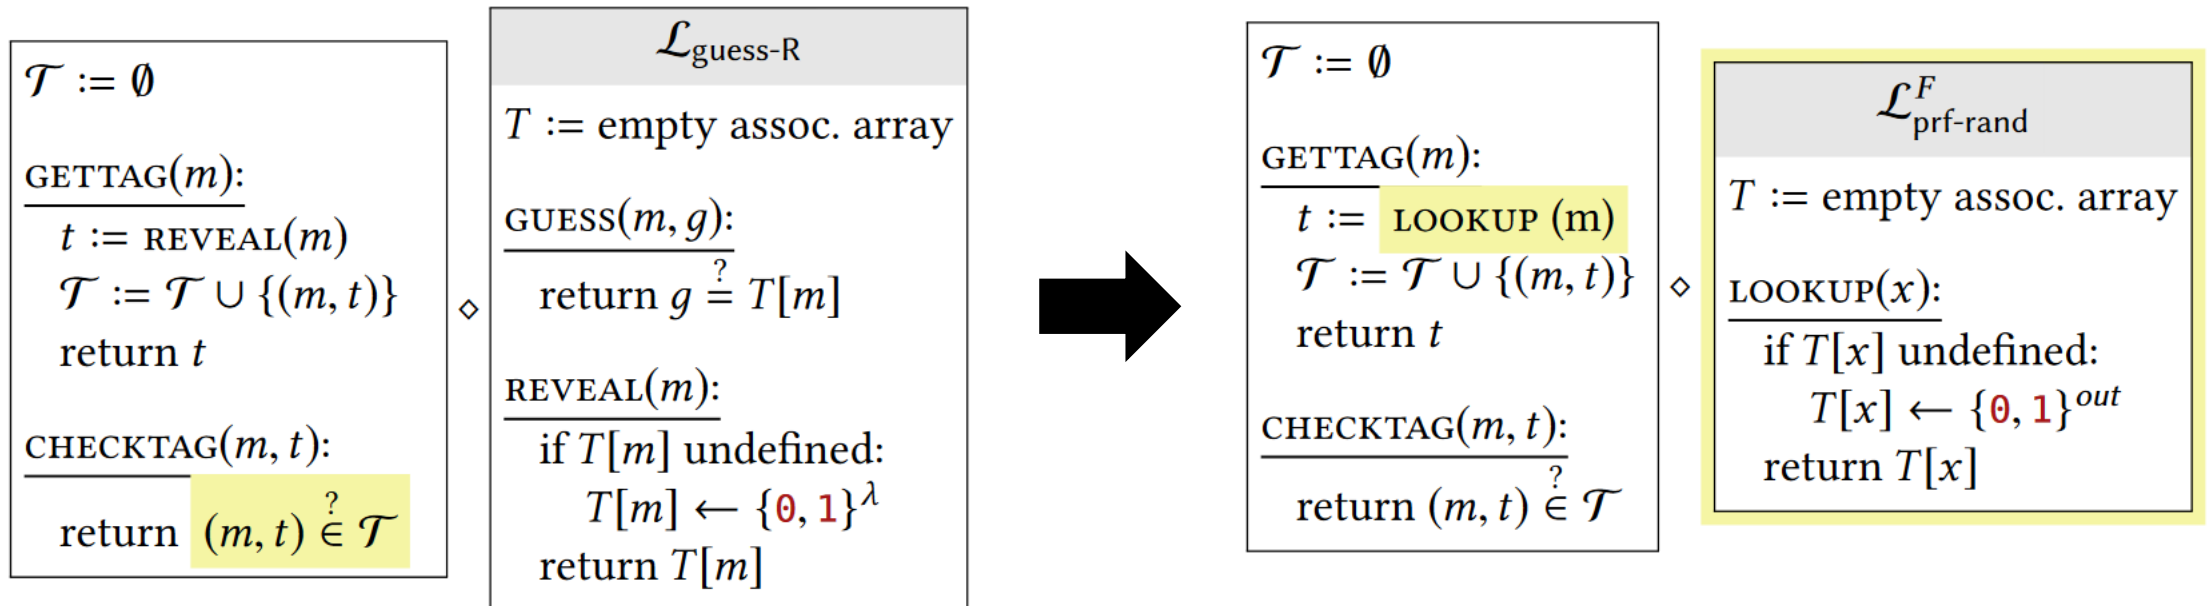


A PRF is a MAC

Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$



A PRF is a MAC

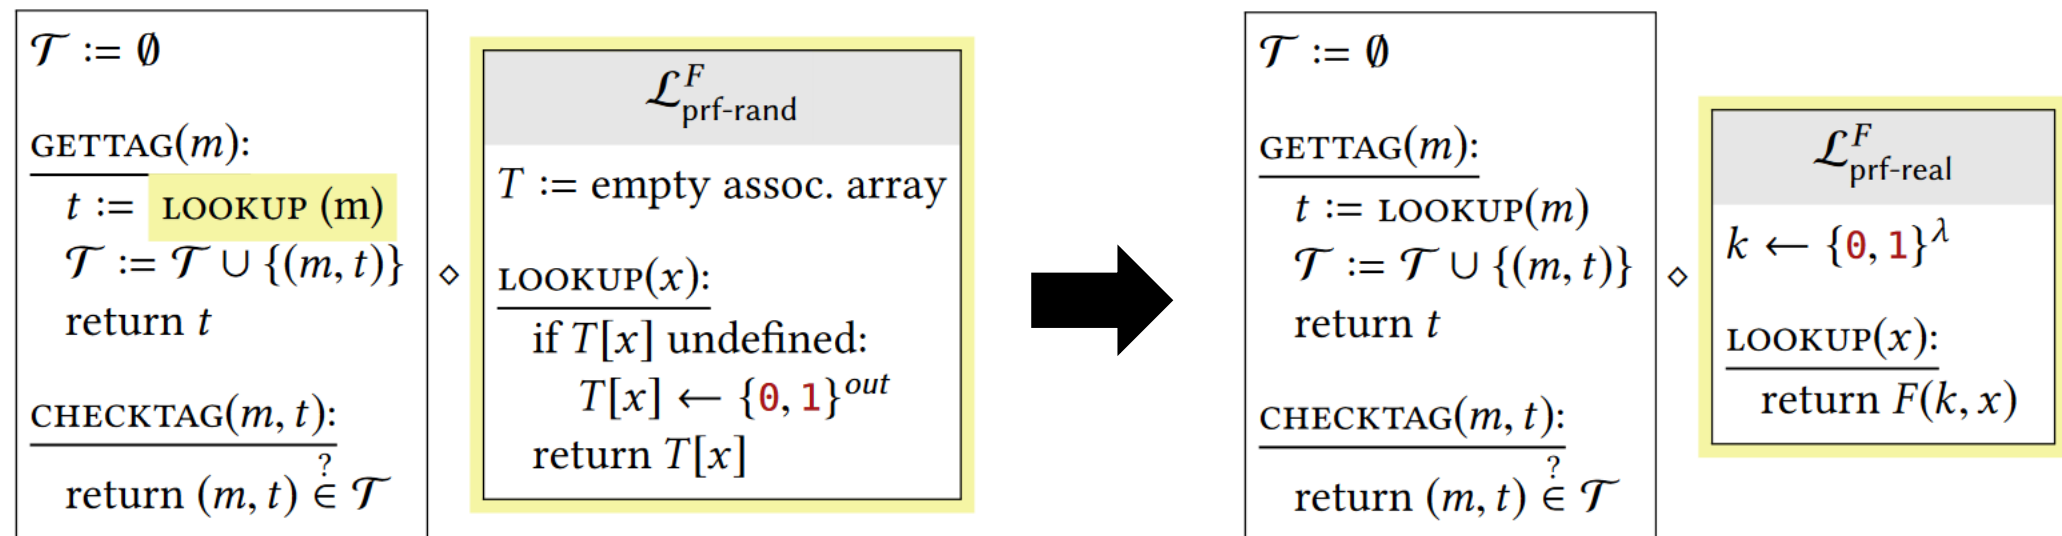
Claim Let F be a secure PRF with input length in and output length $out = \lambda$. Then the scheme $\text{MAC}(k, m) = F(k, m)$ is a secure MAC for message space $\{0, 1\}^{in}$.

proof

We show that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$

Inlining $\mathcal{L}_{\text{prf-real}}$ in the final hybrid, we see that the result is exactly $\mathcal{L}_{\text{mac-fake}}^F$.

Hence, we have shown that $\mathcal{L}_{\text{mac-real}}^F \approx \mathcal{L}_{\text{mac-fake}}^F$



If PRFs are MACs, why do we even need a definition for MACs?

- *Not every PRF is a MAC.* Only sufficiently long random values are hard to guess, so only PRFs with long outputs ($out > \lambda$) are MACs. It is perfectly reasonable to consider a PRF with short outputs.
- *Not every MAC is a PRF.* Something doesn't have to be uniformly random in order to be hard to guess.
 - $MAC'(k, m) = MAC(k, m) || 0^\lambda$.
- It is good practice to know whether you really need something that is pseudorandom or hard to guess.

MACs for Long Messages

- Using a PRF as a MAC is useful **only for short, fixed-length messages**, since most PRFs that exist in practice are limited to such inputs.
- Can we somehow extend a PRF to construct a MAC scheme for **long messages**, similar to how we used block cipher modes to construct encryption for long messages?

MACs for Long Messages

- Let F be a PRF with $in = out = \lambda$. Below is a MAC approach for messages of length 2λ .

ECBMAC($k, m_1 || m_2$):

$t_1 := F(k, m_1)$

$t_2 := F(k, m_2)$

return $t_1 || t_2$

- What is the problem?
- It does nothing to authenticate that m_1 is the first block but m_2 is the second one.

\mathcal{A} :

$t_1 || t_2 := \text{GETTAG}(\mathbf{0}^\lambda || \mathbf{1}^\lambda)$

return $\text{CHECKTAG}(\mathbf{1}^\lambda || \mathbf{0}^\lambda, t_2 || t_1)$

MACs for Long Messages

- Let F be a PRF with $in = \lambda + 1$ and $out = \lambda$. Below is a MAC approach for messages of length 2λ .

ECB++MAC($k, m_1 \| m_2$):

$t_1 := F(k, \mathbf{0} \| m_1)$

$t_2 := F(k, \mathbf{1} \| m_2)$

return $t_1 \| t_2$

- What is the problem?
- This construction doesn't authenticate the fact that this particular m_1 and m_2 belong together.

\mathcal{A} :

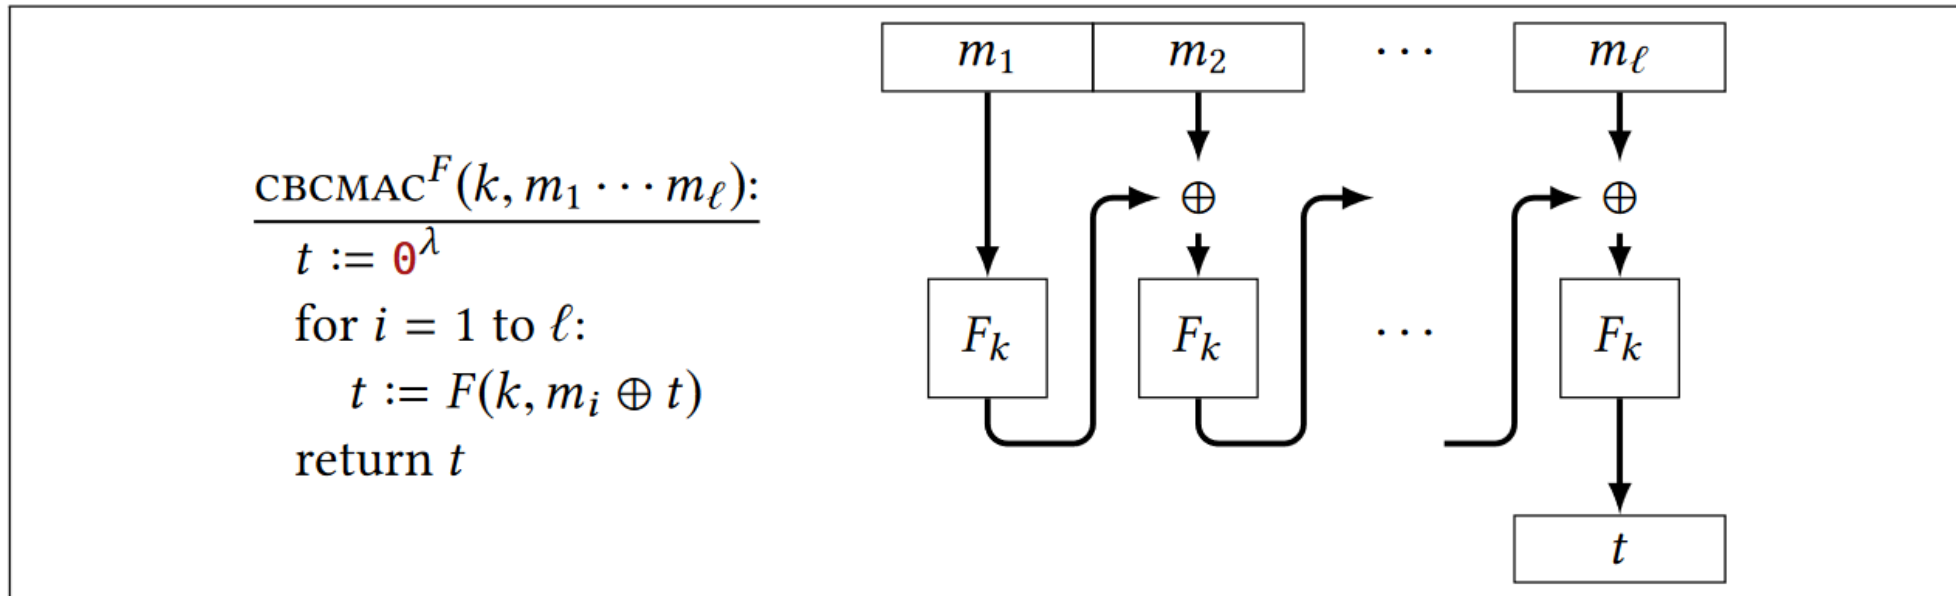
$t_1 \| t_2 := \text{GETTAG}(\mathbf{0}^{2\lambda})$

$t'_1 \| t'_2 := \text{GETTAG}(\mathbf{1}^{2\lambda})$

return $\text{CHECKTAG}(\mathbf{0}^\lambda \| \mathbf{1}^\lambda, t_1 \| t'_2)$

How to do it: CBC-MAC

- Let F be a PRF with $in = out = \lambda$. CBC-MAC refers to the following MAC scheme:



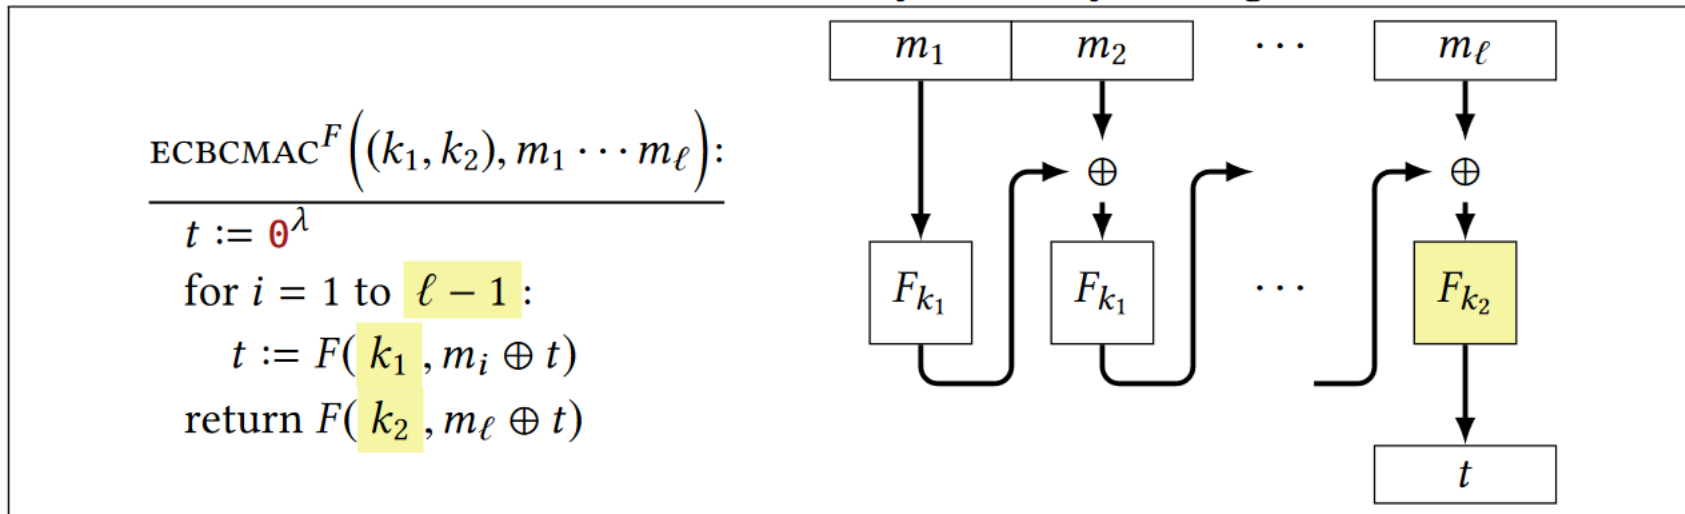
Unlike CBC encryption, CBC-MAC uses **no initialization vector** (or, you can think of it as using the all-zeroes IV), and it **outputs only the last block**.

How to do it: CBC-MAC

- **Theorem** If F is a secure PRF with $in = out = \lambda$, then for any **fixed** ℓ , CBC-MAC is a secure MAC when used with message space $\mathcal{M} = \{0, 1\}^{\lambda\ell}$.
- If you only ever authenticate 4-block messages, CBC-MAC is secure. If you only ever authenticate 24-block messages, CBC-MAC is secure. However, if you want to authenticate **both** 4-block and 24-block messages (i.e., **under the same key**), then CBC-MAC is not secure.

More Robust CBC-MAC

Let F be a PRF with $in = out = \lambda$. ECBC-MAC refers to the following scheme:



Theorem If F is a secure PRF with $in = out = \lambda$, then **ECBC-MAC** is a secure MAC for message space $\mathcal{M} = (\{0, 1\}^\lambda)^*$.

To extend ECBC-MAC to messages of any length (not necessarily a multiple of the block length), one can use a **padding** scheme as in the case of encryption.

Encrypt-Then-MAC

- Constructing a CCA-secure encryption scheme: add a MAC to a CPA-secure encryption scheme.
- Then the decryption algorithm can raise an error if the MAC is invalid, thereby ensuring that adversarially-generated (or adversarially-modified) ciphertexts are not accepted.
- There are several natural ways to combine a MAC and encryption scheme, but not all are secure!
- The **safest** way is known as encrypt-then-MAC:

Encrypt-Then-MAC

- Let E denote an encryption scheme, and M denote a MAC scheme where $E.\mathcal{C} \subseteq M.\mathcal{M}$ (i.e., the MAC scheme is capable of generating MACs of ciphertexts in the E scheme). Then let EtM denote the encrypt-then-MAC construction given below:

$\mathcal{K} = E.\mathcal{K} \times M.\mathcal{K}$	<u>Enc($(k_e, k_m), m$):</u>
$\mathcal{M} = E.\mathcal{M}$	$c := E.\text{Enc}(k_e, m)$
$\mathcal{C} = E.\mathcal{C} \times M.\mathcal{T}$	$t := M.\text{MAC}(k_m, c)$
	return (c, t)
<u>KeyGen:</u>	
$k_e \leftarrow E.\text{KeyGen}$	<u>Dec($(k_e, k_m), (c, t)$):</u>
$k_m \leftarrow M.\text{KeyGen}$	if $t \neq M.\text{MAC}(k_m, c)$:
return (k_e, k_m)	return err
	return $E.\text{Dec}(k_e, c)$

Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.

proof

$\mathcal{L}_{cca-L}^{EtM}$

$k_e \leftarrow E.\text{KeyGen}$
 $k_m \leftarrow M.\text{KeyGen}$

 $\mathcal{S} := \emptyset$

$\text{EAVESDROP}(m_L, m_R):$
 if $|m_L| \neq |m_R|$
 return null


$c := E.\text{Enc}(k_e, m_L)$
 $t \leftarrow M.\text{MAC}(k_m, c)$

 $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$
 return (c, t)

$\text{DEC}(c, t):$
 if $(c, t) \in \mathcal{S}$ return null

if $t \neq M.\text{MAC}(k_m, c):$
 return **err**

 return $E.\text{Dec}(k_e, c)$



$k_e \leftarrow E.\text{KeyGen}$
 $\mathcal{S} := \emptyset$

$\text{EAVESDROP}(m_L, m_R):$
 if $|m_L| \neq |m_R|$
 return null
 $c := E.\text{Enc}(k_e, m_L)$

$t := \text{GETTAG}(c)$

 $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$
 return (c, t)

$\text{DEC}(c, t):$
 if $(c, t) \in \mathcal{S}$
 return null

if **not** $\text{CHECKTAG}(c, t):$
 return **err**

 return $E.\text{Dec}(k_e, c)$



$\mathcal{L}_{mac-real}^M$

$k_m \leftarrow M.\text{KeyGen}$

$\text{GETTAG}(c):$
 return $M.\text{MAC}(k_m, c)$

$\text{CHECKTAG}(c, t):$
 return $t \stackrel{?}{=} M.\text{MAC}(k_m, c)$

Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.

```

 $k_e \leftarrow E.\text{KeyGen}$ 
 $\mathcal{S} := \emptyset$ 

EAVESDROP( $m_L, m_R$ ):
  if  $|m_L| \neq |m_R|$ 
    return null
   $c := E.\text{Enc}(k_e, m_L)$ 
   $t := \text{GETTAG}(c)$ 
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$ 
  return  $(c, t)$ 

DEC( $c, t$ ):
  if  $(c, t) \in \mathcal{S}$ 
    return null
  if not CHECKTAG( $c, t$ ):
    return err
  return  $E.\text{Dec}(k_e, c)$ 

```

◇

```

 $\mathcal{L}_{\text{mac-real}}^M$ 

 $k_m \leftarrow M.\text{KeyGen}$ 

GETTAG( $c$ ):
  return  $M.\text{MAC}(k_m, c)$ 

CHECKTAG( $c, t$ ):
  return  $t \stackrel{?}{=} M.\text{MAC}(k_m, c)$ 

```



```

 $k_e \leftarrow E.\text{KeyGen}$ 
 $\mathcal{S} := \emptyset$ 

EAVESDROP( $m_L, m_R$ ):
  if  $|m_L| \neq |m_R|$ 
    return null
   $c := E.\text{Enc}(k_e, m_L)$ 
   $t := \text{GETTAG}(c)$ 
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$ 
  return  $(c, t)$ 

DEC( $c, t$ ):
  if  $(c, t) \in \mathcal{S}$ 
    return null
  if not CHECKTAG( $c, t$ ):
    return err
  return  $E.\text{Dec}(k_e, c)$ 

```

◇

```

 $\mathcal{L}_{\text{mac-fake}}^M$ 

 $k_m \leftarrow M.\text{KeyGen}$ 
 $\mathcal{T} := \emptyset$ 

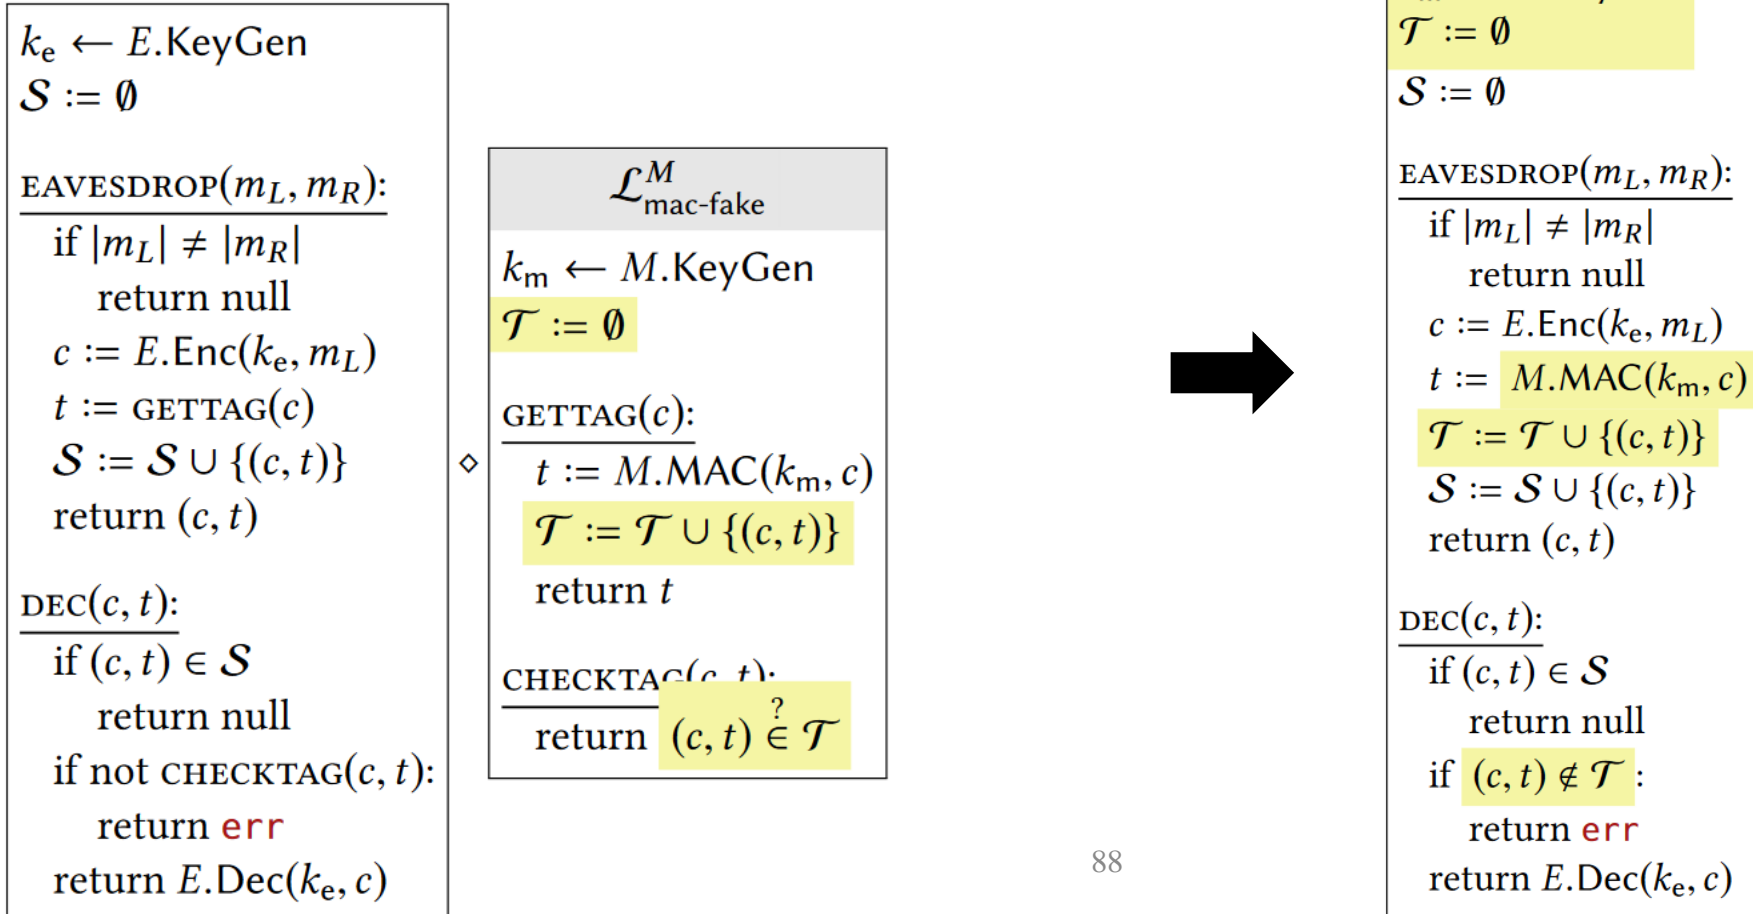
GETTAG( $c$ ):
   $t := M.\text{MAC}(k_m, c)$ 
   $\mathcal{T} := \mathcal{T} \cup \{(c, t)\}$ 
  return  $t$ 

CHECKTAG( $c, t$ ):
  return  $(c, t) \stackrel{?}{\in} \mathcal{T}$ 

```

Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.



Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.

proof

```
 $k_e \leftarrow E.\text{KeyGen}$   
 $k_m \leftarrow M.\text{KeyGen}$   
 $\mathcal{T} := \emptyset$   
 $\mathcal{S} := \emptyset$   
  
EAVESDROP( $m_L, m_R$ ):  
  if  $|m_L| \neq |m_R|$   
    return null  
   $c := E.\text{Enc}(k_e, m_L)$   
   $t := M.\text{MAC}(k_m, c)$   
   $\mathcal{T} := \mathcal{T} \cup \{(c, t)\}$   
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$   
  return  $(c, t)$   
  
DEC( $c, t$ ):  
  if  $(c, t) \in \mathcal{S}$   
    return null  
  if  $(c, t) \notin \mathcal{T}$ :  
    return err  
  return  $E.\text{Dec}(k_e, c)$ 
```




```
 $k_e \leftarrow E.\text{KeyGen}$   
 $k_m \leftarrow M.\text{KeyGen}$   
 $\mathcal{S} := \emptyset$   
  
EAVESDROP( $m_L, m_R$ ):  
  if  $|m_L| \neq |m_R|$   
    return null  
   $c := E.\text{Enc}(k_e, m_L)$   
   $t := M.\text{MAC}(k_m, c)$   
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$   
  return  $(c, t)$   
  
DEC( $c, t$ ):  
  if  $(c, t) \in \mathcal{S}$   
    return null  
  if  $(c, t) \notin \mathcal{S}$ :  
    return err  
  // unreachable
```

Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.

proof

```
 $k_e \leftarrow E.\text{KeyGen}$   
 $k_m \leftarrow M.\text{KeyGen}$   
 $\mathcal{S} := \emptyset$   
  
EAVESDROP( $m_L, m_R$ ):  
  if  $|m_L| \neq |m_R|$   
    return null  
   $c := E.\text{Enc}(k_e, m_L)$   
   $t := M.\text{MAC}(k_m, c)$   
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$   
  return  $(c, t)$   
  
DEC( $c, t$ ):  
  if  $(c, t) \in \mathcal{S}$   
    return null  
  if  $(c, t) \notin \mathcal{S}$ :  
    return err  
  // unreachable
```



```
 $k_m \leftarrow M.\text{KeyGen}$   
 $\mathcal{S} := \emptyset$   
  
EAVESDROP( $m_L, m_R$ ):  
  if  $|m_L| \neq |m_R|$   
    return null  
   $c := \text{CPA.EAVESDROP}(m_L, m_R)$   
   $t := M.\text{MAC}(k_m, c)$   
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$   
  return  $(c, t)$   
  
DEC( $c, t$ ):  
  if  $(c, t) \in \mathcal{S}$   
    return null  
  if  $(c, t) \notin \mathcal{S}$ :  
    return err
```



$\mathcal{L}_{\text{cpa-L}}^E$
$k_e \leftarrow E.\text{KeyGen}$
<u>CPA.EAVESDROP(m_L, m_R):</u>
$c := E.\text{Enc}(k_e, m_L)$
return c

Encrypt-Then-MAC

Claim If E has CPA security and M is a secure MAC, then EtM Construction has CCA security.

proof

```
 $k_m \leftarrow M.\text{KeyGen}$   
 $\mathcal{S} := \emptyset$   
  
EAVESDROP( $m_L, m_R$ ):  
  if  $|m_L| \neq |m_R|$   
    return null  
   $c := \text{CPA.EAVESDROP}(m_L, m_R)$   
   $t := M.\text{MAC}(k_m, c)$   
   $\mathcal{S} := \mathcal{S} \cup \{(c, t)\}$   
  return  $(c, t)$   
  
DEC( $c, t$ ):  
  if  $(c, t) \in \mathcal{S}$   
    return null  
  if  $(c, t) \notin \mathcal{S}$ :  
    return err
```

◇

```
 $\mathcal{L}_{\text{cpa-L}}^E$   
  
 $k_e \leftarrow E.\text{KeyGen}$   
  
CPA.EAVESDROP( $m_L, m_R$ ):  
   $c := E.\text{Enc}(k_e, m_L)$   
  return  $c$ 
```

We have now reached the **half-way** point of the proof.

The proof proceeds by replacing $\mathcal{L}_{\text{cpa-L}}$ with $\mathcal{L}_{\text{cpa-R}}$ (so that m_R rather than m_L is encrypted), applying the same modifications as before (but in reverse order), to finally arrive at $\mathcal{L}_{\text{cca-R}}$.

Hash Functions

Cryptographic Hash Function

- $H: \{0, 1\}^* \rightarrow \{0, 1\}^n$
- **Collision resistance.** It should be **hard** to compute any collision x, x' such that $H(x) = H(x')$.
- **Second-preimage resistance.** Given x , it should be **hard** to compute any collision involving x . In other words, it should be **hard** to compute $x' \neq x$ such that $H(x) = H(x')$.

Brute Force Attacks on Hash Functions

- Let's make a simplifying assumption that for any $m > n$, the following distribution is roughly uniform over $\{0, 1\}^n$:

$$x \leftarrow \{0, 1\}^m$$
$$\text{return } H(x)$$

Brute Force Attacks on Hash Functions

- Let's make a simplifying assumption that for any $m > n$, the following distribution is roughly uniform over $\{0, 1\}^n$:

```
Collision brute force:  
 $\mathcal{A}_{\text{cr}}()$ :  
  for  $i = 1, \dots$ :  
     $x_i \leftarrow \{0, 1\}^m$   
     $y_i := H(x_i)$   
    if there is some  $j < i$  with  $x_i \neq x_j$   
      but  $y_i = y_j$ :  
      return  $(x_i, x_j)$ 
```

Since each $y_i \in \{0, 1\}^n$, the probability of finding a repeated value after q times through the main loop is roughly $\Theta(q^2/2^n)$ by the birthday bound. While in the worst case it could take 2^n steps to find a collision in H , the birthday bound implies that it takes only $2^{\frac{n}{2}}$ attempts to find a collision with 99% probability (or any constant probability).

Brute Force Attacks on Hash Functions

- Let's make a simplifying assumption that for any $m > n$, the following distribution is roughly uniform over $\{0, 1\}^n$:

Second preimage brute force:

$\mathcal{A}_{2\text{pi}}(x)$:

while true:

$x' \leftarrow \{0, 1\}^m$

$y' := H(x')$

if $y' = H(x)$: return x'

It will therefore take $\Theta(2^n)$ attempts in expectation to terminate successfully.

Brute Force Attacks on Hash Functions

- Let's make a simplifying assumption that for any $m > n$, the following distribution is roughly uniform over $\{0, 1\}^n$:

Collision brute force:

```
 $\mathcal{A}_{\text{cr}}()$ :  
for  $i = 1, \dots$ :  
   $x_i \leftarrow \{0, 1\}^m$   
   $y_i := H(x_i)$   
  if there is some  $j < i$  with  $x_i \neq x_j$   
    but  $y_i = y_j$ :  
    return  $(x_i, x_j)$ 
```

Second preimage brute force:

```
 $\mathcal{A}_{2\text{pi}}(x)$ :  
while true:  
   $x' \leftarrow \{0, 1\}^m$   
   $y' := H(x')$   
  if  $y' = H(x)$ : return  $x'$ 
```

This [difference](#) explains why you will typically see cryptographic hash functions in practice that have 256- to 512-bit output length (but not 128-bit output length), while you only typically see block ciphers with 128-bit or 256-bit keys.

In order to make brute force attacks cost 2^n , [a block cipher needs only an \$n\$ -bit key](#) while [a collision-resistant hash function needs a \$2n\$ -bit output](#).

Hash Function Security in Theory

TEST(x, x'):

if $x \neq x'$ and $H(x) = H(x')$: return true
else: return false

\approx

TEST(x, x'):

return false

- What is the problem?
- With exponential time, we could find such an (x, x') pair and **write down the code of an attacker** (the values x and x' are **hard-coded** into \mathcal{A}):

\mathcal{A} :

return TEST(x, x')

Hash Function Security in Theory

The way around this technical issue is to **introduce some randomness into the libraries and into the inputs of H** . We define hash functions to take two arguments: a randomly chosen, public value s called a **salt**, and an adversarially chosen input x .

Definition A hash function H is collision-resistant if $\mathcal{L}_{\text{cr-real}}^{\mathcal{H}} \approx \mathcal{L}_{\text{cr-fake}}^{\mathcal{H}}$, where:

$\mathcal{L}_{\text{cr-real}}^{\mathcal{H}}$
$s \leftarrow \{0, 1\}^{\lambda}$
<u>GETSALT():</u> return s
<u>TEST($x, x' \in \{0, 1\}^*$):</u> if $x \neq x'$ and $H(s, x) = H(s, x')$: return true return false

$\mathcal{L}_{\text{cr-fake}}^{\mathcal{H}}$
$s \leftarrow \{0, 1\}^{\lambda}$
<u>GETSALT():</u> return s
<u>TEST($x, x' \in \{0, 1\}^*$):</u> return false

Think of salt as an extra value that “personalizes” the hash function for a given application.

Salts in Practice

- Hash functions are **commonly** used to store passwords. A server may store user records of the form (username, $h = H(\text{password})$). When a user attempts to login with password p' , the server computes $H(p')$ and compares it to h .
- Storing hashed passwords means that, in the event that the password file is stolen, **an attacker would need to find a preimage of h in order to impersonate the user.**
- Best practice is to use a **separate salt** for each user. Instead of storing (username, $H(\text{password})$), choose a random salt s for each user and store (username, s , $H(s, \text{password})$).

Merkle-Damgård Construction

- Building a hash function, especially one that accepts inputs of **arbitrary length**, seems like a challenging task.
- Instead of a full-fledged hash function, imagine that we had a collision-resistant function whose inputs were of a **single fixed length**, but **longer than its outputs**.
- In other words, $h: \{0, 1\}^{n+t} \rightarrow \{0, 1\}^n$, where $t > 0$. We call such an h a **compression** function.
- We want to build a full-fledged hash function (**supporting inputs of arbitrary length**) out of such a compression function

Merkle-Damgård Construction

- Let $h : \{0, 1\}^{n+t} \rightarrow \{0, 1\}^n$ be a compression function. Then the Merkle-Damgård transformation of h is $MD_h : \{0, 1\}^* \rightarrow \{0, 1\}^n$, where:

MDPAD_t(x)

$\ell := |x|$, as length- t binary number

while $|x|$ not a multiple of t :

$x := x \parallel \mathbf{0}$

return $x \parallel \ell$

MD_h(x):

$x_1 \parallel \dots \parallel x_{k+1} := \text{MDPAD}_t(x)$


// each x_i is t bits

$y_0 := \mathbf{0}^n$

for $i = 1$ to $k + 1$:

$y_i := h(y_{i-1} \parallel x_i)$

output y_{k+1}



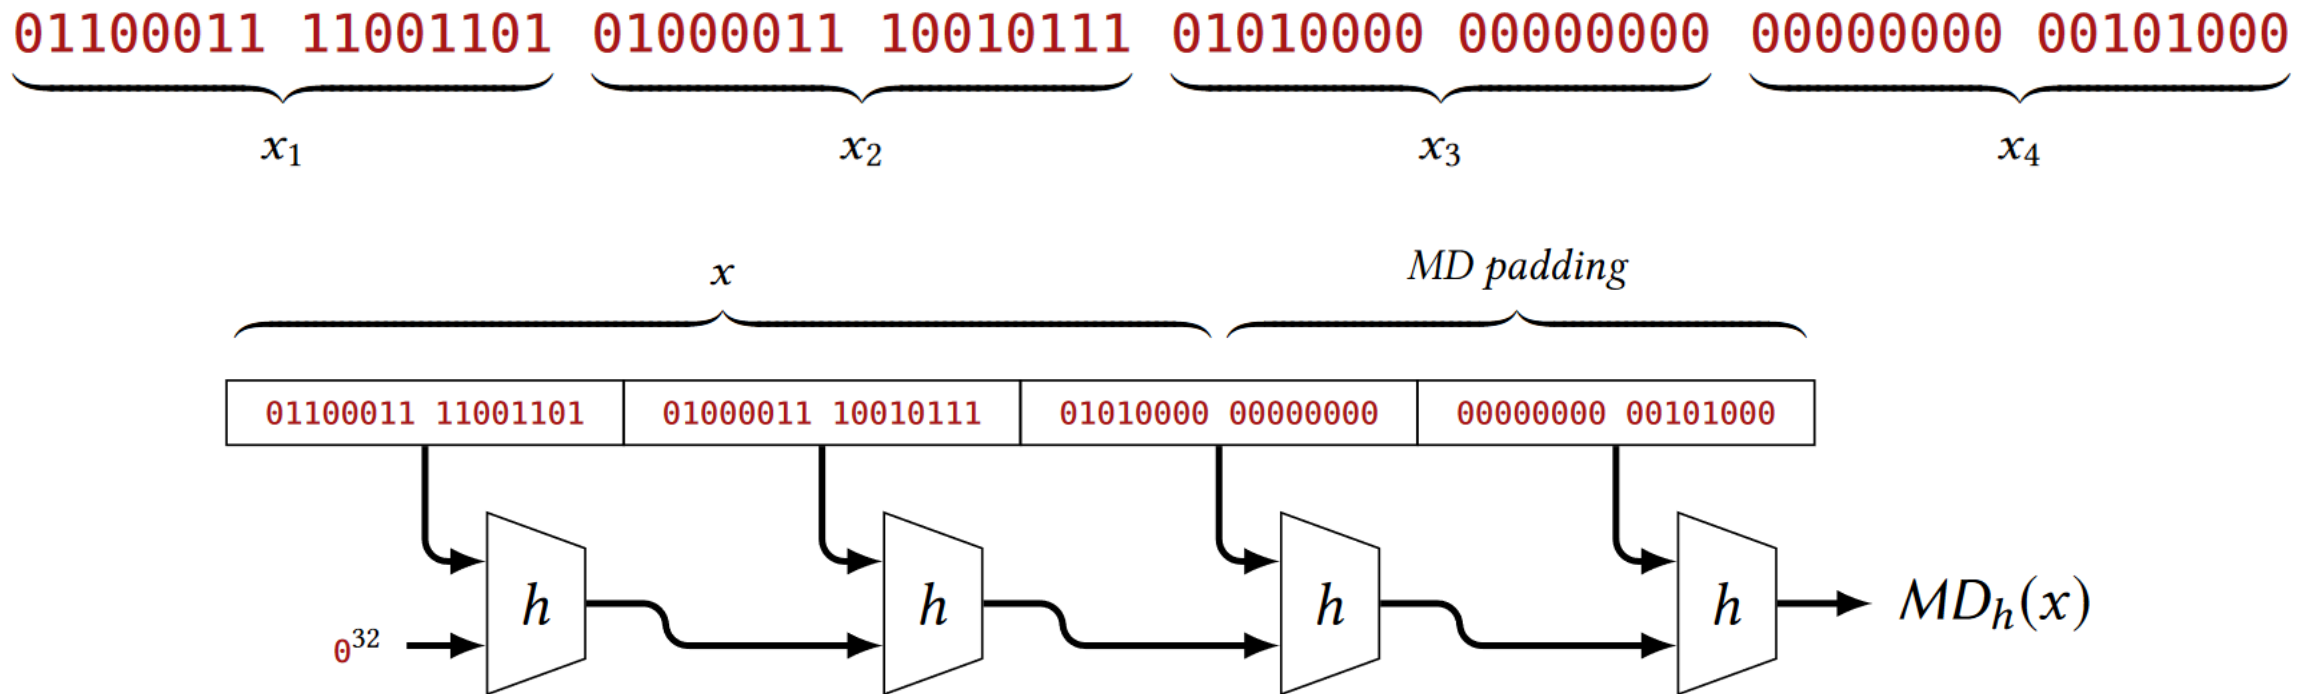
Merkle-Damgård Construction

- Suppose we have a compression function $h : \{0, 1\}^{48} \rightarrow \{0, 1\}^{32}$, so that $t = 16$.
- To compute the hash of the following 5-byte (40-bit) string:
 $x = 01100011\ 11001101\ 01000011\ 10010111\ 01010000$
- We must first pad x appropriately (MDPAD(x)):
 - Since x is not a multiple of $t = 16$ bits, we need to add 8 bits to make it so.
 - Since $|x| = 40$, we need to add an extra 16-bit block that encodes the number 40 in binary (101000)

$\underbrace{01100011\ 11001101}_{x_1} \underbrace{01000011\ 10010111}_{x_2} \underbrace{01010000\ 00000000}_{x_3} \underbrace{00000000\ 00101000}_{x_4}$

Merkle-Damgård Construction

- Suppose we have a compression function $h : \{0, 1\}^{48} \rightarrow \{0, 1\}^{32}$, so that $t = 16$.



Merkle-Damgård Construction

Claim Suppose h is a compression function and MD_h is the Merkle-Damgård construction applied to h . Given a collision x, x' in MD_h , it is easy to find a collision in h . In other words, if it is hard to find a collision in h , then it must also be hard to find a collision in MD_h .

proof

Suppose that x, x' are a **collision** under MD_h . Define the values x_1, \dots, x_{k+1} and y_1, \dots, y_{k+1} as in the computation of $MD_h(x)$.

Similarly, define $x'_1, \dots, x'_{k'+1}$ and $y'_1, \dots, y'_{k'+1}$ as in the computation of $MD_h(x')$. Note that, in general, k may not equal k' .

Merkle-Damgård Construction

Claim Suppose h is a compression function and MD_h is the Merkle-Damgård construction applied to h . Given a collision x, x' in MD_h , it is easy to find a collision in h . In other words, if it is hard to find a collision in h , then it must also be hard to find a collision in MD_h .

proof

$$\begin{aligned} MD_h(x) &= y_{k+1} = h(y_k \| x_{k+1}) \\ MD_h(x') &= y'_{k'+1} = h(y'_{k'} \| x'_{k'+1}) \end{aligned}$$

Since we are assuming $MD_h(x) = MD_h(x')$, we have $y_{k+1} = y'_{k'+1}$.

- Case 1: If $|x| \neq |x'|$, then the padding blocks x_{k+1} and $x'_{k'+1}$ which encode $|x|$ and $|x'|$ are not equal. Hence we have $y_k \| x_{k+1} \neq y'_{k'} \| x'_{k'+1}$, so $y_k \| x_{k+1}$ and $y'_{k'} \| x'_{k'+1}$ are a collision under h and we are done.

Merkle-Damgård Construction

Claim Suppose h is a compression function and MD_h is the Merkle-Damgård construction applied to h . Given a collision x, x' in MD_h , it is easy to find a collision in h . In other words, if it is hard to find a collision in h , then it must also be hard to find a collision in MD_h .

proof

$$\begin{aligned} MD_h(x) &= y_{k+1} = h(y_k \| x_{k+1}) \\ MD_h(x') &= y'_{k'+1} = h(y'_{k'} \| x'_{k'+1}) \end{aligned}$$

Since we are assuming $MD_h(x) = MD_h(x')$, we have $y_{k+1} = y'_{k'+1}$.

- Case 2: If $|x| = |x'|$, then x and x' are broken into the **same** number of blocks.

$$\begin{aligned} y_{k+1} &= h(y_k \| x_{k+1}) \\ &= \\ y'_{k+1} &= h(y'_k \| x'_{k+1}) \end{aligned}$$

If $y_k \| x_{k+1}$ and $y'_k \| x'_{k+1}$ are **not equal**, then they are a collision under h . Otherwise, $y_k = y'_k$ (and $x_{k+1} = x'_{k+1}$, then

$$y_k = h(y_{k-1} \| x_k) = y'_k = h(y'_{k-1} \| x'_k)$$

We **must have** $y_i \neq y'_i$ for some i . Otherwise, it implies that $x_i = x'_i$ for all i . Contradiction.

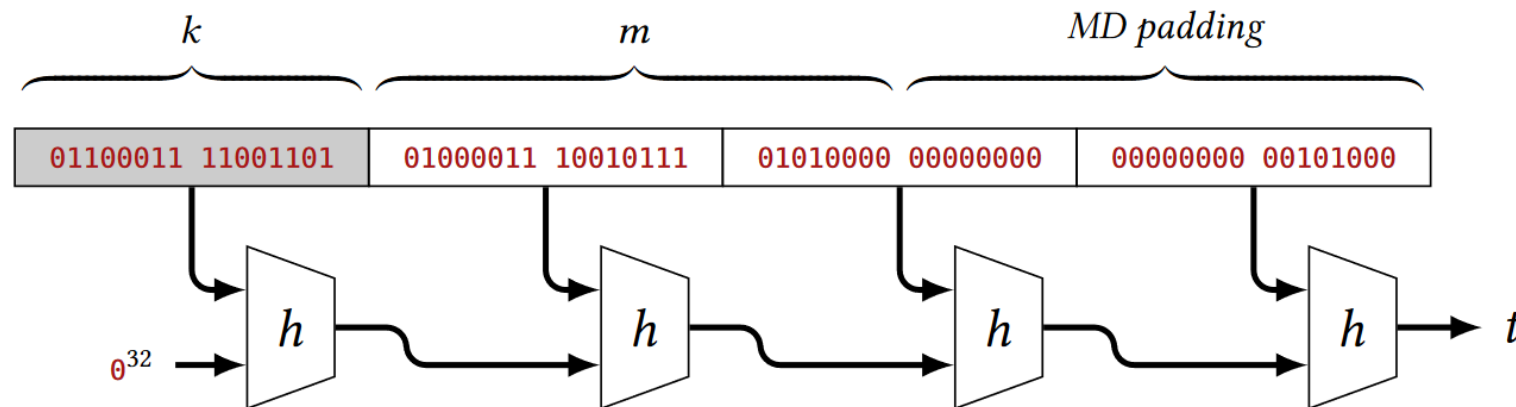
Hash Functions vs. MACs: Length-Extension Attacks

- What happens when we make the salt **private**?
- A hash function with secret salt most closely resembles a MAC. So, do we get a secure MAC by using a hash function with private salt?
- Unfortunately, the answer is **no** in general (although it can be yes in some cases, depending on the hash function). In particular, the method is insecure when H is constructed using the Merkle-Damgård approach.

Knowing $H(x)$ allows you to predict the hash of any string that begin with $MDPAD(x)$.

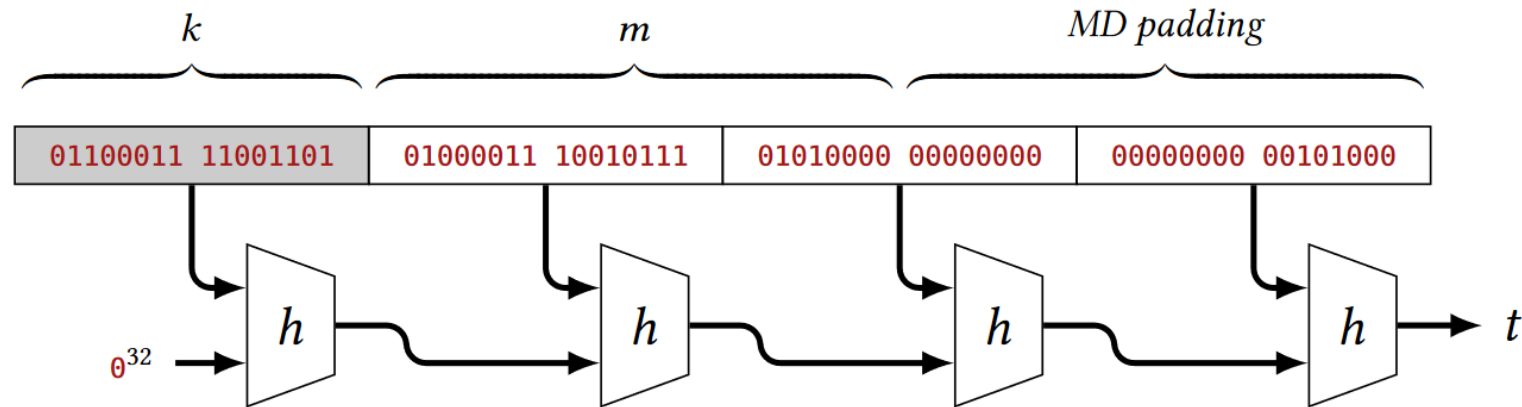
Hash Functions vs. MACs: Length-Extension Attacks

- If we use the construction $MAC(k, m) = H(k || m)$ as a MAC:
 - Suppose the MAC key is chosen as $k = 01100011\ 11001101$, and an attacker sees the MAC tag t of the message $m = 01000011\ 10010111\ 01010000$. Then $t = H(k || m)$ corresponds exactly to the example from before:

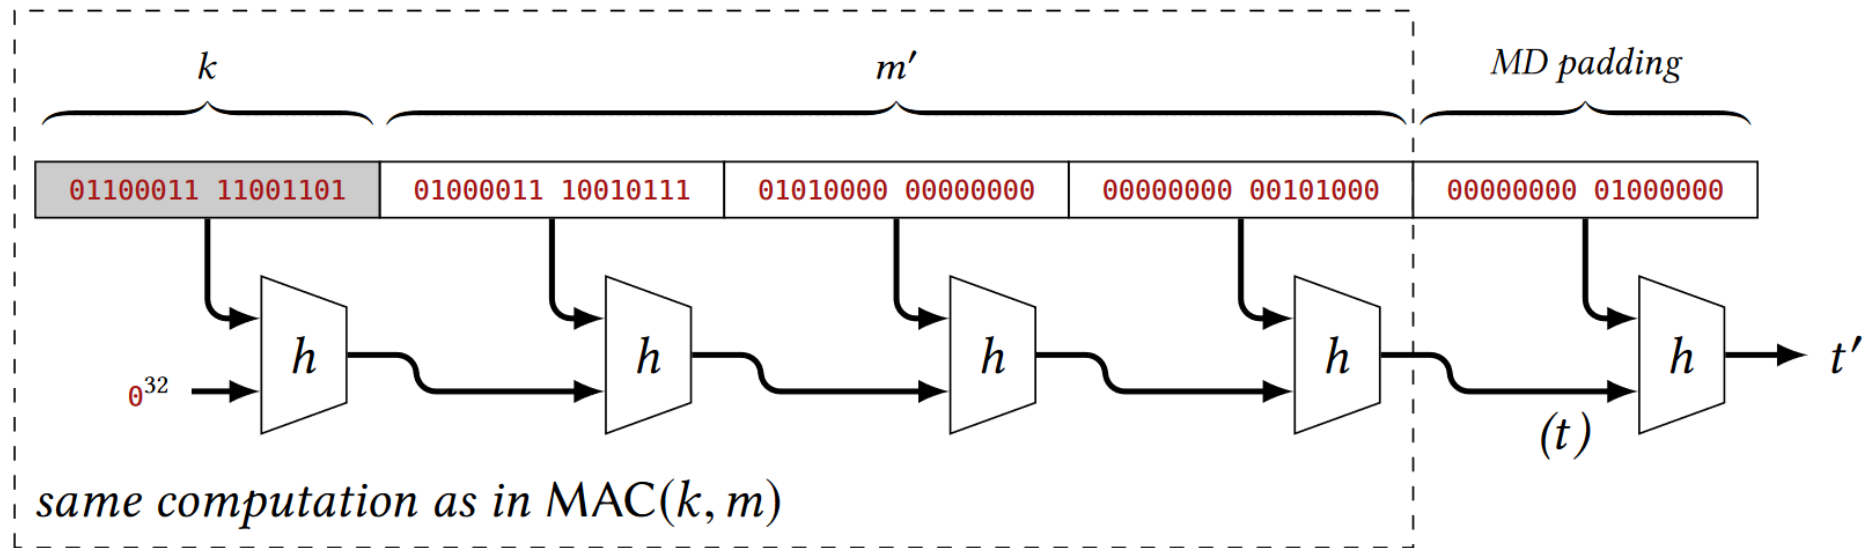


- Let $m' = 01000011\ 10010111\ 01010000\ 00000000\ 00000000\ 00101000$

Hash Functions vs. MACs: Length-Extension Attacks



- Let $m' = 01000011\ 10010111\ 01010000\ 00000000\ 00000000\ 00101000$



- Knowing the hash of $k||m$ allows you to also compute the hash of $k||m||p$.

Hash Functions vs. MACs: Length-Extension Attacks

- The Merkle-Damgård approach suffers from length-extension attacks because it outputs its **entire internal state**.
- The value t is **both** the **output of $H(k||m)$** as well as **the only information about $k||m$ needed to compute the last call to h in the computation $H(k||m||p)$** .
- **Solution 1:** In Merkle-Damgård, we compute $y_i = h(y_{i-1}||x_i)$ until reaching the final output y_{k+1} . Suppose instead that **we only output half of y_{k+1}** (the y_i values may need to be made longer in order for this to make sense). Then **just knowing half of y_{k+1} is not enough to predict what the hash output will be** in a length-extension scenario.
- The hash function SHA-3 was designed in this way (often called a “wide pipe” construction). One of the explicit design criteria of SHA-3 was that $H(k||m)$ would be a secure MAC.

Hash Functions vs. MACs: Length-Extension Attacks

- The Merkle-Damgård approach suffers from length-extension attacks because it outputs its **entire internal state**.
- The value t is **both** the **output of $H(k||m)$** as well as **the only information about $k||m$ needed to compute the last call to h in the computation $H(k||m||p)$** .
- **Solution 2:** by doing $H(k_2||H(k_1||m))$, with independent keys.
- This change is enough to mark the end of the input. This construction is known as NMAC.
- A closely related (and popular) construction called HMAC allows k_1 and k_2 to even be related in some way.