

# Chosen Ciphertext Attacks

# Chosen Ciphertext Attacks

- The CPA security definition considers only the information leaked to the adversary by **honestly-generated** ciphertexts.
- However, it does **not** consider what happens **when an adversary is allowed to inject its own maliciously crafted ciphertexts into an honest system.**
- If that happens, then **even a CPA-secure encryption scheme can fail in spectacular ways.**

# Padding Oracle Attacks

- Imagine a webserver that receives CBC-encrypted ciphertexts for processing.
- When receiving a ciphertext, the webserver decrypts it under the appropriate key and then **checks** whether the plaintext has **valid** X.923 padding (Data is padded with null bytes, except for the last byte of padding which indicates how many padding bytes there are. )

01 34 11 d9 81 88 05 57 1d 73 c3 00 00 00 00 05  $\Rightarrow$  *valid*

95 51 05 4a d6 5a a3 44 af b3 85 00 00 00 00 03  $\Rightarrow$  *valid*

71 da 77 5a 5e 77 eb a8 73 c5 50 b5 81 d5 96 01  $\Rightarrow$  *valid*

5b 1c 01 41 5d 53 86 4e e4 94 13 e8 7a 89 c4 71  $\Rightarrow$  *invalid*

d4 0d d8 7b 53 24 c6 d1 af 5f d6 f6 00 c0 00 04  $\Rightarrow$  *invalid*

# Padding Oracle Attacks

- No matter how the attacker comes by this information, we say that the attacker has access to a padding oracle:

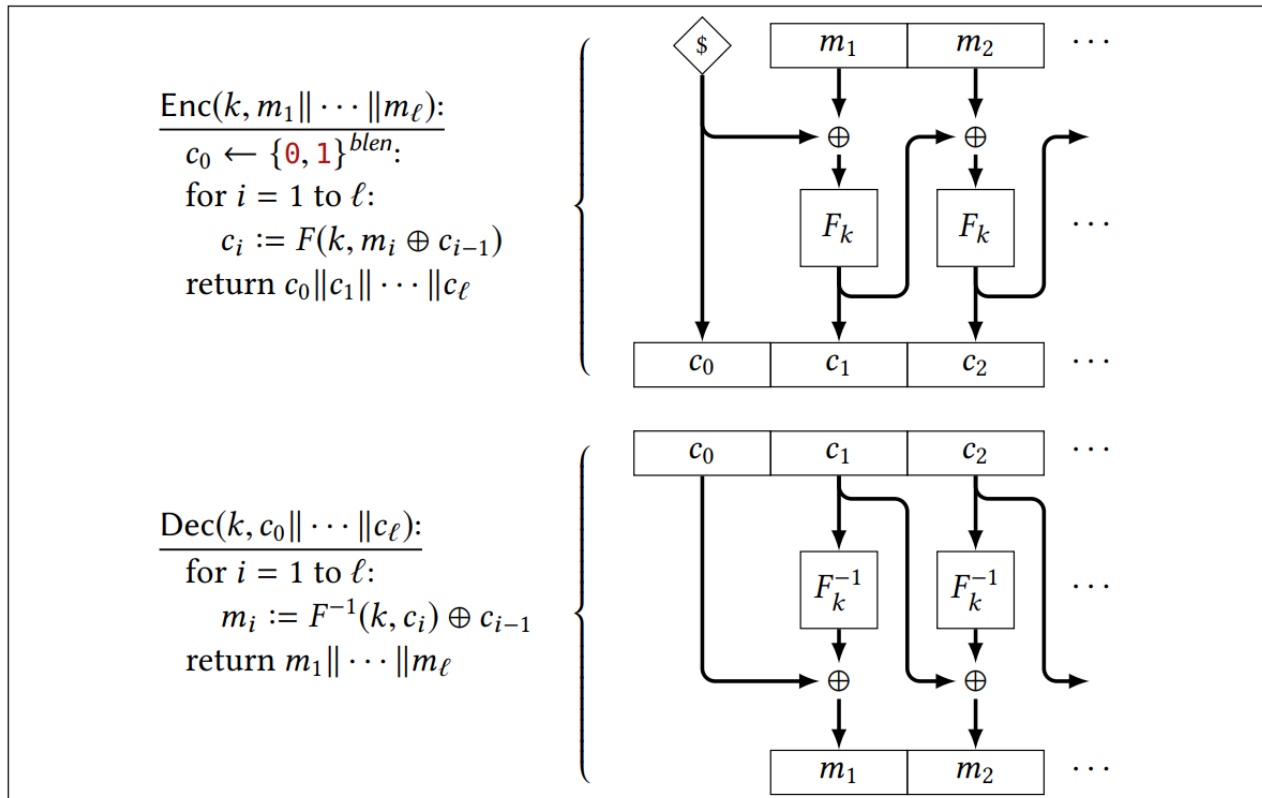
$\begin{array}{l} \text{PADDINGORACLE}(c): \\ \hline m := \text{Dec}(k, c) \\ \text{return VALIDPAD}(m) \end{array}$
--

- We call this a **padding oracle** because it **answers only one specific kind of question about the input**. In this case, the answer that it gives is always a single boolean value.
- We can show that an attacker who doesn't know the encryption key  $k$  can use a padding oracle alone to **decrypt any ciphertext of its choice!**

# Malleability of CBC Encryption

- Recall the definition of CBC decryption. If the ciphertext is  $c = c_0 \cdots c_\ell$  then the  $i$ th plaintext block is computed as:

$$m_i := F^{-1}(k, c_i) \oplus c_{i-1}$$



# Malleability of CBC Encryption

- Recall the definition of **CBC** decryption. If the ciphertext is  $c = c_0 \cdots c_\ell$  then the  $i$ th plaintext block is computed as:

$$m_i := F^{-1}(k, c_i) \oplus c_{i-1}$$

- Two consecutive blocks  $(c_{i-1}, c_i)$  taken in isolation are a **valid** encryption of  $m_i$ . This fact allows the attacker to focus on **decrypting a single block at a time**.
- Xoring a ciphertext block** with a known value (say,  $x$ ) has the effect of **xoring the corresponding plaintext block by the same value**. In other words, for all  $x$ , the ciphertext  $(c_{i-1} \oplus x, c_i)$  decrypts to  $m_i \oplus x$ :  
$$\text{Dec}(k, (c_{i-1} \oplus x, c_i)) = F^{-1}(k, c_i) \oplus (c_{i-1} \oplus x) = (F^{-1}(k, c_i) \oplus c_{i-1}) \oplus x = m_i \oplus x$$
- If we send such a ciphertext  $(c_{i-1} \oplus x, c_i)$  to the padding oracle, we would therefore learn whether  $m_i \oplus x$  is a (single block) with **valid** padding.

- Instead of thinking in terms of padding, it might be best to think of the oracle as telling you **whether  $m_i \oplus x$  ends in one of the suffixes** `01` , `00 02` , `00 00 03` , etc.

# Malleability of CBC Encryption

- By carefully choosing different values  $x$  and asking questions of this form to the padding oracle, we will show **how it is possible to learn all of  $m_i$** .

*// suppose  $c$  encrypts an (unknown) plaintext  $m_1 \parallel \dots \parallel m_\ell$   
// does  $m_i \oplus x$  end in one of the valid padding strings?*

CHECKXOR( $c, i, x$ ):

return PADDINGORACLE( $c_{i-1} \oplus x, c_i$ )

- Given a ciphertext  $c$  that encrypts an unknown message  $m$ , we can see that an adversary can **generate another ciphertext** whose contents are related to  $m$  **in a predictable way**.
- This property of an encryption scheme is called **malleability**.

# Learning the Last Byte of a Block

- How to use CHECKXOR to determine the last byte of a plaintext block  $m$ .
- Case 1: second-to-last byte of  $m$  is nonzero.
  - Try every possible byte  $b$  and ask whether  $m \oplus b$  has valid padding.
  - Only  $m \oplus b$  ends in byte 01 has valid padding.
  - Therefore, if  $b$  is the candidate byte that succeeds (i.e.,  $m \oplus b$  has valid padding) then the last byte of  $m$  must be  $b \oplus 01$ .

*Using LEARNLASTBYTE to learn the last byte of a plaintext block:*

$\dots$	$\boxed{a0} \boxed{42} \boxed{??}$	$m = \text{unknown plaintext block}$
$\oplus$	$\dots \boxed{00} \boxed{00} \boxed{b}$	$b = \text{byte that causes oracle to return true}$
<hr/>		
$=$	$\dots \boxed{a0} \boxed{42} \boxed{01}$	$\text{valid padding} \Leftrightarrow \boxed{b} \oplus \boxed{??} = \boxed{01}$
		$\Leftrightarrow \boxed{??} = \boxed{01} \oplus \boxed{b}$



# Learning the Last Byte of a Block

- How to use CHECKXOR to determine the last byte of a plaintext block  $m$ .
- Case 2: second-to-last byte of  $m$  is zero. Then  $m \oplus b$  will have valid padding for several candidate values of  $b$ :

*Using LEARNLASTBYTE to learn the last byte of a plaintext block:*

Diagram illustrating a brute-force attack on a 3-byte ciphertext. The diagram shows two parallel calculations.

**Left Calculation:**

$$\begin{array}{r} \dots \text{a0 } 00 \text{ ??} \\ \oplus \dots 00 \text{ } 00 \text{ } b_1 \\ \hline = \dots \text{a0 } 00 \text{ } 01 \end{array}$$

**Right Calculation:**

$$\begin{array}{r} \dots \text{a0 } 00 \text{ ??} \\ \oplus \dots 00 \text{ } 00 \text{ } b_2 \\ \hline = \dots \text{a0 } 00 \text{ } 02 \end{array}$$

Both results are marked as *true*.

**Second Calculation (Changing the second byte of the ciphertext to 01):**

$$\begin{array}{r} \dots \text{a0 } 01 \text{ ??} \\ \oplus \dots 00 \text{ } 01 \text{ } b_1 \\ \hline = \dots \text{a0 } 01 \text{ } 01 \end{array}$$

**Right Calculation (Changing the second byte of the ciphertext to 01):**

$$\begin{array}{r} \dots \text{a0 } 01 \text{ ??} \\ \oplus \dots 00 \text{ } 01 \text{ } b_2 \\ \hline = \dots \text{a0 } 01 \text{ } 02 \end{array}$$

The result for  $b_1$  is marked as *true*, and the result for  $b_2$  is marked as *false*.

Conclusion: *only one causes oracle to return true*

$\Rightarrow ?? = b_1 \oplus 01$

# Learning the Last Byte of a Block

- How to use CHECKXOR to determine the last byte of a plaintext block  $m$ .
- Case 2: second-to-last byte of  $m$  is zero. Then  $m \oplus b$  will have valid padding for several candidate values of  $b$ :
  - Whenever **more than one** candidate  $b$  value yields valid padding, we know that the **second-to-last byte of  $m$  is zero** (in fact, by counting **the number of successful candidates**, we can **know exactly how many zeroes precede the last byte of  $m$** ).
  - If the second-to-last byte of  $m$  is zero, then the second-to-last byte of  $m \oplus \text{01 } b$  is **nonzero**.
  - The only way for both strings  $m \oplus \text{01 } b$  and  $m \oplus b$  to have valid padding is when  $m \oplus b$  ends in byte **01**.

# Learning Other Bytes of a Block

- Suppose we **know the last 3 bytes of a plaintext block**. We would like to use the padding oracle to discover **the 4th-to-last byte**.
- In the worst case, this subroutine makes 256 queries to the padding oracle.

*Using `LEARNPREVBYTE` to learn the 4th-to-last byte when the last 3 bytes of the block are already known.*

	...	?? a0 42 3c	$m = \text{partially unknown plaintext block}$
$\oplus$	...	00 00 00 04	$p = \text{string ending in } 04$
$\oplus$	...	00 a0 42 3c	$s = \text{known bytes of } m$
$\oplus$	...	b 00 00 00	$y = \text{candidate byte } b \text{ shifted into place}$
<hr/>			
=	...	00 00 00 04	$\text{valid padding} \Leftrightarrow ?? = b$

# What Went Wrong?

- CBC encryption (in fact, every encryption scheme we've seen so far) has a property called **malleability**. Given an encryption  $c$  of an **unknown** plaintext  $m$ , **it is possible to generate another ciphertext  $c'$  whose contents are related to  $m$  in a predictable way.**
  - In the case of CBC encryption, if ciphertext  $c_0 || \dots || c_\ell$  encrypts a plaintext  $m_1 || \dots || m_\ell$ , then ciphertext  $(c_{i-1} \oplus x, c_i)$  encrypts the related plaintext  $m_i \oplus x$ .
- Decryption **has no impact on CPA security!** But the padding oracle setting **involved the Dec algorithm**.
- The attack makes 256 queries per byte of plaintext, so it costs about  **$256\ell$  queries for a plaintext of  $\ell$  bytes**. **Brute-forcing** the entire plaintext would cost  **$256^\ell$**  since that's how many  $\ell$ -byte plaintexts there are. So the attack is **exponentially better than brute force**.

# Defining CCA Security

- How can we possibly **anticipate** every kind of partial information that might make its way to the adversary in every possible usage of the encryption scheme?
- Let's just allow the adversary to **totally decrypt** arbitrary ciphertexts of its choice.
- Simply providing **unrestricted** Dec access to the adversary **cannot** lead to a **reasonable security definition**.
- Allow the adversary to ask for the decryption of **any** ciphertext, **except those produced in response to eavesdrop queries**.

# Defining CCA Security

**Definition** Let  $\Sigma$  be an encryption scheme. We say that  $\Sigma$  has security against chosen-ciphertext attacks (CCA security) if  $\mathcal{L}_{\text{cca-L}}^\Sigma \approx \mathcal{L}_{\text{cca-R}}^\Sigma$ , where:

$\mathcal{L}_{\text{cca-L}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$ $\mathcal{S} := \emptyset$
<u>EAVESDROP(<math>m_L, m_R \in \Sigma.\mathcal{M}</math>):</u> if $ m_L  \neq  m_R $ return <b>err</b> $c := \Sigma.\text{Enc}(k, m_L)$ $\mathcal{S} := \mathcal{S} \cup \{c\}$ return $c$
<u>DECRYPT(<math>c \in \Sigma.C</math>):</u> if $c \in \mathcal{S}$ return <b>err</b> return $\Sigma.\text{Dec}(k, c)$

$\mathcal{L}_{\text{cca-R}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$ $\mathcal{S} := \emptyset$
<u>EAVESDROP(<math>m_L, m_R \in \Sigma.\mathcal{M}</math>):</u> if $ m_L  \neq  m_R $ return <b>err</b> $c := \Sigma.\text{Enc}(k, m_R)$ $\mathcal{S} := \mathcal{S} \cup \{c\}$ return $c$
<u>DECRYPT(<math>c \in \Sigma.C</math>):</u> if $c \in \mathcal{S}$ return <b>err</b> return $\Sigma.\text{Dec}(k, c)$

# An Example

- Consider the adversary below attacking the CCA security of CBC mode (with block length  $blen$ )

$\mathcal{A}$
$c = c_0 \  c_1 \  c_2 := \text{EAVESDROP}(\textcolor{red}{0}^{2blen}, \textcolor{red}{1}^{2blen})$
$m := \text{DECRYPT}(c_0 \  c_1)$
return $m \stackrel{?}{=} \textcolor{red}{0}^{blen}$

- If  $c_0 \| c_1 \| c_2$  encrypts  $m_1 \| m_2$ , then  $c_0 \| c_1$  encrypts  $m_1$ .

# Pseudorandom Ciphertexts

- **Definition** Let  $\Sigma$  be an encryption scheme. We say that  $\Sigma$  has pseudorandom ciphertexts in the presence of chosen-ciphertext attacks (CCA\$ security) if

$\mathcal{L}_{\text{cca\$-real}}^\Sigma \approx \mathcal{L}_{\text{cca\$-rand}}^\Sigma$ , where:

Just like for CPA security, if a scheme has CCA\$ security, then it also has CCA security, but not vice-versa.

$\mathcal{L}_{\text{cca\$-real}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$ $\mathcal{S} := \emptyset$
$\text{CTXT}(m \in \Sigma.\mathcal{M}):$ <hr/> $c := \Sigma.\text{Enc}(k, m)$ $\mathcal{S} := \mathcal{S} \cup \{c\}$ return $c$
$\text{DECRYPT}(c \in \Sigma.C):$ <hr/> if $c \in \mathcal{S}$ return <b>err</b> return $\Sigma.\text{Dec}(k, c)$

$\mathcal{L}_{\text{cca\$-rand}}^\Sigma$
$k \leftarrow \Sigma.\text{KeyGen}$ $\mathcal{S} := \emptyset$
$\text{CTXT}(m \in \Sigma.\mathcal{M}):$ <hr/> $c \leftarrow \Sigma.C( m )$ $\mathcal{S} := \mathcal{S} \cup \{c\}$ return $c$
$\text{DECRYPT}(c \in \Sigma.C):$ <hr/> if $c \in \mathcal{S}$ return <b>err</b> return $\Sigma.\text{Dec}(k, c)$



# A Simple CCA-Secure Scheme

- Let  $F$  be a strong pseudorandom permutation with block length  $blen = n + \lambda$ . Define the following encryption scheme with message space  $\mathcal{M} = \{0, 1\}^n$  :

KeyGen:

$k \leftarrow \{0, 1\}^\lambda$

return  $k$

Enc( $k, m$ ):

$r \leftarrow \{0, 1\}^\lambda$

return  $F(k, m||r)$

Dec( $k, c$ ):

$v := F^{-1}(k, c)$

return first  $n$  bits of  $v$

- We can informally reason about the security of this scheme as follows:
  - As long as the random value  $r$  does not repeat, all inputs to the PRP are distinct, and thus its outputs will therefore all look independently uniform.
  - For any other value  $c'$  that the adversary asks to be decrypted, the guarantee of a strong PRP is that the result will look independently random. In particular, the result will not depend on the choice of plaintexts used to generate challenge ciphertexts.

# Advanced Cryptography

(Provable Security)

Yi LIU

# A Simple CCA-Secure Scheme

KeyGen:

$k \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$

return  $k$

Enc( $k, m$ ):

$r \leftarrow \{\mathbf{0}, \mathbf{1}\}^\lambda$

return  $F(k, m\|r)$

Dec( $k, c$ ):

$v := F^{-1}(k, c)$

return first  $n$  bits of  $v$

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*

$\mathcal{L}_{\text{cca\$-real}}^\Sigma$

$$k \leftarrow \{0, 1\}^\lambda$$

$$\mathcal{S} := \emptyset$$

CTXT( $m$ ):

$$r \leftarrow \{0, 1\}^\lambda$$

$$c := F(k, m \| r)$$

$$\mathcal{S} := \mathcal{S} \cup \{c\}$$


$$\text{return } c$$

DECRYPT( $c \in \Sigma.C$ ):

$$\text{if } c \in \mathcal{S} \text{ return error}$$

$$\text{return first } n \text{ bits of } F^{-1}(k, c)$$

strong PRP security



$$\mathcal{S} := \emptyset$$

$$T, T_{\text{inv}} := \text{empty assoc. arrays}$$

CTXT( $m$ ):

$$r \leftarrow \{0, 1\}^\lambda$$

$$\text{if } T[m \| r] \text{ undefined:}$$

$$c \leftarrow \{0, 1\}^{\text{blen}} \setminus T.\text{values}$$

$$T[m \| r] := c; T_{\text{inv}}[c] := m \| r$$

$$c := T[m \| r]$$

$$\mathcal{S} := \mathcal{S} \cup \{c\}$$

$$\text{return } c$$

DECRYPT( $c \in \Sigma.C$ ):

$$\text{if } c \in \mathcal{S} \text{ return error}$$

$$\text{if } T_{\text{inv}}[c] \text{ undefined:}$$

$$m \| r \leftarrow \{0, 1\}^{\text{blen}} \setminus T_{\text{inv}}.\text{values}$$

$$T_{\text{inv}}[c] := m \| r; T[m \| r] := c$$

$$\text{return first } n \text{ bits of } T_{\text{inv}}[c]$$

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*

$\mathcal{S} := \emptyset$

$T, T_{inv} := \text{empty assoc. arrays}$

CTXT( $m$ ):

$r \leftarrow \{0, 1\}^\lambda$

if  $T[m||r]$  undefined:

$c \leftarrow \{0, 1\}^{blen} \setminus T.\text{values}$

$T[m||r] := c; T_{inv}[c] := m||r$

$c := T[m||r]$

$\mathcal{S} := \mathcal{S} \cup \{c\}$

return  $c$

DECRYPT( $c \in \Sigma.C$ ):

if  $c \in \mathcal{S}$  return **err**

if  $T_{inv}[c]$  undefined:

$m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv}.\text{values}$

$T_{inv}[c] := m||r; T[m||r] := c$

return first  $n$  bits of  $T_{inv}[c]$

To prove CCA\$-security, we must reach a hybrid in which the responses of CTXT are **uniform**.

In the current hybrid there are two properties in the way of this goal:

- The ciphertext values  $c$  are sampled from  $\{0, 1\}^{blen} \setminus T.\text{values}$ , **rather than**  $\{0, 1\}^{blen}$ .
- To show CCA\$ security, we must **remove the dependence** of DECRYPT on previous values given to CTXT.

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*


```

 $\mathcal{S} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
   $r \leftarrow \{0, 1\}^\lambda$ 
  if  $T[m||r]$  undefined:
     $c \leftarrow \{0, 1\}^{blen} \setminus T.\text{values}$ 
     $T[m||r] := c; T_{inv}[c] := m||r$ 
   $c := T[m||r]$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv}.\text{values}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
  
```

Add some book-keeping that is not used anywhere.



```

 $\mathcal{S} := \emptyset; \mathcal{R} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
   $r \leftarrow \{0, 1\}^\lambda$ 
  if  $T[m||r]$  undefined:
     $c \leftarrow \{0, 1\}^{blen} \setminus T.\text{values}$ 
     $T[m||r] := c; T_{inv}[c] := m||r$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
   $c := T[m||r]$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv}.\text{values}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
  
```

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*

$\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$   
 $T, T_{inv} := \text{empty assoc. arrays}$


CTXT( $m$ ):

$r \leftarrow \{0, 1\}^\lambda$   
 if  $T[m||r]$  undefined:  
 $c \leftarrow \{0, 1\}^{blen} \setminus T.\text{values}$   
 $T[m||r] := c; T_{inv}[c] := m||r$   
 $\mathcal{R} := \mathcal{R} \cup \{r\}$   
 $c := T[m||r]$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
 return  $c$

DECRYPT( $c \in \Sigma.C$ ):

if  $c \in \mathcal{S}$  return **err**  
 if  $T_{inv}[c]$  undefined:  
 $m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv}.\text{values}$   
 $T_{inv}[c] := m||r; T[m||r] := c$   
 $\mathcal{R} := \mathcal{R} \cup \{r\}$   
 return first  $n$  bits of  $T_{inv}[c]$

Apply replacement vs  
 without replacement three  
 separate times.



$\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$

$T, T_{inv} := \text{empty assoc. arrays}$

CTXT( $m$ ):

$r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$   
 if  $T[m||r]$  undefined:  
 $c \leftarrow \{0, 1\}^{blen}$   
 $T[m||r] := c; T_{inv}[c] := m||r$   
 $\mathcal{R} := \mathcal{R} \cup \{r\}$   
 $c := T[m||r]$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
 return  $c$

DECRYPT( $c \in \Sigma.C$ ):

if  $c \in \mathcal{S}$  return **err**  
 if  $T_{inv}[c]$  undefined:  
 $m||r \leftarrow \{0, 1\}^{blen}$   
 $T_{inv}[c] := m||r; T[m||r] := c$   
 $\mathcal{R} := \mathcal{R} \cup \{r\}$   
 return first  $n$  bits of  $T_{inv}[c]$

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*


```

 $\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
   $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
  if  $T[m||r]$  undefined:
     $c \leftarrow \{0, 1\}^{blen}$ 
     $T[m||r] := c; T_{inv}[c] := m||r$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
   $c := T[m||r]$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
  
```

The if-statement in CTXT  
is always taken.



```

 $\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
   $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
   $c \leftarrow \{0, 1\}^{blen}$ 
   $T[m||r] := c; T_{inv}[c] := m||r$ 
   $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
  
```



# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*

```
 $\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$   
 $T, T_{inv} := \text{empty assoc. arrays}$ 
```

CTXT( $m$ ):

```
 $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
```

```
 $c \leftarrow \{0, 1\}^{blen}$ 
```

```
 $T[m||r] := c; T_{inv}[c] := m||r$ 
```

```
 $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
```

```
 $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
```

```
return  $c$ 
```

DECRYPT( $c \in \Sigma.C$ ):

```
if  $c \in \mathcal{S}$  return err
```

```
if  $T_{inv}[c]$  undefined:
```


```
 $m||r \leftarrow \{0, 1\}^{blen}$ 
```

```
 $T_{inv}[c] := m||r; T[m||r] := c$ 
```

```
 $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
```

```
return first  $n$  bits of  $T_{inv}[c]$ 
```

- **No** line of code ever reads from  $T$
- The first line of DECRYPT returns **err** for  $c \in \mathcal{S}$ . So  $T_{inv}$  is not read.



```
 $\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$ 
```

```
 $T, T_{inv} := \text{empty assoc. arrays}$ 
```

CTXT( $m$ ):

```
 $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
```

```
 $c \leftarrow \{0, 1\}^{blen}$ 
```

```
//  $T[m||r] := c; T_{inv}[c] := m||r$ 
```

```
 $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
```

```
 $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
```

```
return  $c$ 
```

DECRYPT( $c \in \Sigma.C$ ):

```
if  $c \in \mathcal{S}$  return err
```

```
if  $T_{inv}[c]$  undefined:
```

```
 $m||r \leftarrow \{0, 1\}^{blen}$ 
```

```
 $T_{inv}[c] := m||r; T[m||r] := c$ 
```

```
 $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
```

```
return first  $n$  bits of  $T_{inv}[c]$ 
```

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*


```

 $\mathcal{S} := \emptyset; \quad \mathcal{R} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
   $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
   $c \leftarrow \{0, 1\}^{blen}$ 
  //  $T[m||r] := c; T_{inv}[c] := m||r$ 
   $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
     $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
    
```

It has no effect to simply remove all lines that refer to variable  $\mathcal{R}$ .



```

 $\mathcal{S} := \emptyset; \quad // \mathcal{R} := \emptyset$ 
 $T, T_{inv} := \text{empty assoc. arrays}$ 

CTXT( $m$ ):
  //  $r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$ 
   $c \leftarrow \{0, 1\}^{blen}$ 
  //  $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
   $\mathcal{S} := \mathcal{S} \cup \{c\}$ 
  return  $c$ 

DECRYPT( $c \in \Sigma.C$ ):
  if  $c \in \mathcal{S}$  return err
  if  $T_{inv}[c]$  undefined:
     $m||r \leftarrow \{0, 1\}^{blen}$ 
     $T_{inv}[c] := m||r; T[m||r] := c$ 
    //  $\mathcal{R} := \mathcal{R} \cup \{r\}$ 
  return first  $n$  bits of  $T_{inv}[c]$ 
    
```

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*


```
 $\mathcal{S} := \emptyset;$      $// \mathcal{R} := \emptyset$   
 $T, T_{inv} :=$  empty assoc. arrays
```

CTXT( $m$ ):

```
 $// r \leftarrow \{0, 1\}^\lambda \setminus \mathcal{R}$   
 $c \leftarrow \{0, 1\}^{blen}$   
 $// \mathcal{R} := \mathcal{R} \cup \{r\}$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
return  $c$ 
```

DECRYPT( $c \in \Sigma.C$ ):

```
if  $c \in \mathcal{S}$  return err  
if  $T_{inv}[c]$  undefined:  
   $m||r \leftarrow \{0, 1\}^{blen}$   
   $T_{inv}[c] := m||r; T[m||r] := c$   
   $// \mathcal{R} := \mathcal{R} \cup \{r\}$   
return first  $n$  bits of  $T_{inv}[c]$ 
```



```
 $\mathcal{S} := \emptyset$ 
```

```
 $T, T_{inv} :=$  empty assoc. arrays
```

CTXT( $m$ ):

```
 $c \leftarrow \{0, 1\}^{blen}$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
return  $c$ 
```

DECRYPT( $c \in \Sigma.C$ ):


```
if  $c \in \mathcal{S}$  return err  
if  $T_{inv}[c]$  undefined:  
   $m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv}.values$   
   $T_{inv}[c] := m||r; T[m||r] := c$   
return first  $n$  bits of  $T_{inv}[c]$ 
```

# A Simple CCA-Secure Scheme

**Claim** If  $F$  is a strong PRP then the construction has CCA\$ security (and therefore CCA security).

*proof*

$\mathcal{S} := \emptyset$   
 $T, T_{inv} := \text{empty assoc. arrays}$   
CTXT( $m$ ):  
 $c \leftarrow \{0, 1\}^{blen}$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
 return  $c$   
DECRYPT( $c \in \Sigma.C$ ):  
 if  $c \in \mathcal{S}$  return **err**  
 if  $T_{inv}[c]$  undefined:  
 $m||r \leftarrow \{0, 1\}^{blen} \setminus T_{inv.values}$   
 $T_{inv}[c] := m||r; T[m||r] := c$   
 return first  $n$  bits of  $T_{inv}[c]$



$\mathcal{L}_{cca\$-rand}^\Sigma$   
 $k \leftarrow \{0, 1\}^\lambda$   
 $\mathcal{S} := \emptyset$   
CTXT( $m$ ):  
 $c \leftarrow \{0, 1\}^{blen}$   
 $\mathcal{S} := \mathcal{S} \cup \{c\}$   
 return  $c$   
DECRYPT( $c \in \Sigma.C$ ):  
 if  $c \in \mathcal{S}$  return **err**  
 return first  $n$  bits of  $F^{-1}(k, c)$

# One-Way Function

# One-Way Function

- A one-way function  $f: \{0, 1\}^* \rightarrow \{0, 1\}^*$  is **easy to compute**, yet **hard to invert**.
- **Easy to compute:**  $f$  is **computable in polynomial time**
- **Hard to invert:** It is **infeasible** for **any probabilistic polynomial-time algorithm** to invert  $f$ —that is, to find a preimage of a given value  $y$ —**except with negligible probability**.

# One-Way Function

Let  $f: \{0, 1\}^* \rightarrow \{0, 1\}^*$  be a function. Consider the following experiment defined for any algorithm  $\mathcal{A}$  and any value  $\lambda$  for the security parameter:

The inverting experiment **Invert** $_{\mathcal{A},f}(\lambda)$

1. Choose **uniform**  $x \in \{0, 1\}^\lambda$ , and compute  $y := f(x)$ .
  2. The algorithm  $\mathcal{A}$  is given  $1^\lambda$  and  $y$  as input, and outputs  $x'$ .
  3. The output of the experiment is defined to be **1** if  $f(x') = y$ , and **0** otherwise.
- We stress that  $\mathcal{A}$  need not find the original preimage  $x$ ; it suffices for  $\mathcal{A}$  to find any value  $x'$  for which  $f(x') = y = f(x)$ .
  - The security parameter  $1^\lambda$  is given to  $\mathcal{A}$  in the second step to stress that  $\mathcal{A}$  may **run in time polynomial in the security parameter  $\lambda$** , regardless of the length of  $y$ .

# One-Way Function

A function  $f: \{0, 1\}^* \rightarrow \{0, 1\}^*$  is one-way if the following two conditions hold:

1. **(Easy to compute:)** There **exists** a **polynomial-time** algorithm  $M_f$  computing  $f$ ; that is,  $M_f(x) = f(x)$  for all  $x$ .
2. **(Hard to invert:)** For **every probabilistic polynomial-time** algorithm  $\mathcal{A}$ , there is a **negligible** function  $\text{negl}$  such that

$$\Pr[\mathbf{Invert}_{\mathcal{A},f}(\lambda) = 1] \leq \text{negl}(\lambda)$$

$\Pr[\mathbf{Invert}_{\mathcal{A},f}(\lambda) = 1] \leq \text{negl}(\lambda)$  can be rewritten as

$$\Pr_{x \leftarrow \{0,1\}^\lambda} [\mathcal{A}(1^\lambda, f(x)) \in f^{-1}(f(x))] \leq \text{negl}(\lambda)$$



# Candidate One-Way Functions

- $f(p, q) = p \times q$  for **large equal-length** primes  $p$  and  $q$ .
- subset-sum problem
  - $f_{ss}(x_1, \dots, x_n, J) = (x_1, \dots, x_n, [\sum_{j \in J} x_j \bmod 2^n])$
  - where each  $x_i$  is an  $n$ -bit string interpreted as an integer, and  $J$  is an  $n$ -bit string interpreted as specifying a subset of  $\{1, \dots, n\}$
- $f_{p,g}(x) = [g^x \bmod p]$ 
  - $p$  is an  $\lambda$ -bit prime.  $g$  is the generator of  $\mathbb{Z}_p^*$
- SHA-2 or AES under the assumption that they are **collision resistant** or a **pseudorandom permutation**, respectively.

# Hard-Core Predicates

- One-way function: given  $y = f(x)$ , the value  $x$  cannot be computed in its entirety by **any polynomial-time** algorithm (except with negligible probability)
- Does this mean that **nothing** about  $x$  can be determined from  $f(x)$  in polynomial time?
- **No**. It is possible for  $f(x)$  to “**leak**” a lot of information about  $x$  even if  $f$  is one-way.
  - let  $g$  be a one-way function and define  $f(x_1, x_2) = (x_1, g(x_2))$ , where  $|x_1| = |x_2|$ .
  - It is **easy** to show that  $f$  is also a one-way function, even though it reveals half its input.
- For our applications, we will need to identify a **specific piece** of information about  $x$  that is “**hidden**” by  $f(x)$ . This motivates the notion of a ***hard-core predicate***.

# Hard-Core Predicates

A **hard-core predicate**  $hc : \{0, 1\}^* \rightarrow \{0, 1\}$  of a function  $f$  has the property that  $hc(x)$  is **hard** to compute with probability **significantly better than 1/2 given  $f(x)$** . (Since  $hc$  is a **boolean** function, it is always possible to compute  $hc(x)$  with probability 1/2 by random guessing.)

**Definition** A function  $hc : \{0, 1\}^* \rightarrow \{0, 1\}$  is a **hard-core predicate** of a function  $f$  if  $hc$  **can be computed in polynomial time**, and for **every probabilistic polynomial-time algorithm**  $\mathcal{A}$  there is a negligible function  $\text{negl}$  such that

$$\Pr_{x \leftarrow \{0,1\}^\lambda} [A(1^\lambda, f(x)) = hc(x)] \leq 1/2 + \text{negl}(\lambda)$$

where the probability is taken over the **uniform** choice of  $x$  in  $\{0, 1\}^\lambda$  and the **randomness** of  $\mathcal{A}$ .

We stress that  $hc(x)$  is **efficiently computable given  $x$**  (since the function  $hc$  can be computed in polynomial time); the definition requires that  $hc(x)$  is **hard to compute given  $f(x)$** .

# Simple ideas don't work

Consider the predicate  $hc(x) = \bigoplus_{i=1}^{\lambda} x_i$  where  $x_1, \dots, x_{\lambda}$  denote the **bits of  $x$** . Is this a hard-core predicate of **any** one-way function  $f$ ?

- If  $f$  **cannot** be inverted, then  $f(x)$  must **hide at least one of the bits  $x_i$**  of its preimage  $x$ , which would seem to imply that the XOR of all of the bits of  $x$  is hard to compute.
- This argument is **incorrect**
  - Let  $g$  be a one-way function and define  $f(x) = (g(x), \bigoplus_{i=1}^{\lambda} x_i)$ . It is not hard to show that  $f$  is **one-way**. However, it is clear that  $f(x)$  **does not hide** the value of  $hc(x) = \bigoplus_{i=1}^{\lambda} x_i$  because this is part of its output. Therefore,  $hc(x)$  is **not** a hard-core predicate of  $f$ .
  - For any fixed predicate  $hc$ , **there is a one-way function  $f$  for which  $hc$  is not a hard-core predicate of  $f$** .

# Trivial hard-core predicates

- Some functions have “trivial” hard-core predicates.
- Let  $f$  be the function that drops the last bit of its input (i.e.,  $f(x_1 \cdots x_\lambda) = x_1 \cdots x_{\lambda-1}$ ). It is hard to determine  $x_\lambda$  given  $f(x)$  since  $x_\lambda$  is independent of the output; thus,  $\text{hc}(x) = x_\lambda$  is a hard-core predicate of  $f$ .
- However,  $f$  is not one-way.
- Trivial hard-core predicates of this sort are of no use.

# Hard-Core Predicate from One-Way Functions

**Theorem** (Goldreich–Levin theorem) Assume one-way functions (resp., permutations) exist. Then there **exists** a one-way function (resp., permutation)  $g$  and a hard-core predicate  $gl$  of  $g$ .

Let  $f$  be a one-way function. Functions  $g$  and  $gl$  are constructed as follows: set  $g(x, r) = (f(x), r)$ , for  $|x| = |r|$ , and define

$$gl(x, r) = \bigoplus_{i=1}^{\lambda} x_i \cdot r_i$$

where  $x_i$  (resp.,  $r_i$ ) denotes the  $i$ th bit of  $x$  (resp.,  $r$ ).

Notice that if  $r$  is **uniform**, then  $gl(x, r)$  outputs the **XOR of a random subset of the bits of  $x$** . (When  $r_i = 1$  the bit  $x_i$  is included in the XOR, and otherwise it is not.) The Goldreich–Levin theorem thus states that if  $f$  is a one-way function then  $f(x)$  **hides the XOR of a random subset of the bits of  $x$** .

# Hard-Core Predicate from One-Way Functions

**Theorem** Let  $f$  be a one-way function and define  $g(x, r) = (f(x), r)$ , where  $|x| = |r|$ , and  $\text{gl}(x, r) = \bigoplus_{i=1}^{\lambda} x_i \cdot r_i$ . Then  $\text{gl}$  is a hard-core predicate of  $g$ .

We first show that if there exists a polynomial-time adversary  $\mathcal{A}$  that always correctly computes  $\text{gl}(x, r)$  given  $g(x, r) = (f(x), r)$ , then it is possible to invert  $f$  in polynomial time.

# Hard-Core Predicate from One-Way Functions

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a polynomial-time algorithm  $\mathcal{A}$  such that  $\mathcal{A}(f(x), r) = gl(x, r)$  for all  $\lambda$  and all  $x, r \in \{0, 1\}^\lambda$ , then there exists a polynomial-time algorithm  $\mathcal{A}'$  such that  $\mathcal{A}'(1^\lambda, f(x)) = x$  for all  $n$  and all  $x \in \{0, 1\}^\lambda$ .

If  $f$  is one-way, it is **impossible** for any probabilistic polynomial-time algorithm to invert  $f$  with **non-negligible probability**. Thus, we conclude that there is **no** polynomial-time algorithm that **always** correctly computes  $gl(x, r)$  from  $(f(x), r)$ .

This is a rather **weak** result that is **very far from our ultimate goal** of showing that  $gl(x, r)$  cannot be computed with probability **significantly better than  $1/2$**  given  $(f(x), r)$ .



# A Simple Case

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a polynomial-time algorithm  $\mathcal{A}$  such that  $\mathcal{A}(f(x), r) = gl(x, r)$  for all  $n$  and all  $x, r \in \{0, 1\}^\lambda$ , then there exists a polynomial-time algorithm  $\mathcal{A}'$  such that  $\mathcal{A}'(1^\lambda, f(x)) = x$  for all  $n$  and all  $x \in \{0, 1\}^\lambda$ .

*proof*

$\mathcal{A}'(1^\lambda, y)$  computes  $x_i = \mathcal{A}(y, e^i)$  for  $i = 1, \dots, \lambda$ , where  $e^i$  denotes the  $\lambda$ -bit string with 1 in the  $i$ th position and 0 everywhere else. Then  $\mathcal{A}'$  outputs  $x = x_1 \cdots x_\lambda$ . Clearly,  $\mathcal{A}'$  runs in polynomial time.

In the execution of  $\mathcal{A}'(1^\lambda, f(\hat{x}))$ ,

$$x_i = \mathcal{A}(f(\hat{x}), e^i) = gl(\hat{x}, e^i) = \bigoplus_{j=1}^{\lambda} \hat{x}_j \cdot e_j^i = \hat{x}_i$$

Thus  $x_i = \hat{x}_i$  for all  $i$ . The output of  $\mathcal{A}$  is the correct inverse.

# A More Involved Case

- We now show that it is **hard** for any **probabilistic polynomial-time** algorithm  $\mathcal{A}$  to compute  $\text{gl}(x, r)$  from  $(f(x), r)$  with probability **significantly better than  $3/4$** .
- We will again show that any such  $\mathcal{A}$  would imply the **existence** of a polynomial-time algorithm  $\mathcal{A}'$  that inverts  $f$  **with non-negligible probability**.
- For the simple case strategy:
  - It may be that  $\mathcal{A}$  **never** succeeds when  $r = e^i$  (although it may succeed, say, on all other values of  $r$ ).
  - $\mathcal{A}'$  **does not know** if the result  $\mathcal{A}(f(x), r)$  is equal to  $\text{gl}(x, r)$  or not.

# A More Involved Case

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a probabilistic polynomial-time algorithm  $\mathcal{A}$  and a polynomial  $p(\cdot)$  such that

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = gl(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$$

for infinitely many values of  $\lambda$ , then there exists a probabilistic polynomial-time algorithm  $\mathcal{A}'$  such that

$$\Pr_{x \leftarrow \{0,1\}^\lambda} \left[ \mathcal{A}' \left( 1^\lambda, f(x) \right) \in f^{-1}(f(x)) \right] \geq \frac{1}{4 \cdot p(\lambda)}$$

for infinitely many values of  $\lambda$ .

# Proof for the More Involved Case

The main observation underlying the proof of this proposition is that for **every**  $r \in \{0, 1\}^\lambda$ , the values  $\text{gl}(x, r \oplus e^i)$  and  $\text{gl}(x, r)$  **together** can be used to compute the  $i$ th bit of  $x$ .

$$\text{gl}(x, r) \oplus \text{gl}(x, r \oplus e^i) = \left( \bigoplus_{j=1}^{\lambda} x_j \cdot r_j \right) \oplus \left( \bigoplus_{j=1}^{\lambda} x_j \cdot (r_j \oplus e_j^i) \right) = x_i \cdot r_i \oplus (x_i \cdot \bar{r}_i) = x_i$$

If  $\mathcal{A}$  answers **correctly** on **both**  $(f(x), r)$  and  $(f(x), r \oplus e^i)$ , then  $\mathcal{A}'$  **can correctly compute**  $x_i$ .

$\mathcal{A}'$  knows only that  $\mathcal{A}$  answers correctly **with “high” probability**.

For this reason,  $\mathcal{A}'$  will use **multiple random values of**  $r$ , using each one to obtain an **estimate** of  $x_i$ , and then take the estimate occurring a **majority** of the time as its final guess for  $x_i$ .

As a preliminary step, we show that for many  $x$ 's the probability that  $\mathcal{A}$  answers **correctly** for **both**  $(f(x), r)$  and  $(f(x), r \oplus e^i)$ , when  $r$  is uniform, is sufficiently high.

# Proof for the More Involved Case

**Claim** Let  $\lambda$  be such that

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$$

Then there exists a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size **at least**  $\frac{1}{2p(\lambda)} \cdot 2^\lambda$  such that for every  $x \in S_\lambda$  it holds that for every  $x \in S_\lambda$ , it holds that

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{2p(\lambda)}$$

*proof*

Let  $\varepsilon(\lambda) = 1/p(\lambda)$ , and define  $S_\lambda \subseteq \{0, 1\}^\lambda$  to be the set of **all**  $x$ 's for which

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{\varepsilon(\lambda)}{2}$$

# Proof for the More Involved Case

**Claim** Let  $\lambda$  be such that  $\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$ . Then there exists a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size at least  $\frac{1}{2p(\lambda)} \cdot 2^\lambda$  such that for every  $x \in S_\lambda$  it holds that for every  $x \in S_\lambda$ , it holds that

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{2p(\lambda)}$$

$$\begin{aligned} \text{proof } \Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] &= \frac{1}{2^\lambda} \sum_{x \in \{0,1\}^\lambda} \Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \\ &= \frac{1}{2^\lambda} \sum_{x \in S_\lambda} \Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] + \frac{1}{2^\lambda} \sum_{x \notin S_\lambda} \Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \\ &\leq \frac{|S_\lambda|}{2^\lambda} + \frac{1}{2^\lambda} \cdot \sum_{x \notin S_\lambda} \left( \frac{3}{4} + \frac{\varepsilon(\lambda)}{2} \right) \leq \frac{|S_\lambda|}{2^\lambda} + \left( \frac{3}{4} + \frac{\varepsilon(\lambda)}{2} \right) \end{aligned}$$

$$\frac{3}{4} + \varepsilon(\lambda) \leq \frac{|S_\lambda|}{2^\lambda} + \left( \frac{3}{4} + \frac{\varepsilon(\lambda)}{2} \right)$$

Therefore,  $|S_\lambda| \geq \frac{\varepsilon(\lambda)}{2} \cdot 2^\lambda$ .

# Proof for the More Involved Case

**Claim** Let  $\lambda$  be such that  $\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$ . Then there exists a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size at least  $\frac{1}{2p(\lambda)} \cdot 2^\lambda$  such that for every  $x \in S_\lambda$  it holds that for every  $x \in S_\lambda$  and every  $i$ , it holds that  $\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r) \wedge \mathcal{A}(f(x), r \oplus e^i) = \text{gl}(x, r \oplus e^i)] \geq \frac{1}{2} + \frac{1}{p(\lambda)}$

*proof*

Let  $\varepsilon(\lambda) = 1/p(\lambda)$ , and take  $S_\lambda$  to be the set guaranteed by the previous claim. We have

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) \neq \text{gl}(x, r)] \leq \frac{1}{4} - \frac{\varepsilon(\lambda)}{2}$$
$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r \oplus e^i) \neq \text{gl}(x, r \oplus e^i)] \leq \frac{1}{4} - \frac{\varepsilon(\lambda)}{2}$$

# Union Bound

- $\Pr[E_1 \vee E_2] \leq \Pr[E_1] + \Pr[E_2]$
- $\Pr[\vee_{i=1}^k E_i] \leq \sum_{i=1}^k \Pr[E_i]$



# Proof for the More Involved Case

**Claim** Let  $\lambda$  be such that  $\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$ . Then there exists a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size at least  $\frac{1}{2p(\lambda)} \cdot 2^\lambda$  such that for every  $x \in S_\lambda$  it holds that for every  $x \in S_\lambda$  and **every**  $i$ , it holds that  $\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r) \wedge \mathcal{A}(f(x), r \oplus e^i) = \text{gl}(x, r \oplus e^i)] \geq \frac{1}{2} + \frac{1}{p(\lambda)}$

*proof*

Let  $\varepsilon(\lambda) = 1/p(\lambda)$ , and take  $S_\lambda$  to be the set guaranteed by the previous claim. We have

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) \neq \text{gl}(x, r)] \leq \frac{1}{4} - \frac{\varepsilon(\lambda)}{2}$$
$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r \oplus e^i) \neq \text{gl}(x, r \oplus e^i)] \leq \frac{1}{4} - \frac{\varepsilon(\lambda)}{2}$$

The probability that  $\mathcal{A}$  is **incorrect** on **either**  $\text{gl}(x, r)$  or  $\text{gl}(x, r \oplus e^i)$  is **at most**  $\frac{1}{2} - \varepsilon(\lambda)$ . The probability that **both are correct** is  $\frac{1}{2} + \varepsilon(\lambda)$ .

# Proof for the More Involved Case

For the rest of the proof we set  $\varepsilon(\lambda) = 1/p(\lambda)$  and consider only those values of  $\lambda$  for which

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{3}{4} + \varepsilon(\lambda)$$

The previous claim states that for an  $\varepsilon(n)/2$  fraction of inputs  $x$  (a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size at least  $\frac{1}{2p(\lambda)} \cdot 2^\lambda$ ), and any  $i$ , algorithm  $\mathcal{A}$  answers correctly on both  $(f(x), r)$  and  $(f(x), r \oplus e^i)$  with probability at least  $1/2 + \varepsilon(n)$  over uniform choice of  $r$ .

From now on we focus only on such values of  $x$ . We construct a probabilistic polynomial-time algorithm  $\mathcal{A}'$  that inverts  $f(x)$  with probability at least  $1/2$  when  $x \in S_\lambda$ .

$$\begin{aligned} & \Pr_{x \leftarrow \{0,1\}^\lambda} \left[ \mathcal{A}'(1^\lambda, f(x)) \in f^{-1}(f(x)) \right] \\ & \geq \Pr_{x \leftarrow \{0,1\}^\lambda} \left[ \mathcal{A}'(1^\lambda, f(x)) \in f^{-1}(f(x)) \mid x \in S_\lambda \right] \cdot \Pr_{x \leftarrow \{0,1\}^\lambda} [x \in S_\lambda] \geq \frac{1}{4 \cdot p(\lambda)} \end{aligned}$$

# Proof for the More Involved Case

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a probabilistic polynomial-time algorithm  $\mathcal{A}$  and a polynomial  $p(\cdot)$  such that

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = gl(x, r)] \geq \frac{3}{4} + \frac{1}{p(\lambda)}$$

for infinitely many values of  $\lambda$ , then there exists a probabilistic polynomial-time algorithm  $\mathcal{A}'$  such that

$$\Pr_{x \leftarrow \{0,1\}^\lambda} [\mathcal{A}'(1^\lambda, f(x)) \in f^{-1}(f(x))] \geq \frac{1}{4 \cdot p(\lambda)}$$

for infinitely many values of  $\lambda$ .

$$\begin{aligned} & \Pr_{x \leftarrow \{0,1\}^\lambda} [\mathcal{A}'(1^\lambda, f(x)) \in f^{-1}(f(x))] \\ & \geq \Pr_{x \leftarrow \{0,1\}^\lambda} [\mathcal{A}'(1^\lambda, f(x)) \in f^{-1}(f(x)) \mid x \in S_\lambda] \cdot \Pr_{x \leftarrow \{0,1\}^\lambda} [x \in S_\lambda] \geq \frac{1}{4 \cdot p(\lambda)} \end{aligned}$$

# Proof for the More Involved Case

Algorithm  $\mathcal{A}'$ , given as input  $1^\lambda$  and  $y$ , works as follows:

1. For  $i = 1, \dots, \lambda$  do:

**Repeatedly** choose a uniform  $r \in \{0, 1\}^\lambda$  and compute  $\mathcal{A}(y, r) \oplus \mathcal{A}(y, r \oplus e^i)$  as an “estimate” for the  $i$ th bit of the preimage of  $y$ . After doing this sufficiently many times, **let  $x_i$  be the “estimate” that occurs a majority of the time.**

2. Output  $x = x_1 \cdots x_\lambda$ .

By obtaining **sufficiently many estimates** and letting  $x_i$  be the **majority value**,  $\mathcal{A}'$  can ensure that  $x_i$  is equal to  $\text{gl}(\hat{x}, e^i)$  with probability at least  **$1 - \frac{1}{2\lambda}$** .

# Chernoff bound

**Proposition** Fix  $\varepsilon > 0$  and  $b \in \{0, 1\}$ , and let  $\{X_i\}_{i=1,\dots,m}$  be independent 0/1-random variables with  $\Pr[X_i = b] = \frac{1}{2} + \varepsilon$  for all  $i$ . The probability that their majority value is **not**  $b$  is at most  $e^{-\varepsilon^2 m/2}$ .

# Proof for the More Involved Case

Algorithm  $\mathcal{A}'$ , given as input  $1^\lambda$  and  $y$ , works as follows:

1. For  $i = 1, \dots, \lambda$  do:

Repeatedly choose a uniform  $r \in \{0, 1\}^\lambda$  and compute  $\mathcal{A}(y, r) \oplus \mathcal{A}(y, r \oplus e^i)$  as an “estimate” for the  $i$ th bit of the preimage of  $y$ . After doing this sufficiently many times, let  $x_i$  be the “estimate” that occurs a majority of the time.

2. Output  $x = x_1 \cdots x_\lambda$ .

By obtaining sufficiently many estimates and letting  $x_i$  be the majority value,  $\mathcal{A}'$  can ensure that  $x_i$  is equal to  $\text{gl}(\hat{x}, e^i)$  with probability **at least**  $1 - \frac{1}{2\lambda}$ .

- polynomially many estimates suffice

We have that for each  $i$  the value  $x_i$  computed by  $\mathcal{A}'$  is **incorrect** with probability **at most**  $\frac{1}{2\lambda}$ .

A union bound thus shows that  $\mathcal{A}'$  is **incorrect** for some  $i$  with probability at most  $\lambda \cdot 1/2\lambda = 1/2$ , and thus correctly inverts  $y$ —with probability **at least**  $1/2$

# The Full Proof

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a probabilistic polynomial-time algorithm  $\mathcal{A}$  and a polynomial  $p(\cdot)$  such that

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = gl(x, r)] \geq \frac{1}{2} + \frac{1}{p(\lambda)}$$

for infinitely many values of  $\lambda$ , then there exists a probabilistic polynomial-time algorithm  $\mathcal{A}'$  and a polynomial  $p'(\cdot)$  such that

$$\Pr_{x \leftarrow \{0,1\}^\lambda} \left[ \mathcal{A}' \left( 1^\lambda, f(x) \right) \in f^{-1}(f(x)) \right] \geq \frac{1}{p'(\lambda)}$$

for infinitely many values of  $\lambda$ .

# The Full Proof

**Proposition** Let  $f$  and  $gl$  be as before. If there exists a probabilistic polynomial-time algorithm  $\mathcal{A}$  and a polynomial  $p(\cdot)$  such that

$$\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = gl(x, r)] \geq \frac{1}{2} + \frac{1}{p(\lambda)}$$

for infinitely many values of  $\lambda$ , then there exists a probabilistic polynomial-time algorithm  $\mathcal{A}'$  and a polynomial  $p'(\cdot)$  such that

$$\Pr_{x \leftarrow \{0,1\}^\lambda} \left[ \mathcal{A}' \left( 1^\lambda, f(x) \right) \in f^{-1}(f(x)) \right] \geq \frac{1}{p'(\lambda)}$$

for infinitely many values of  $\lambda$ .

*proof* Once again we set  $\varepsilon(n) = 1/p(\lambda)$  and consider only those values of  $\lambda$  for

which  $\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = gl(x, r)] \geq \frac{1}{2} + \varepsilon(\lambda)$



# The Full Proof

The following is **analogous** to the previous claim and is proved **in the same way**.

**Claim** Let  $\lambda$  be such that  $\Pr_{x,r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{1}{2} + \varepsilon(\lambda)$ .

Then there exists a set  $S_\lambda \subseteq \{0, 1\}^\lambda$  of size at least  $\frac{\varepsilon(\lambda)}{2} \cdot 2^\lambda$  such that for every  $x \in S_\lambda$  it holds that for every  $x \in S_\lambda$ , it holds that

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r)] \geq \frac{1}{2} + \frac{\varepsilon(\lambda)}{2}$$

# The Full Proof

If we start by trying to prove an **analogue** of claim for  $\mathcal{A}(f(x), r) = \text{gl}(x, r) \wedge \mathcal{A}(f(x), r \oplus e^i) = \text{gl}(x, r \oplus e^i)$ , we have for any  $i$

$$\Pr_{r \leftarrow \{0,1\}^\lambda} [\mathcal{A}(f(x), r) = \text{gl}(x, r) \wedge \mathcal{A}(f(x), r \oplus e^i) = \text{gl}(x, r \oplus e^i)] \geq \varepsilon(\lambda)$$

All we can claim is that this estimate will be correct with probability **at least  $\varepsilon(\lambda)$** , **which may not be any better than taking a random guess!** We **cannot** claim that flipping the result **gives a good estimate**, either.

Instead, we design  $\mathcal{A}'$  so that it computes  $\text{gl}(x, r)$  and  $\text{gl}(x, r \oplus e^i)$  by invoking  $\mathcal{A}$  **only once**.

We do this by having  $\mathcal{A}'$  run  $\mathcal{A}(f(x), r \oplus e^i)$ , and then **simply “guessing” the value  $\text{gl}(x, r)$  itself**.

# The Full Proof

Instead, we design  $\mathcal{A}'$  so that it computes  $\text{gl}(x, r)$  and  $\text{gl}(x, r \oplus e^i)$  by invoking  $\mathcal{A}$  only once.

We do this by having  $\mathcal{A}'$  run  $A(f(x), r \oplus e^i)$ , and then simply “guessing” the value  $\text{gl}(x, r)$  itself.

The **naïve** way to do this would be to choose the  $r$ ’s independently, as before, and to have  $\mathcal{A}'$  make an independent guess of  $\text{gl}(x, r)$  for each value of  $r$ . But then the probability that all such guesses are correct would be negligible because polynomially many  $r$ ’s are used.

**Solution:**  $\mathcal{A}'$  can generate the  $r$ ’s in a pairwise-independent manner and make its guesses in a particular way so that with non-negligible probability all its guesses are correct.

# The Full Proof

**Solution:**  $\mathcal{A}'$  can generate the  $r$ 's in a **pairwise-independent** manner and make its guesses in a particular way so that **with non-negligible probability all its guesses are correct**.

In order to generate  $m$  values of  $r$ , we have  $\mathcal{A}'$  select  $\ell = \lceil \log(m + 1) \rceil$  **independent** and **uniformly** distributed strings  $s^1, \dots, s^\ell \in \{0, 1\}^\lambda$ . Then, for every nonempty subset  $I \subseteq \{1, \dots, \ell\}$ , we set  $r^I = \bigoplus_{i \in I} s^i$ .

Since there are  $2^\ell - 1$  **nonempty subsets**, this defines a collection of  $2^{\lceil \log(m+1) \rceil} - 1 \geq m$  **strings**.

The strings are **not independent**, but **they are pairwise independent**.

For every two subsets  $I \neq J$  there is an index  $j \in I \cup J$  such that  $j \notin I \cap J$ . Without loss of generality, assume  $j \notin I$ .

Then the value of  $s^j$  is **uniform** and **independent** of the value of  $r^I$ . Since  $s^j$  is included in the XOR that defines  $r^J$ , **this implies that  $r^J$  is uniform and independent of  $r^I$  as well**.

# The Full Proof

**Solution:**  $\mathcal{A}'$  can generate the  $r$ 's in a **pairwise-independent** manner and make its guesses in a particular way so that **with non-negligible probability** all its guesses are correct.

In order to generate  $m$  values of  $r$ , we have  $\mathcal{A}'$  select  $\ell = \lceil \log(m + 1) \rceil$  **independent** and **uniformly** distributed strings  $s^1, \dots, s^\ell \in \{0, 1\}^\lambda$ . Then, for every nonempty subset  $I \subseteq \{1, \dots, \ell\}$ , we set  $r^I = \bigoplus_{i \in I} s^i$ .

We now have the following **two** important observations:

- Given  $\text{gl}(x, s^1), \dots, \text{gl}(x, s^\ell)$ , it is possible to compute  $\text{gl}(x, r^I)$  for every subset  $I \subseteq \{1, \dots, \ell\}$ . This is because

$$\text{gl}(x, r^I) = \text{gl}\left(x, \bigoplus_{i \in I} s^i\right) = \bigoplus_{i \in I} \text{gl}(x, s^i)$$

- If  $\mathcal{A}'$  simply guesses the values of  $\text{gl}(x, s^1), \dots, \text{gl}(x, s^\ell)$  by choosing a uniform bit for each, then **all these guesses will be correct with probability  $1/2^\ell$** . If  $m$  is polynomial in the security parameter  $\lambda$ , then  **$1/2^\ell$  is not negligible**, and so **with non-negligible probability  $\mathcal{A}'$  correctly guesses all the values  $\text{gl}(x, s^1), \dots, \text{gl}(x, s^\ell)$** .

# The Full Proof

**Solution:**  $\mathcal{A}'$  can generate the  $r$ 's in a pairwise-independent manner and make its guesses in a particular way so that with non-negligible probability all its guesses are correct.

In order to generate  $m$  values of  $r$ , we have  $\mathcal{A}'$  select  $\ell = \lceil \log(m + 1) \rceil$  independent and uniformly distributed strings  $s^1, \dots, s^\ell \in \{0, 1\}^\lambda$ . Then, for every nonempty subset  $I \subseteq \{1, \dots, \ell\}$ , we set  $r^I = \bigoplus_{i \in I} s^i$ .

We can obtain  $m = \text{poly}(\lambda)$  uniform and pairwise-independent strings  $\{r^I\}$  along with correct values for  $\{\text{gl}(x, r^I)\}$  with non-negligible probability.

# The Full Proof

The full description of an algorithm  $\mathcal{A}'$  that receives inputs  $1^\lambda$ ,  $y$  and tries to compute an inverse of  $y$  as follows:

1. Set  $\ell = \lceil \log(m + 1) \rceil$ , where  $m = 2\lambda/\varepsilon(\lambda)^2$ .
2. Choose **uniform, independent**  $s^1, \dots, s^\ell \in \{0, 1\}^\ell$  and  $\sigma^1, \dots, \sigma^\ell \in \{0, 1\}$ .
3. For every nonempty subset  $I \subseteq \{1, \dots, \ell\}$ , compute  $r^I = \bigoplus_{i \in I} s^i$  and  $\sigma^I = \bigoplus_{i \in I} \sigma^i$ .
4. For  $i = 1, \dots, \lambda$  do:
  1. For every nonempty subset  $I \subseteq \{1, \dots, \ell\}$ , set  $x_i^I = \sigma^I \oplus \mathcal{A}(y, r^I \oplus e^i)$ .
  2. Set  $x_i = \text{majority}_I\{x_i^I\}$
5. Output  $x = x_1 \cdots x_\ell$ .

# The Full Proof

It remains to compute the **probability** that  $\mathcal{A}'$  outputs  $x \in f^{-1}(y)$ .

We focus on  $\lambda$  as before, and assume  $y = f(\hat{x})$  for some  $\hat{x} \in S_\ell$ .

As noted earlier, **with non-negligible probability all these guesses are correct**; we show that **conditioned on this event**,  $\mathcal{A}'$  outputs  $x = \hat{x}$  with probability at least  $1/2$ .

Assume  $\sigma_i = \text{gl}(\hat{x}, s^i)$  for all  $i$ . Then  $\sigma^I = \text{gl}(\hat{x}, r^I)$  for all  $I$ . Fix an index  $i \in \{1, \dots, \lambda\}$  and **consider the probability that  $\mathcal{A}'$  obtains the correct value  $x_i = \hat{x}_i$** .

For any nonempty  $I$  we have  $\mathcal{A}(y, r^I \oplus e^i) = \text{gl}(\hat{x}, r^I \oplus e^i)$  with probability at least  $1/2 + \varepsilon(\lambda)/2$  over choice of  $r$  **because  $\hat{x} \in S_\lambda$  and  $r^I \oplus e^i$  is uniformly distributed**.

Thus, for any nonempty subset  $I$  we have  $\Pr[x_i^I = \hat{x}_i] \geq \frac{1}{2} + \varepsilon(\lambda)/2$ .



# The Full Proof

It remains to compute the **probability** that  $\mathcal{A}'$  outputs  $x \in f^{-1}(y)$ .

Thus, for any nonempty subset  $I$  we have  $\Pr[x_i^I = \hat{x}_i] \geq \frac{1}{2} + \varepsilon(n)/2$ .

Moreover, the  $\{x_i^I\}_{I \subseteq \{1, \dots, \ell\}}$  are **pairwise independent** because the  $\{r^I\}_{I \subseteq \{1, \dots, \ell\}}$  (and hence the  $\{r^I \oplus e^i\}_{I \subseteq \{1, \dots, \ell\}}$ ) are **pairwise independent**.

# A Tool

**Proposition** Fix  $\varepsilon > 0$  and  $b \in \{0, 1\}$ , and let  $\{X_i\}$  be pairwise-independent, 0/1-random variables for which  $\Pr[X_i = b] \geq 1/2 + \varepsilon$  for all  $i$ . Consider the process in which  $m$  values  $X_1, \dots, X_m$  are recorded and  $X$  is set to the value that occurs a strict majority of the time. Then

$$\Pr[X \neq b] \leq \frac{1}{4 \cdot \varepsilon^2 \cdot m}$$

# The Full Proof

It remains to compute the probability that  $\mathcal{A}'$  outputs  $x \in f^{-1}(y)$ .

Thus, for any nonempty subset  $I$  we have  $\Pr[x_i^I = \hat{x}_i] \geq \frac{1}{2} + \varepsilon(n)/2$ .

Moreover, the  $\{x_i^I\}_{I \subseteq \{1, \dots, \ell\}}$  are **pairwise independent** because the  $\{r^I\}_{I \subseteq \{1, \dots, \ell\}}$  (and hence the  $\{r^I \oplus e^i\}_{I \subseteq \{1, \dots, \ell\}}$ ) are **pairwise independent**.

$$\Pr[x_i \neq \hat{x}_i] \leq \frac{1}{4 \cdot \left(\frac{\varepsilon(\lambda)}{2}\right)^2 \cdot (2^\ell - 1)} \leq \frac{1}{4 \cdot \left(\frac{\varepsilon(\lambda)}{2}\right)^2 \cdot \left(\frac{2\lambda}{\varepsilon(\lambda)^2}\right)} \leq 1/2\lambda$$

The above holds for all  $i$ , so by applying a **union bound** we see that the probability that  $x_i \neq \hat{x}_i$  for some  $i$  is at most  $1/2$ . That is,  $x_i = \hat{x}_i$  for all  $i$  (and hence  $x = \hat{x}$ ) with probability at least  $1/2$ .

# The Full Proof

Putting everything together:

With probability at least  $\varepsilon(n)/2$  we have  $\hat{x} \in S_\ell$ . All the guesses  $\sigma_i$  are correct with probability at least  $\frac{1}{2^\ell} \geq \frac{1}{2 \cdot \left(\frac{2\lambda}{\varepsilon(\lambda)^2} + 1\right)} > \frac{\varepsilon(\lambda)^2}{5\lambda}$  for  $\lambda$  sufficiently large.

Conditioned on both the above,  $\mathcal{A}'$  outputs  $x = \hat{x}$  with probability at least  $1/2$ .

The overall probability with which  $\mathcal{A}'$  inverts its input is thus at least  $\varepsilon(\lambda)^3/20\lambda = 1/(20 \lambda p(\lambda)^3)$  for infinitely many  $\lambda$ .

$20 \lambda p(\lambda)^3$  is a **polynomial** in  $\lambda$ , this proves the proposition.